

Architectural documentation, views and tradeoffs, patterns

Claire Le Goues
October 3, 2019

Administrivia

Learning Goals

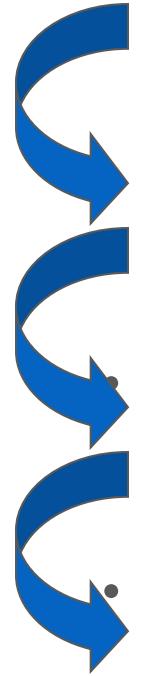
- Distinguish software architecture from (object-oriented) software design
- Use notation and views to describe the architecture suitable to the purpose, and document architectures clearly and without ambiguity.
- Use diagrams to understand systems and reason about tradeoffs.
- Understand the utility of architectural patterns and tactics, and give a couple of examples.

ARCHITECTURE VS OBJECT-LEVEL DESIGN

One slide on a thing from 214: Design

- Design process (analysis, design, implementation)
- Design goals (cohesion, coupling, information hiding, design for reuse, ...)
- Design patterns (what they are, for what they are useful, how they are described)
- Frameworks and libraries (reuse strategies)

Levels of Abstraction

- 
- Requirements
 - high-level “what” needs to be done

Architecture (High-level design)

- high-level “how”, mid-level “what”

OO-Design (Low-level design, e.g. design patterns)

- mid-level “how”, low-level “what”

- Code

- low-level “how”

Design vs. Architecture

Design Questions

- How do I add a menu item in Eclipse?
- How can I make it easy to add menu items in Eclipse?
- What lock protects this data?
- How does Google rank pages?
- What encoder should I use for secure communication?
- What is the interface between objects?

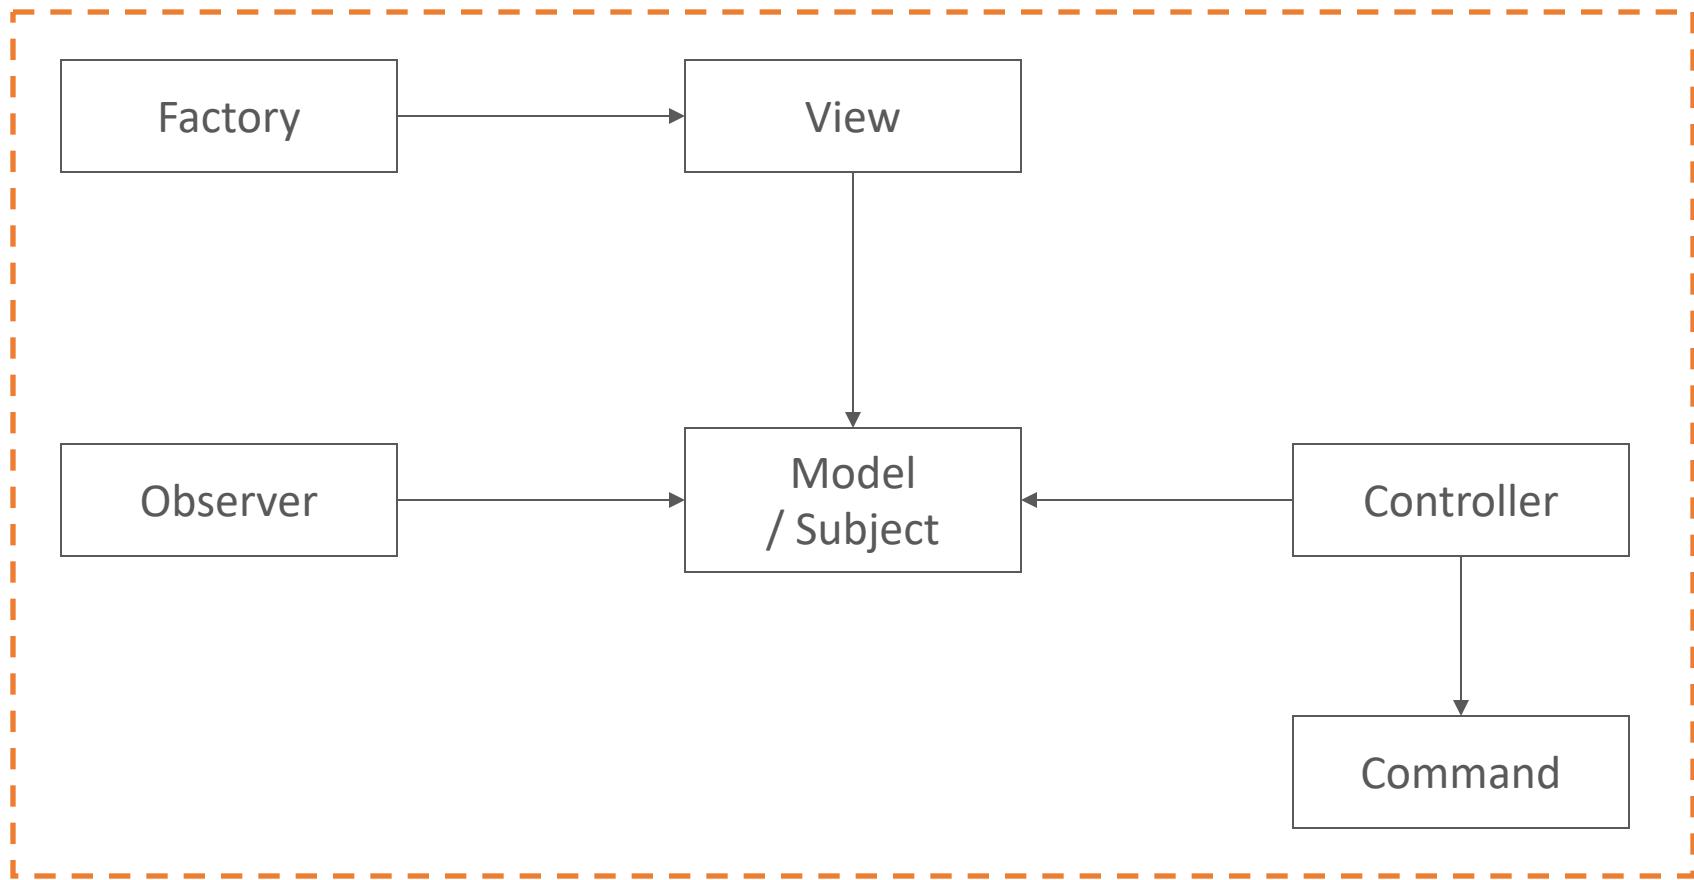
Architectural Questions

- How do I extend Eclipse with a plugin?
- What threads exist and how do they coordinate?
- How does Google scale to billions of hits per day?
- Where should I put my firewalls?
- What is the interface between subsystems?

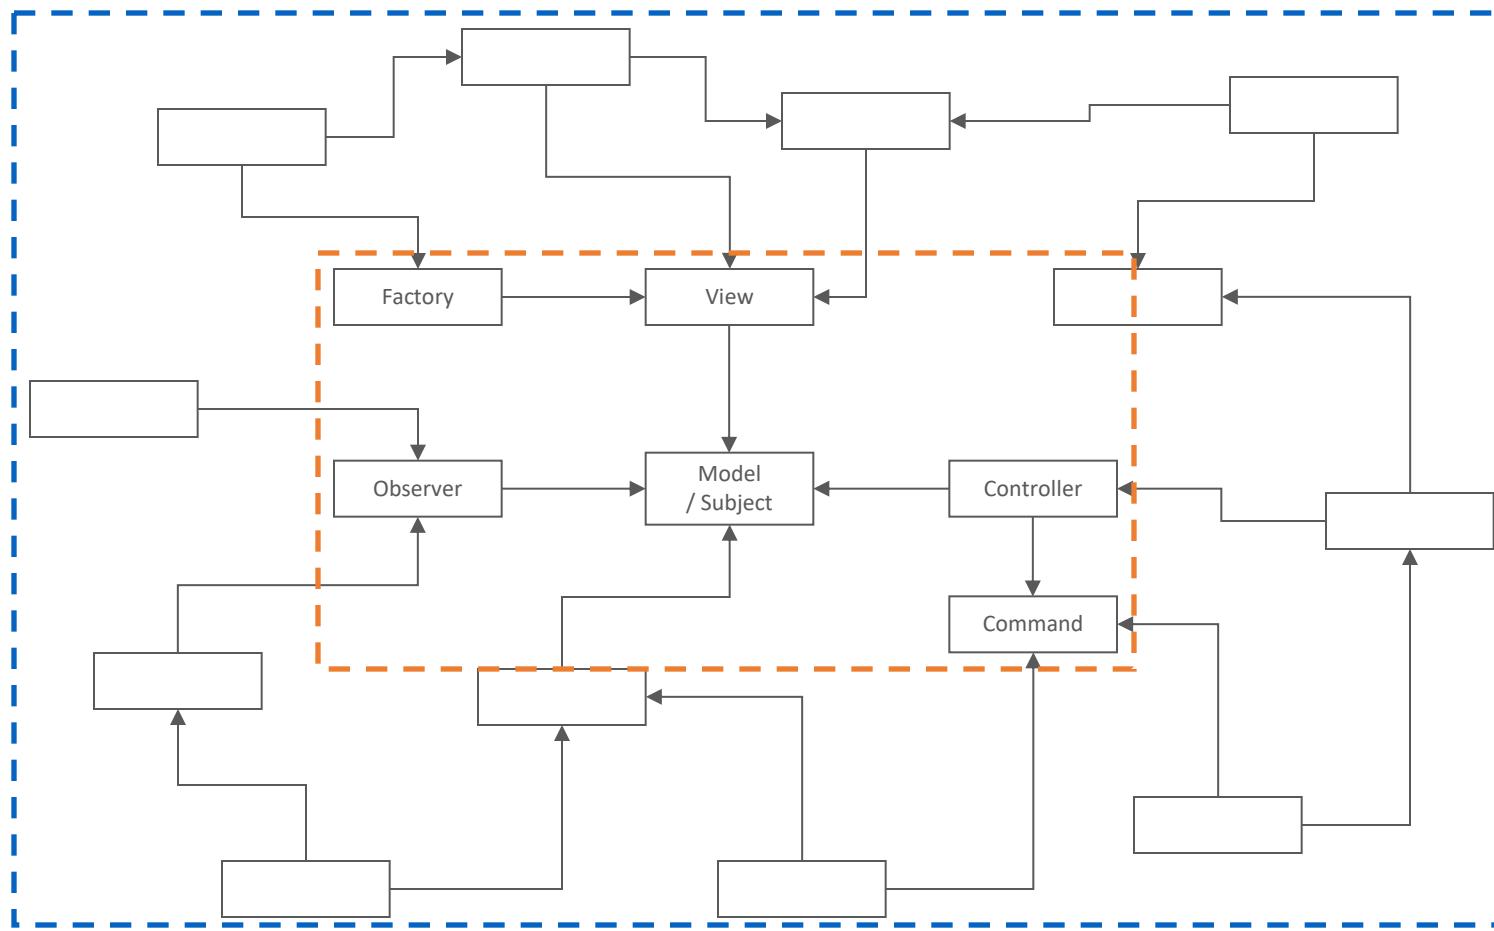
Objects

Model

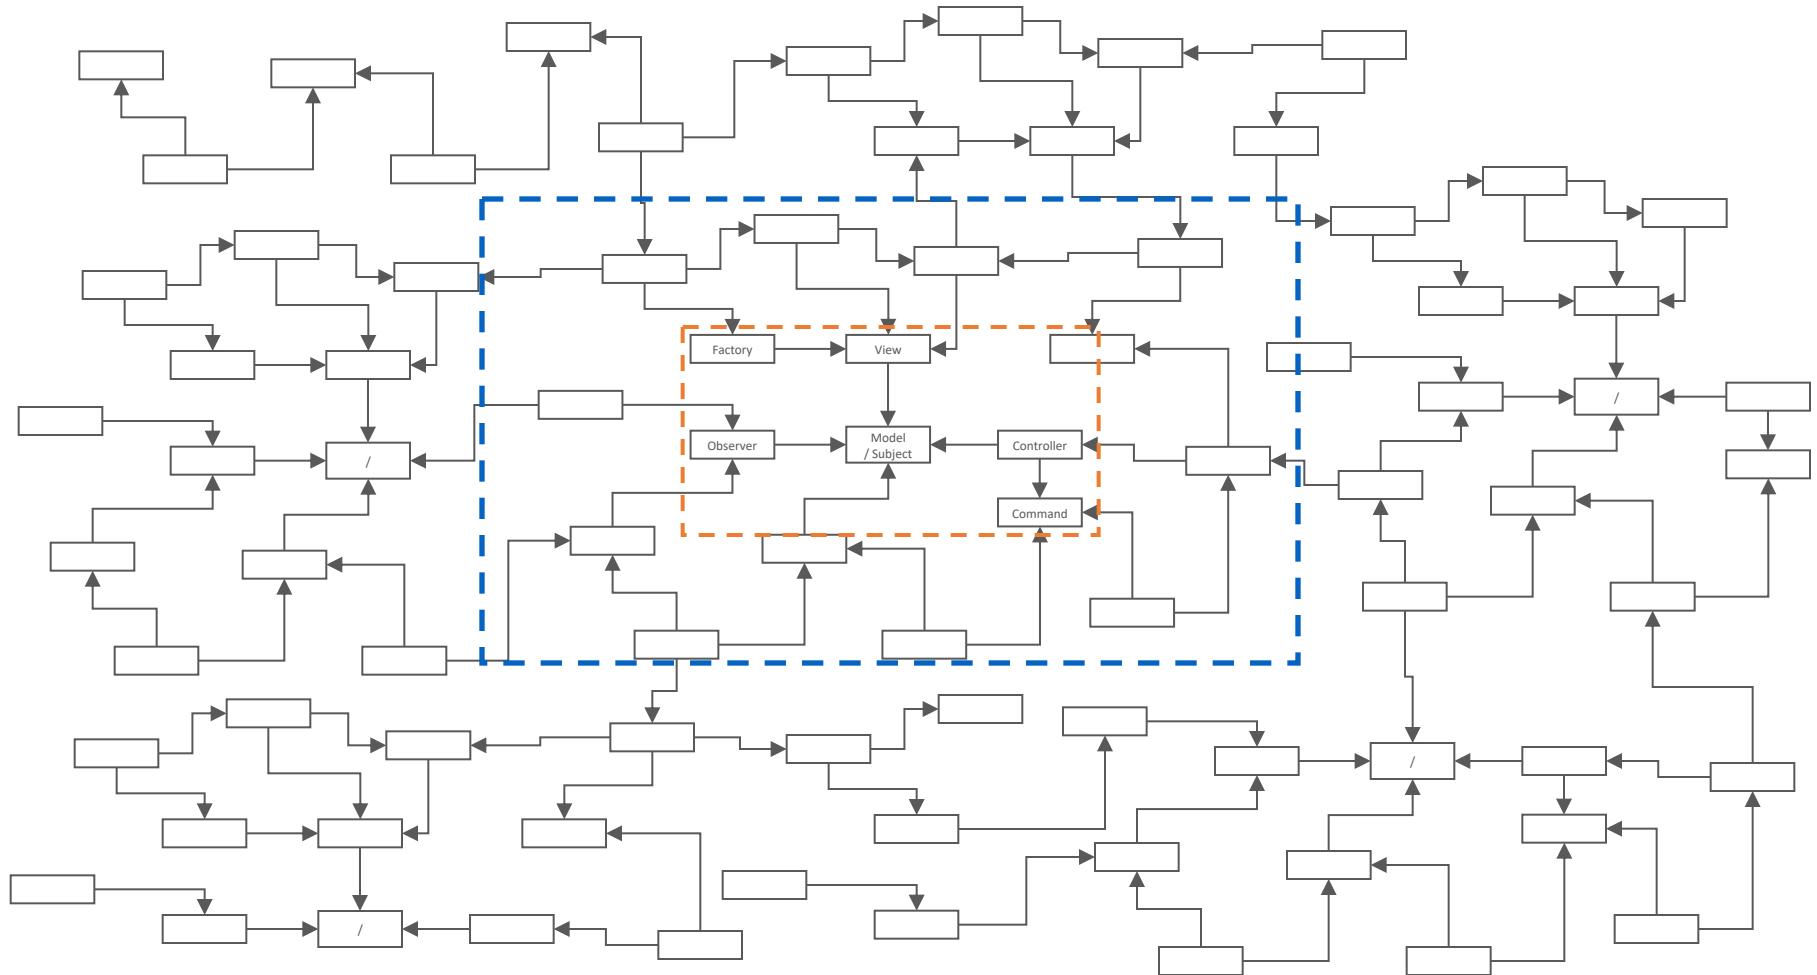
Design Patterns



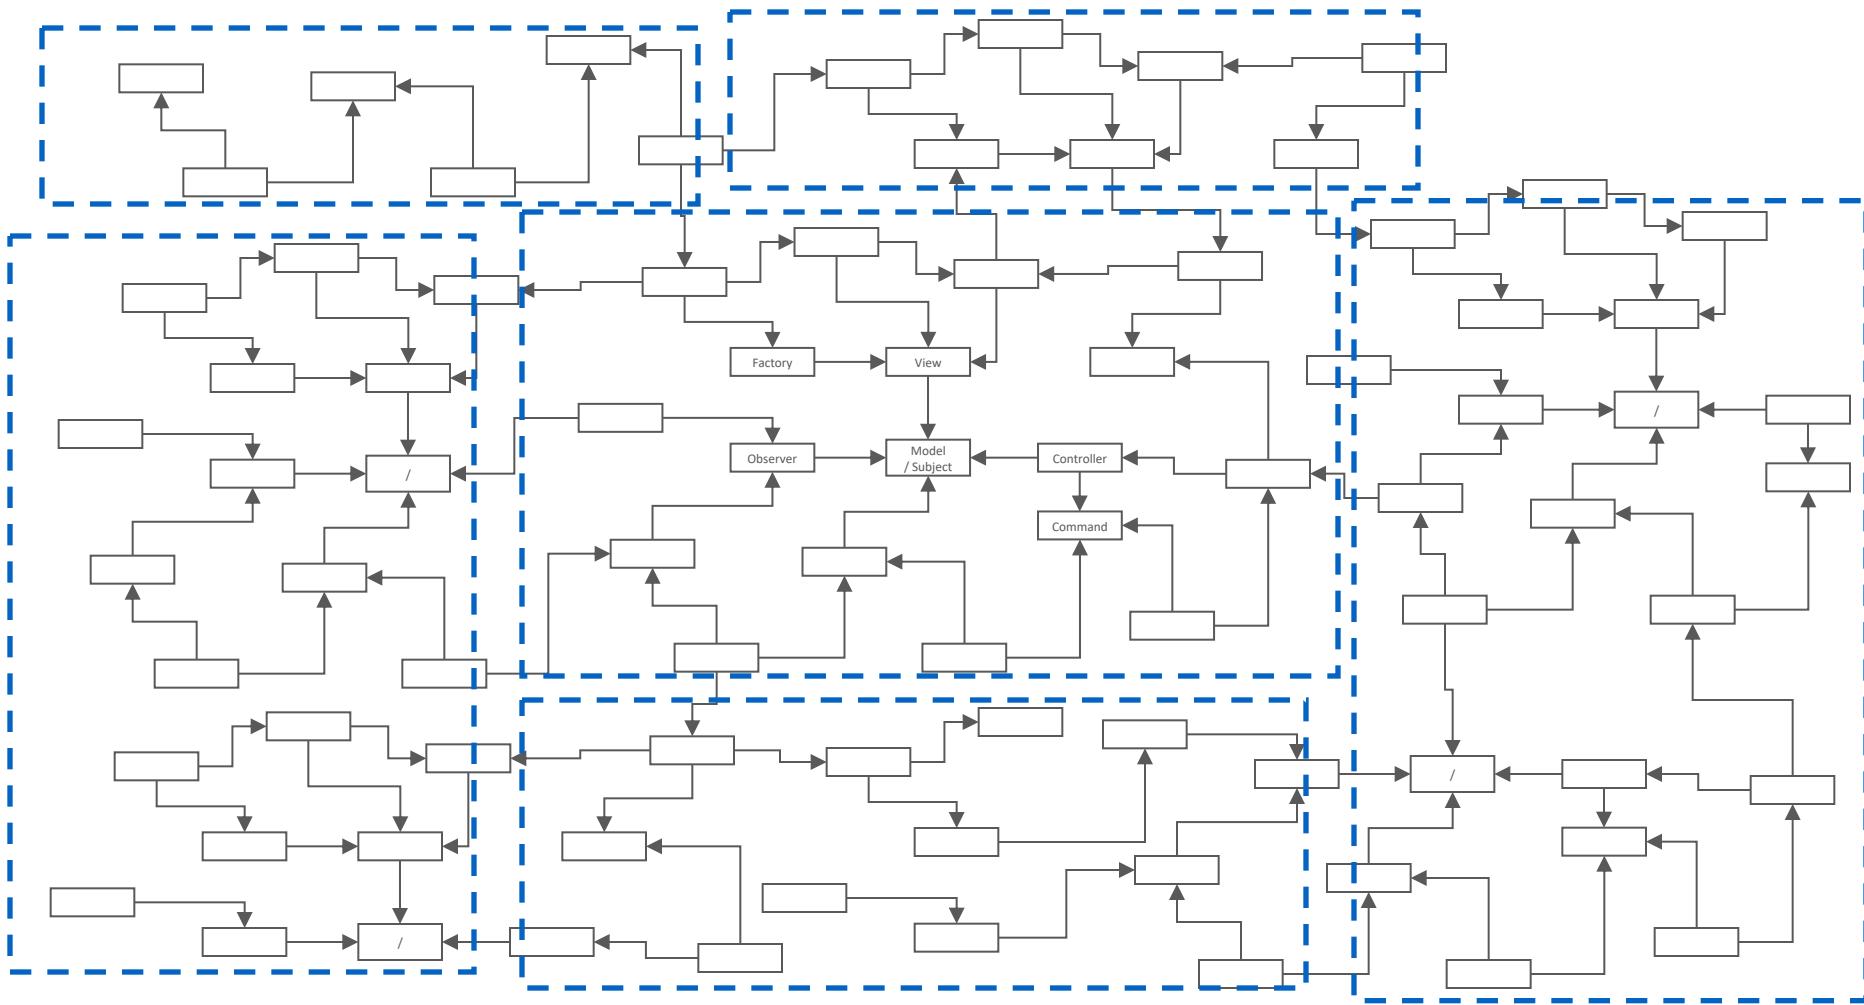
Design Patterns



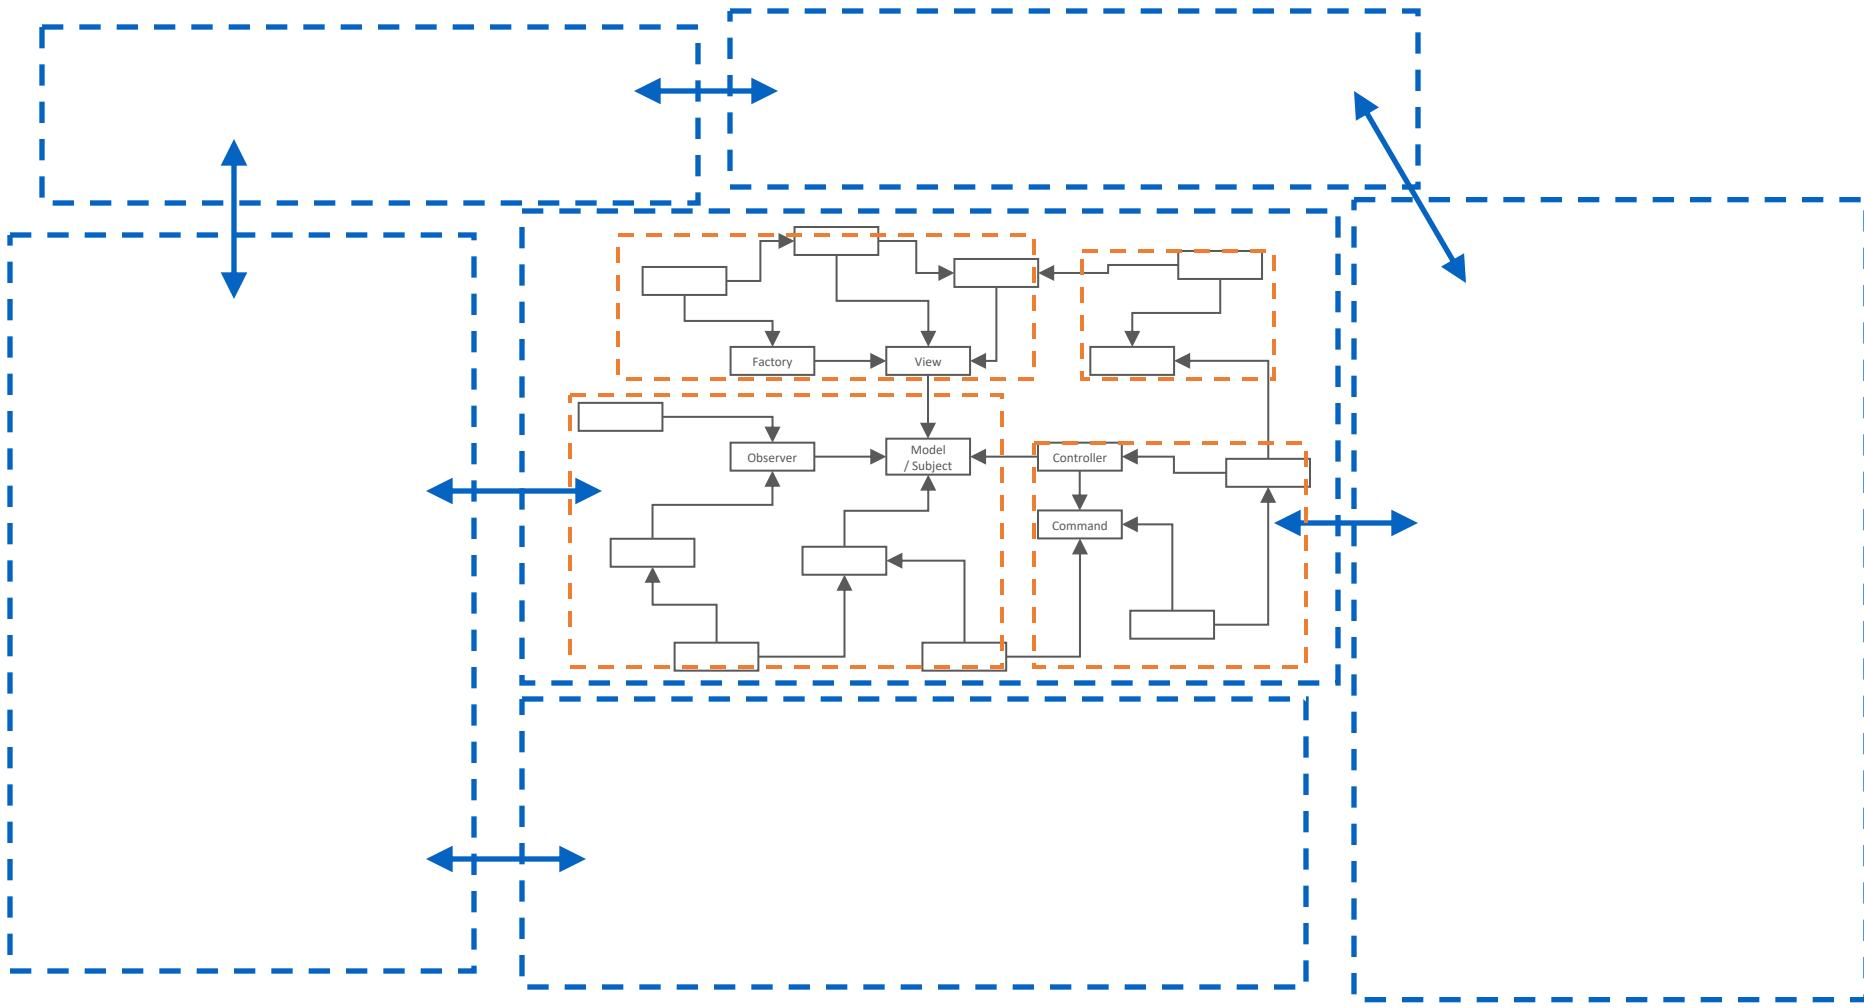
Design Patterns



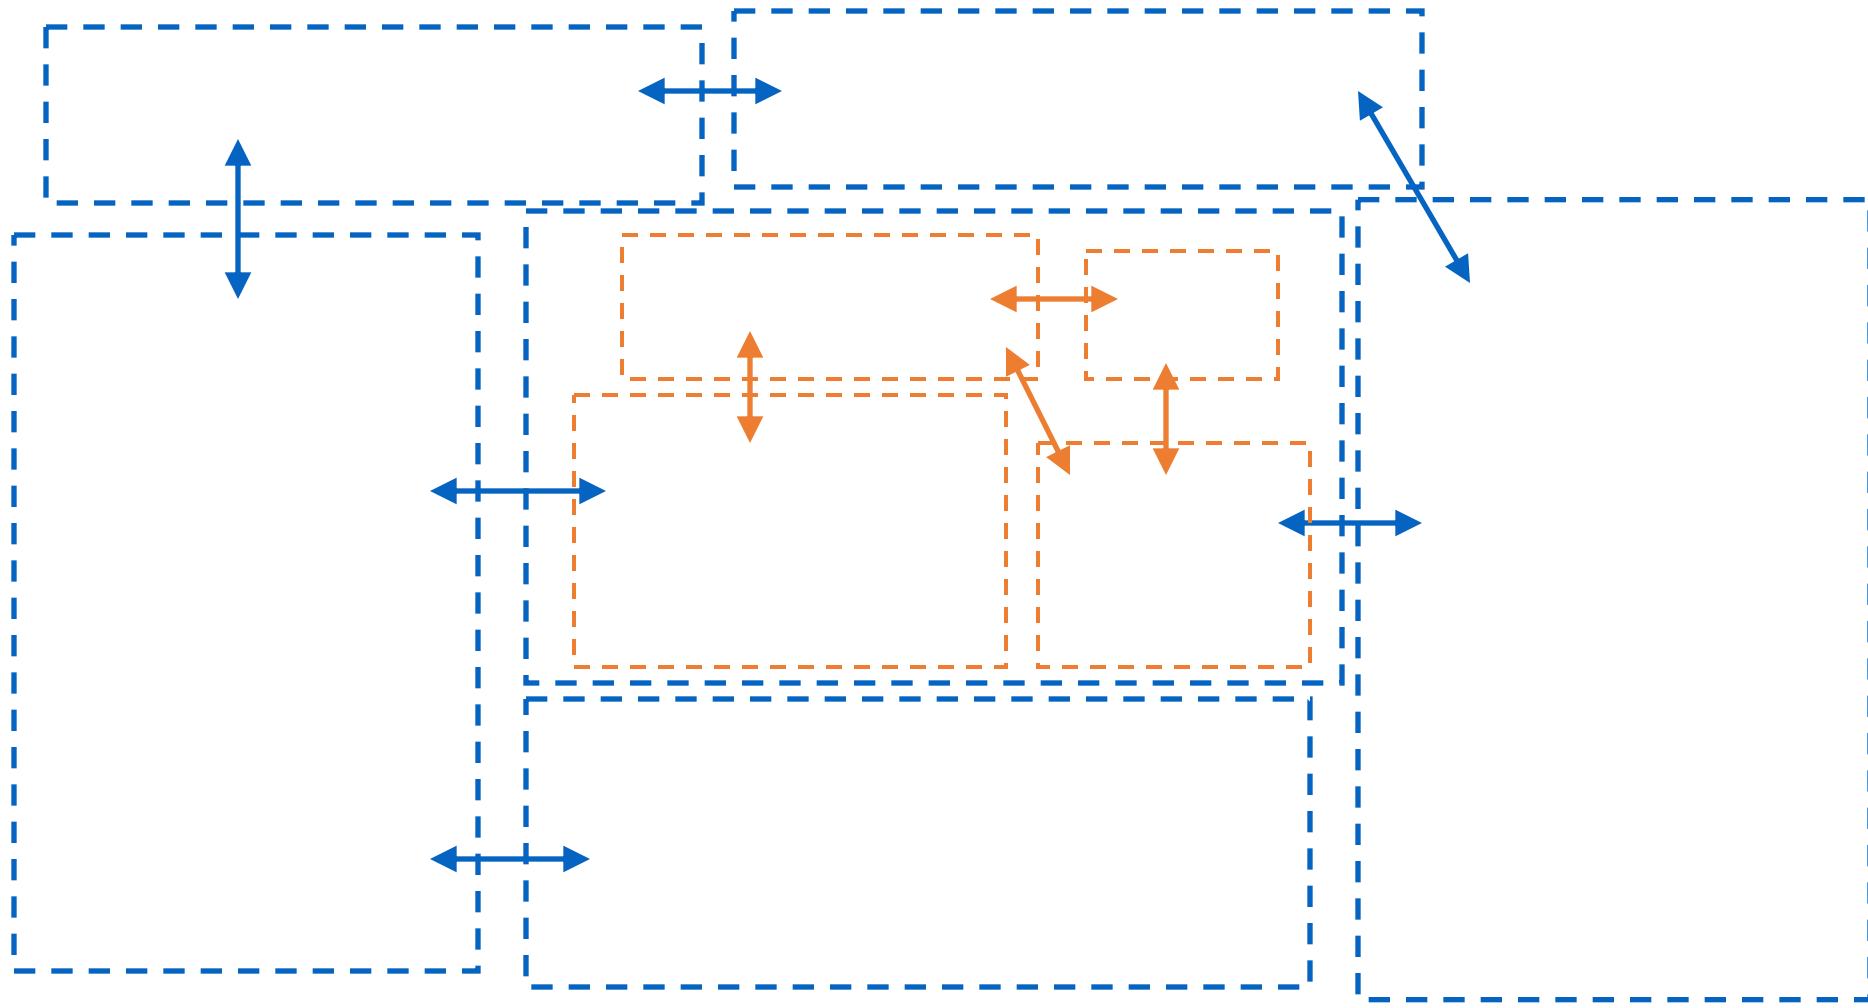
Architecture



Architecture



Architecture



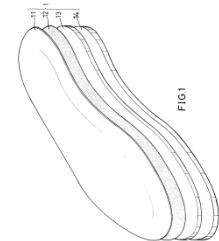
ARCHITECTURE DOCUMENTATION & VIEWS

Architecture Disentangled

Architecture as
structures and relations
(the actual system)



Architecture as
documentation
(representations of the system)

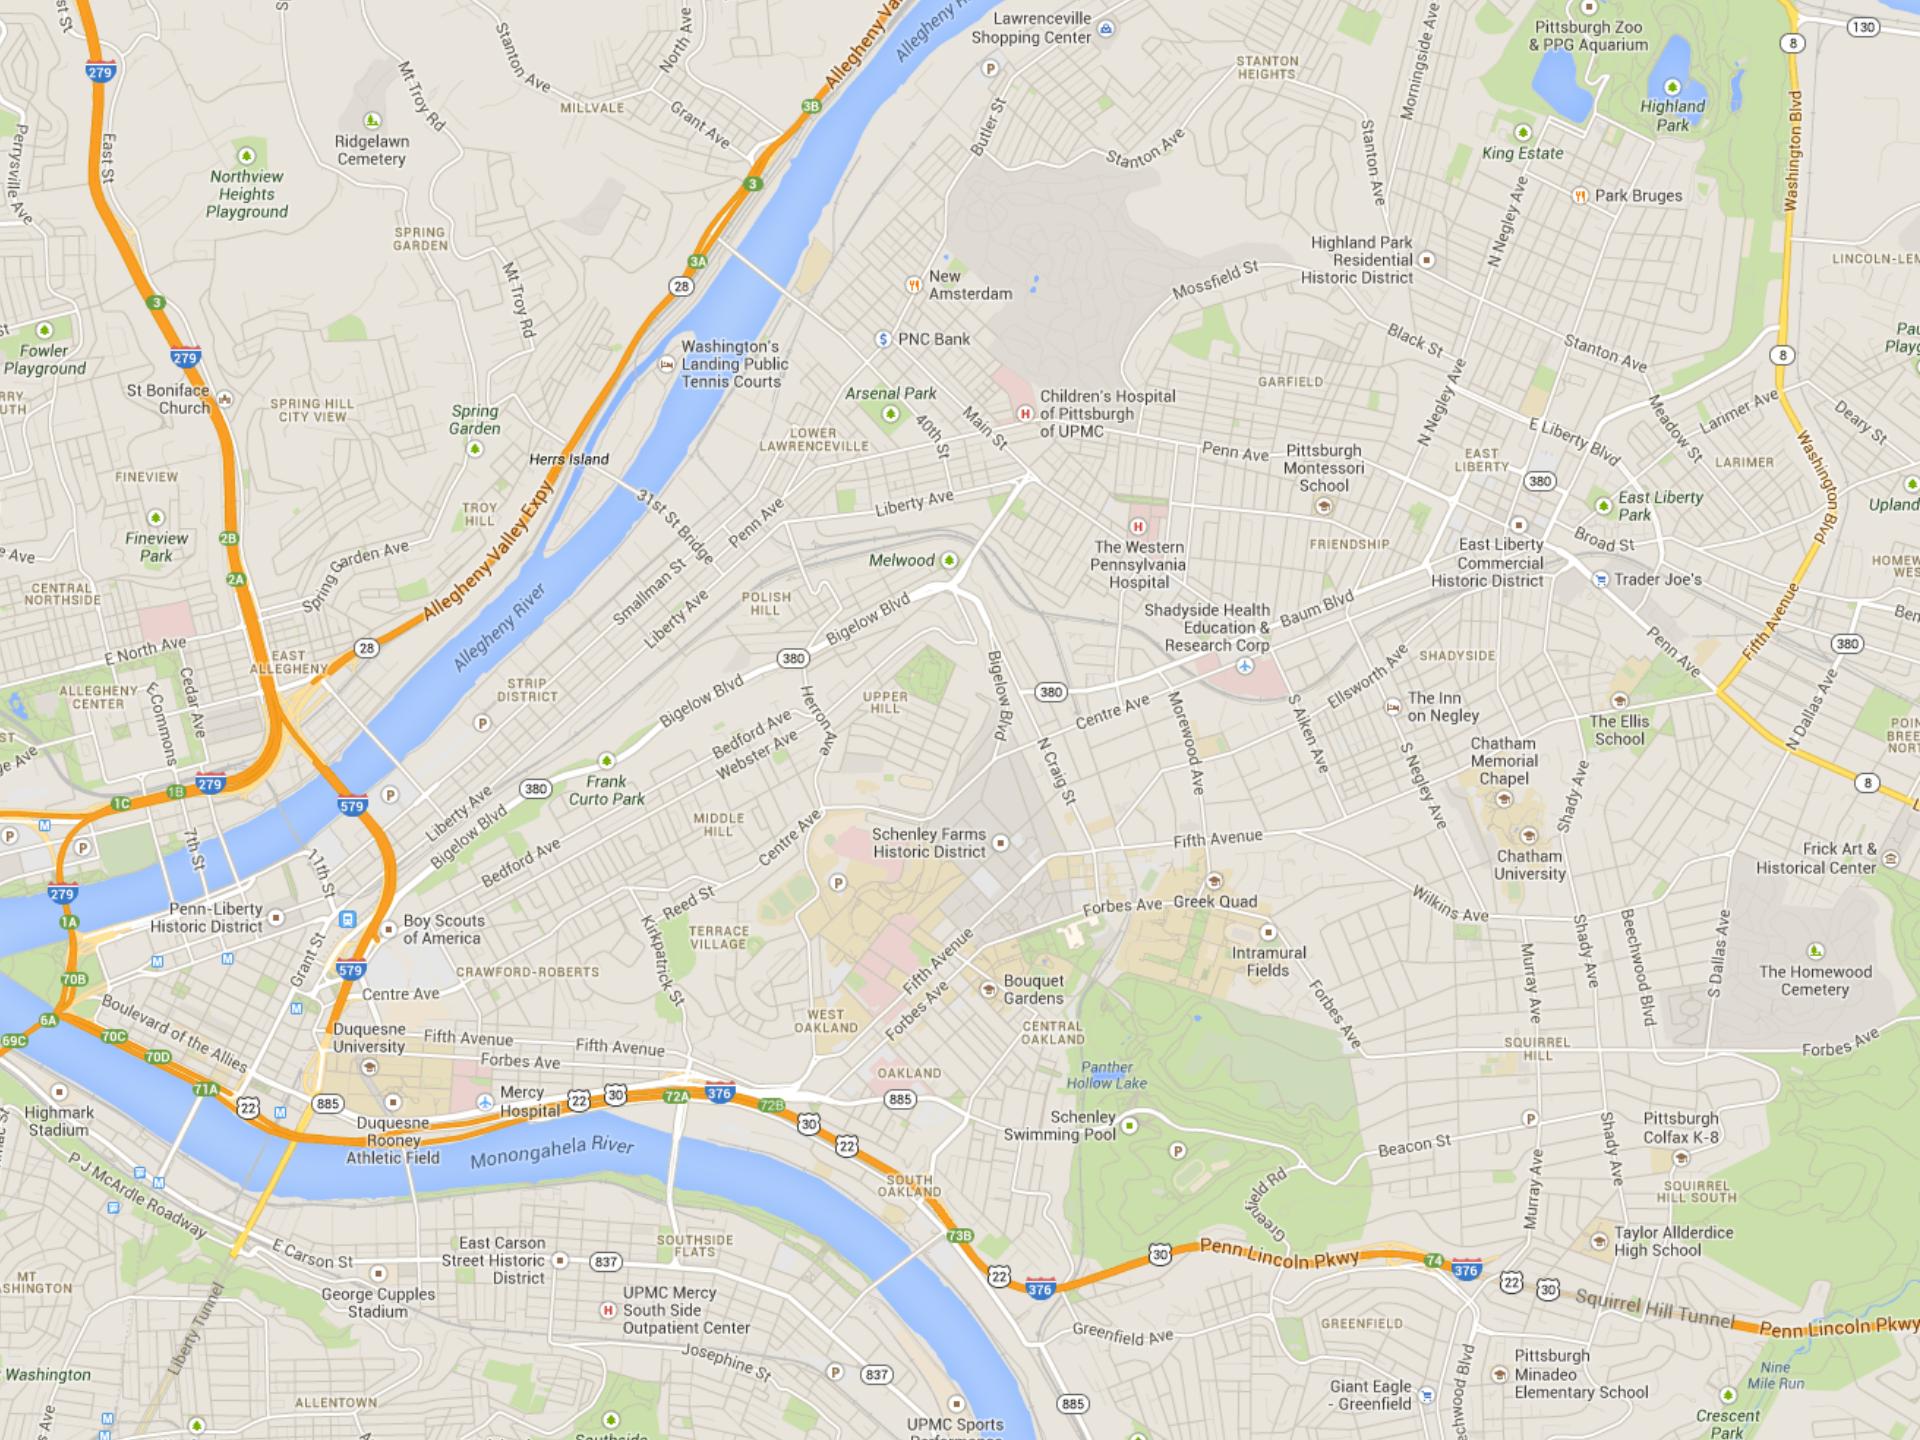


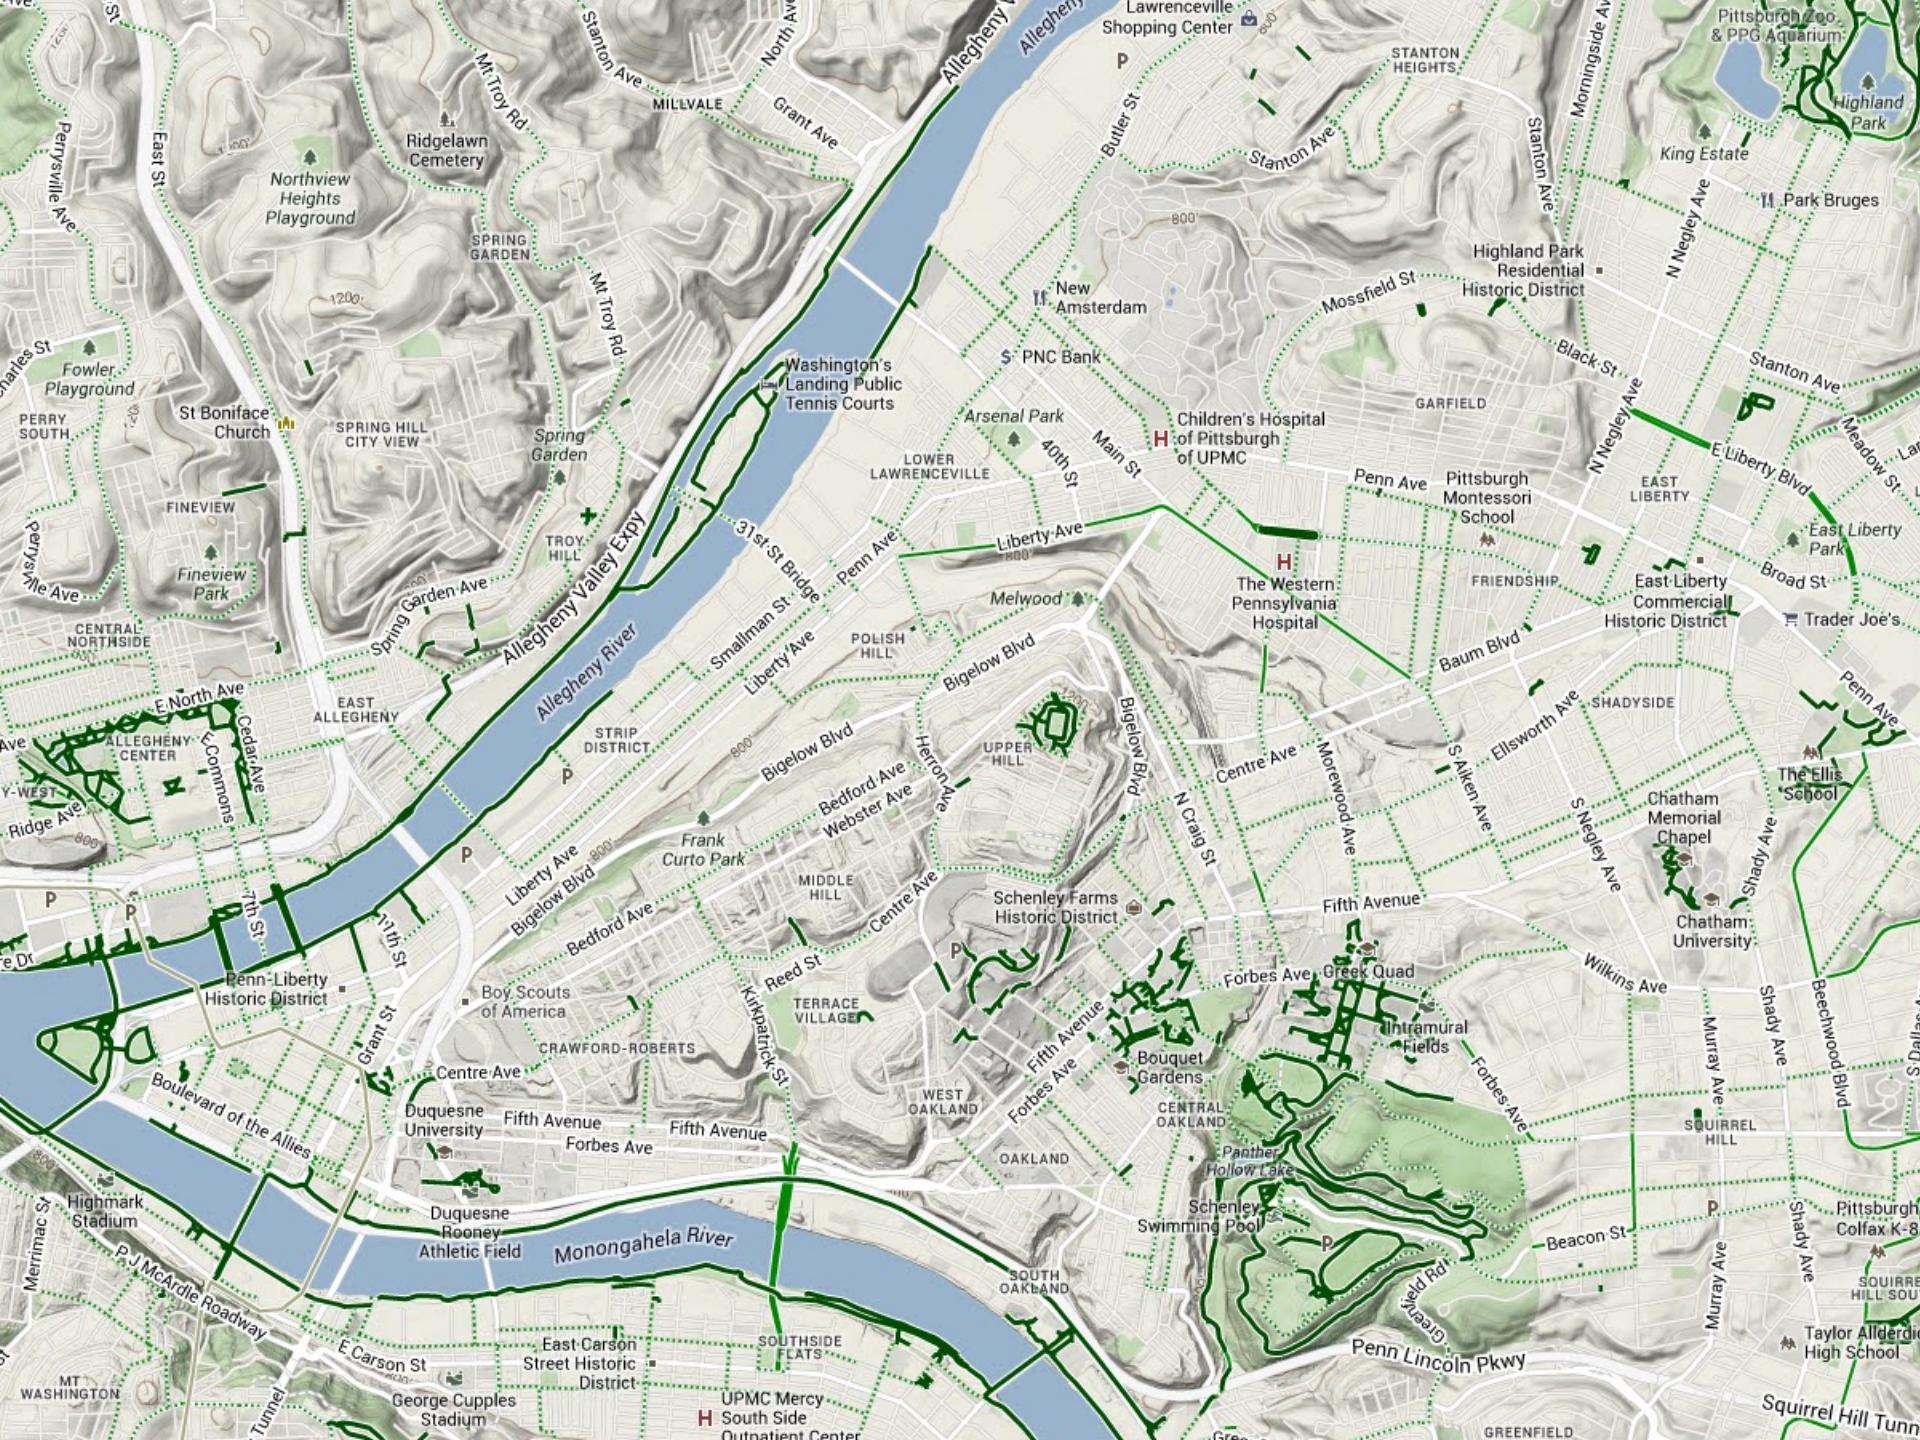
Architecture as (design) process
(activities around the other two)

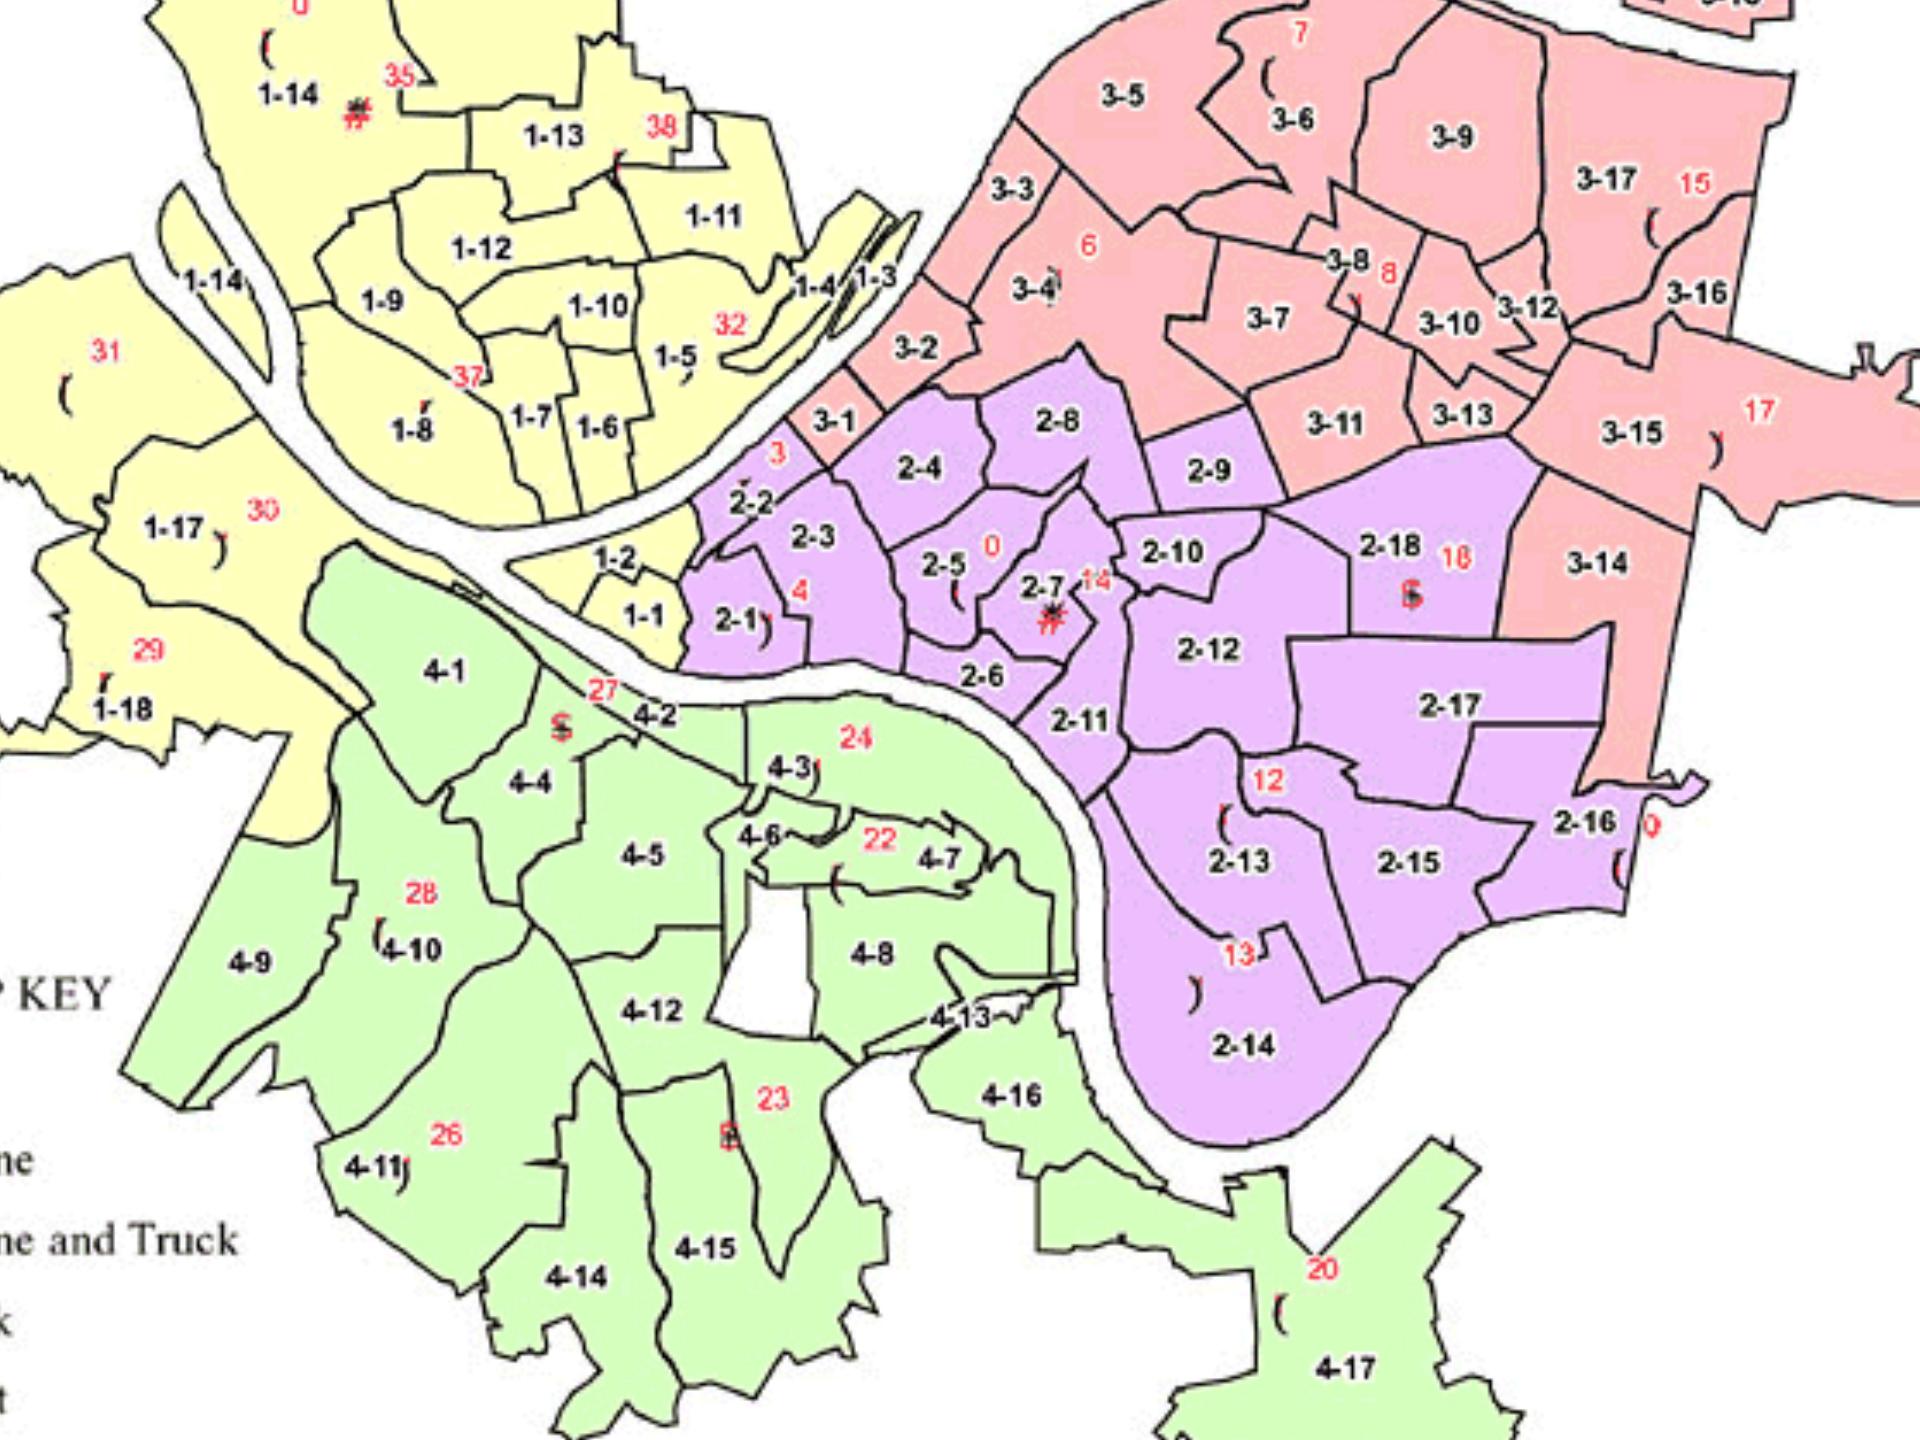


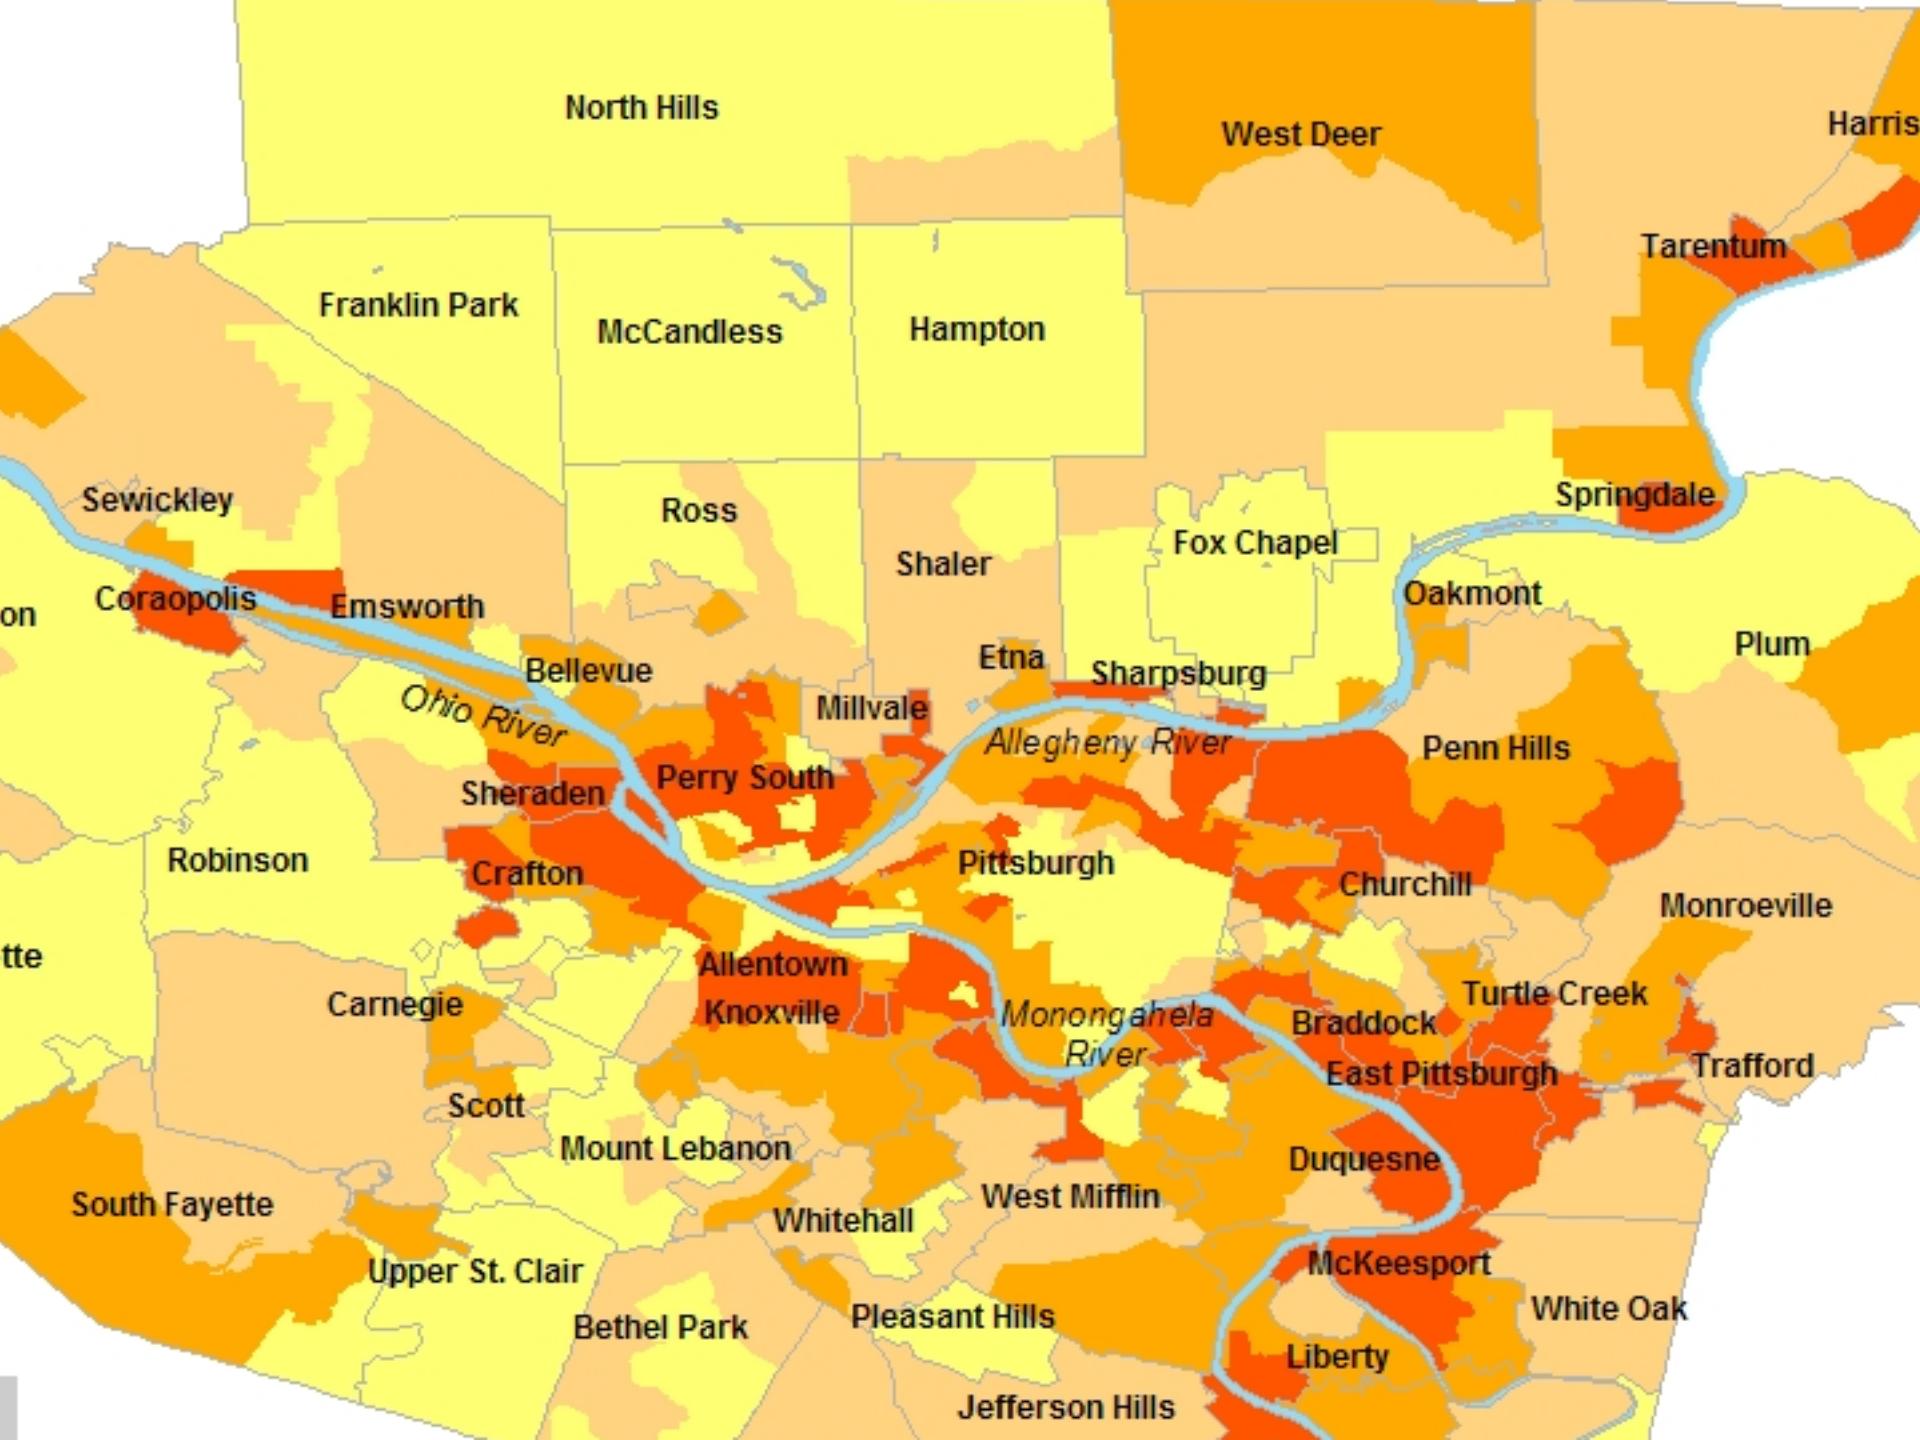
Why Document Architecture?

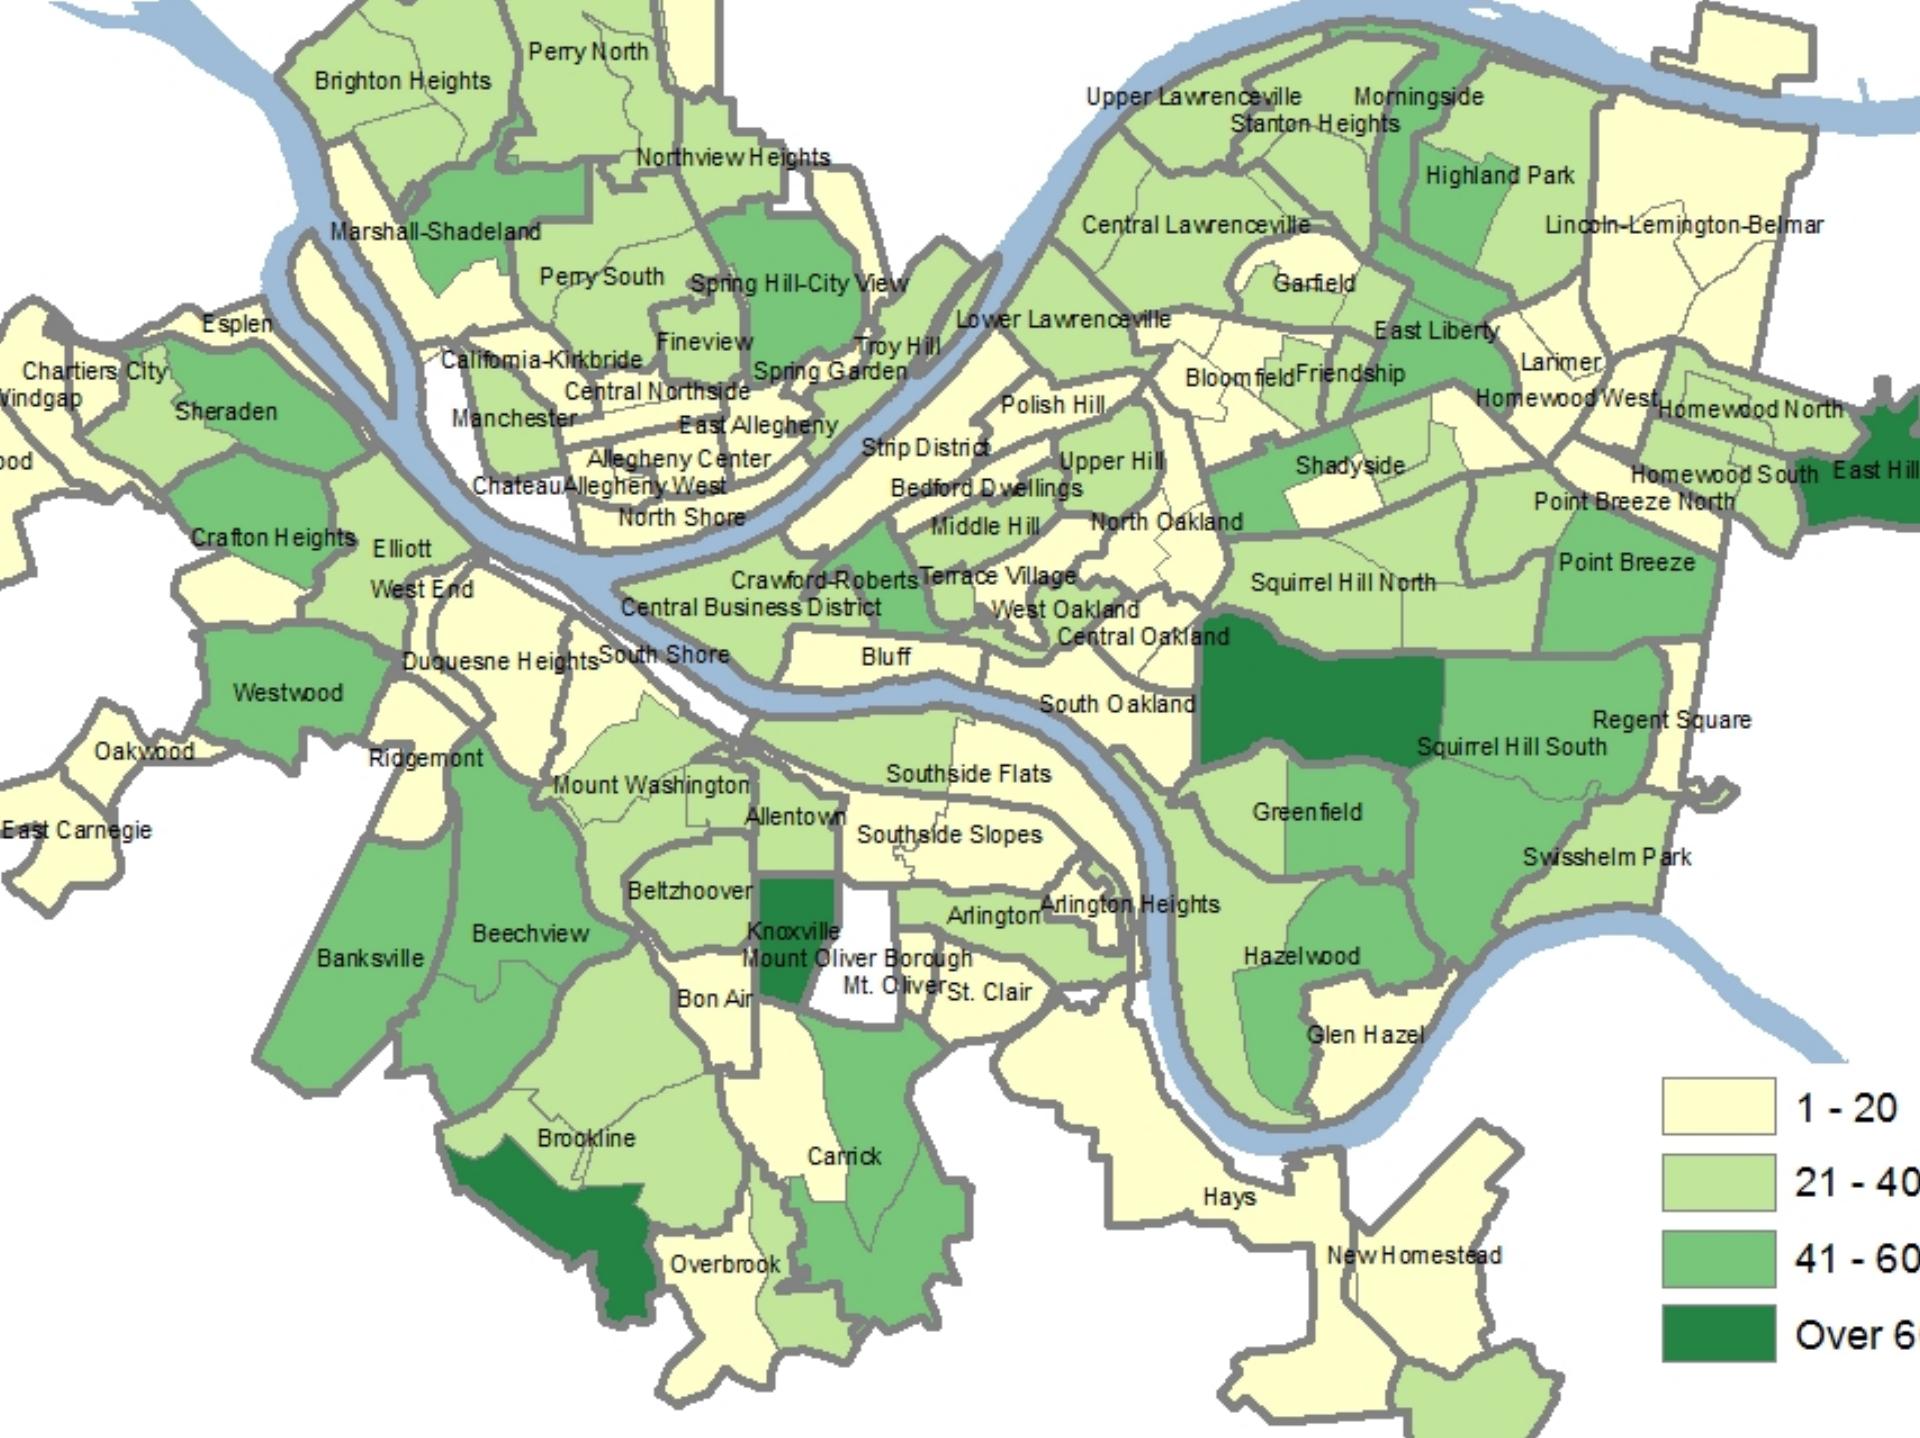
- Blueprint for the system
 - Artifact for early analysis
 - Primary carrier of quality attributes
 - Key to post-deployment maintenance and enhancement
- Documentation speaks for the architect, today and 20 years from today
 - As long as the system is built, maintained, and evolved according to its documented architecture
- Support traceability
- **Reason about quality *tradeoffs*.**













Describing a system's architecture typically requires multiple views or viewpoints.

- Multiple views reflect different perspectives on the design problem.
- Every view should align with a purpose
- Different views are suitable for different reasoning aspects (different quality goals), e.g.,
 - Performance
 - Extensibility
 - Security
 - Scalability
 - ...

Two views of a client-server system

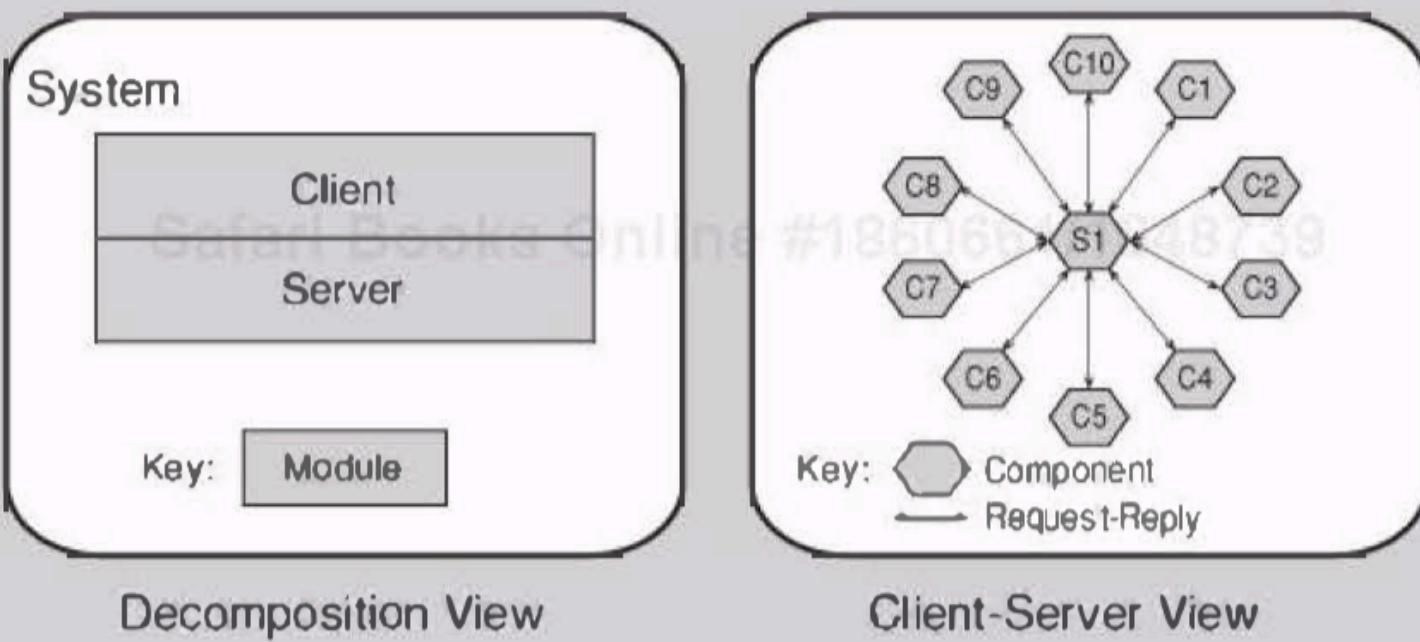
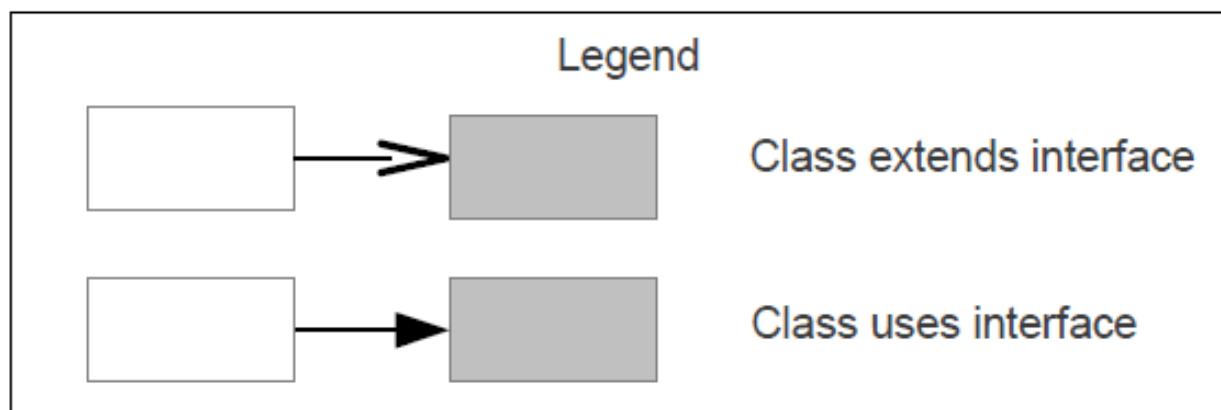
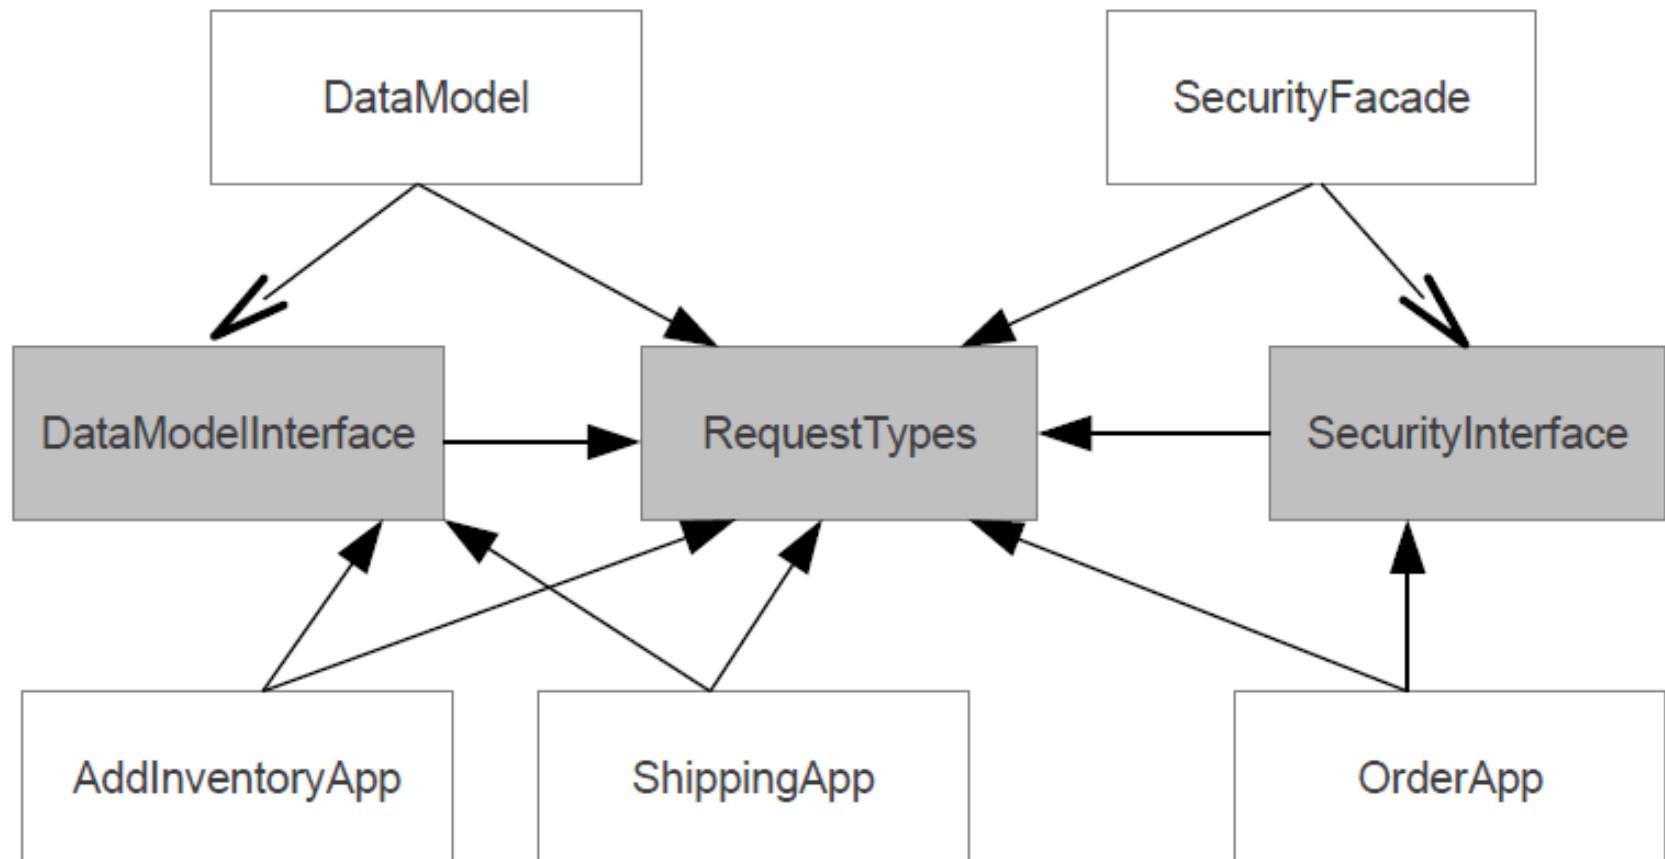
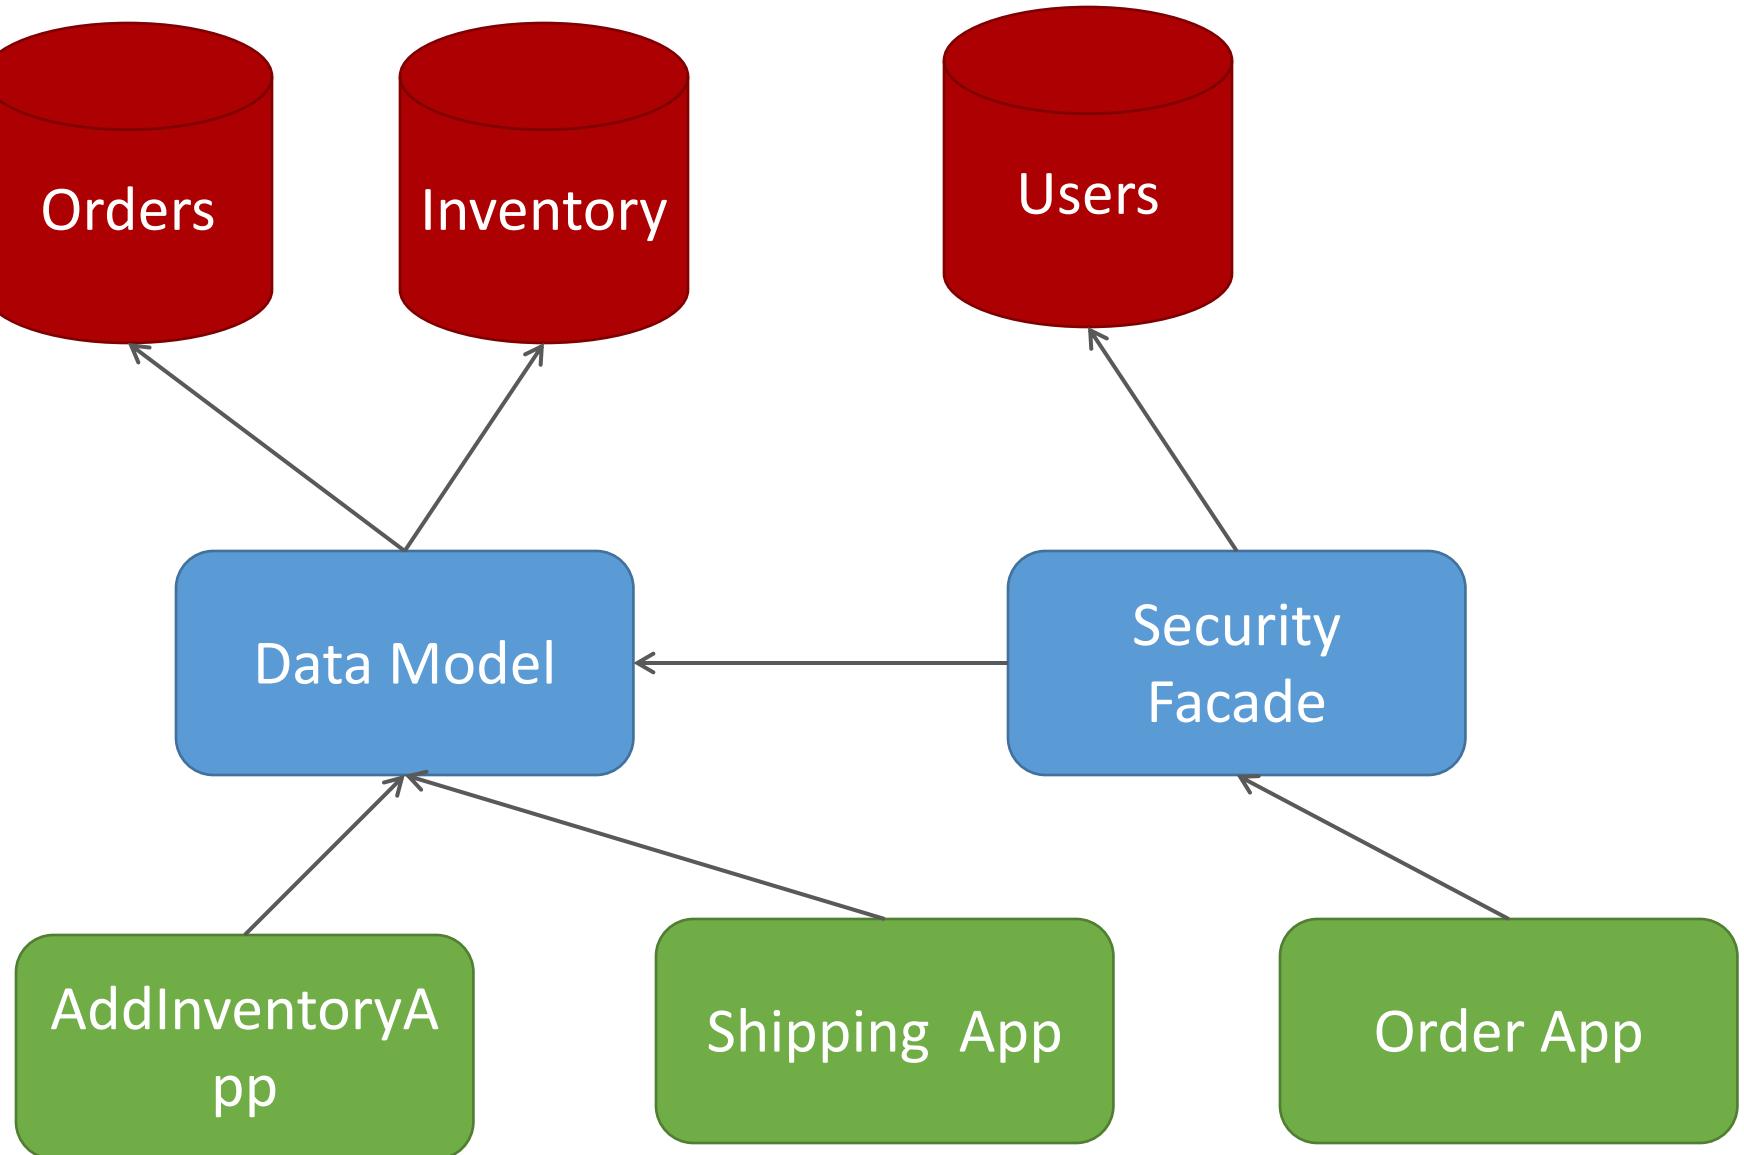


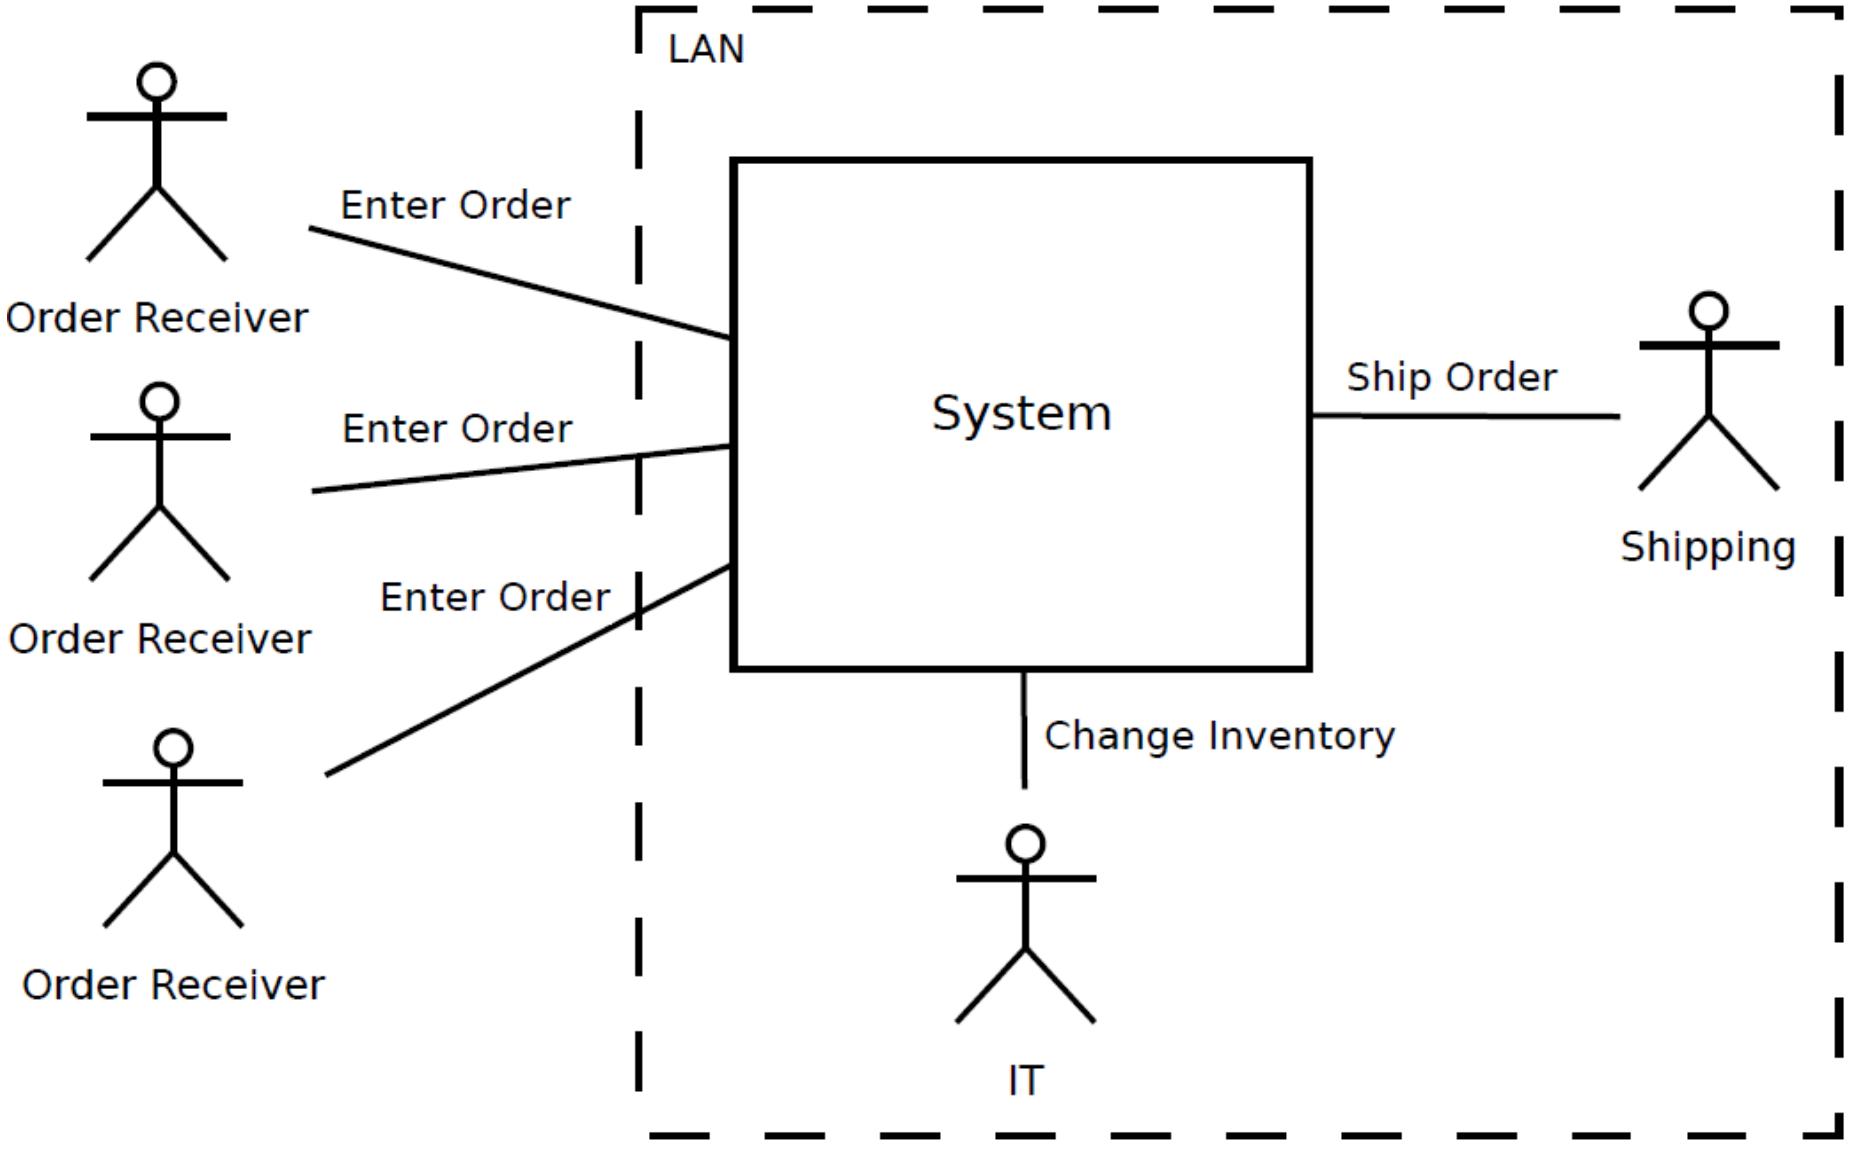
Figure 1.2. Two views of a client-server system

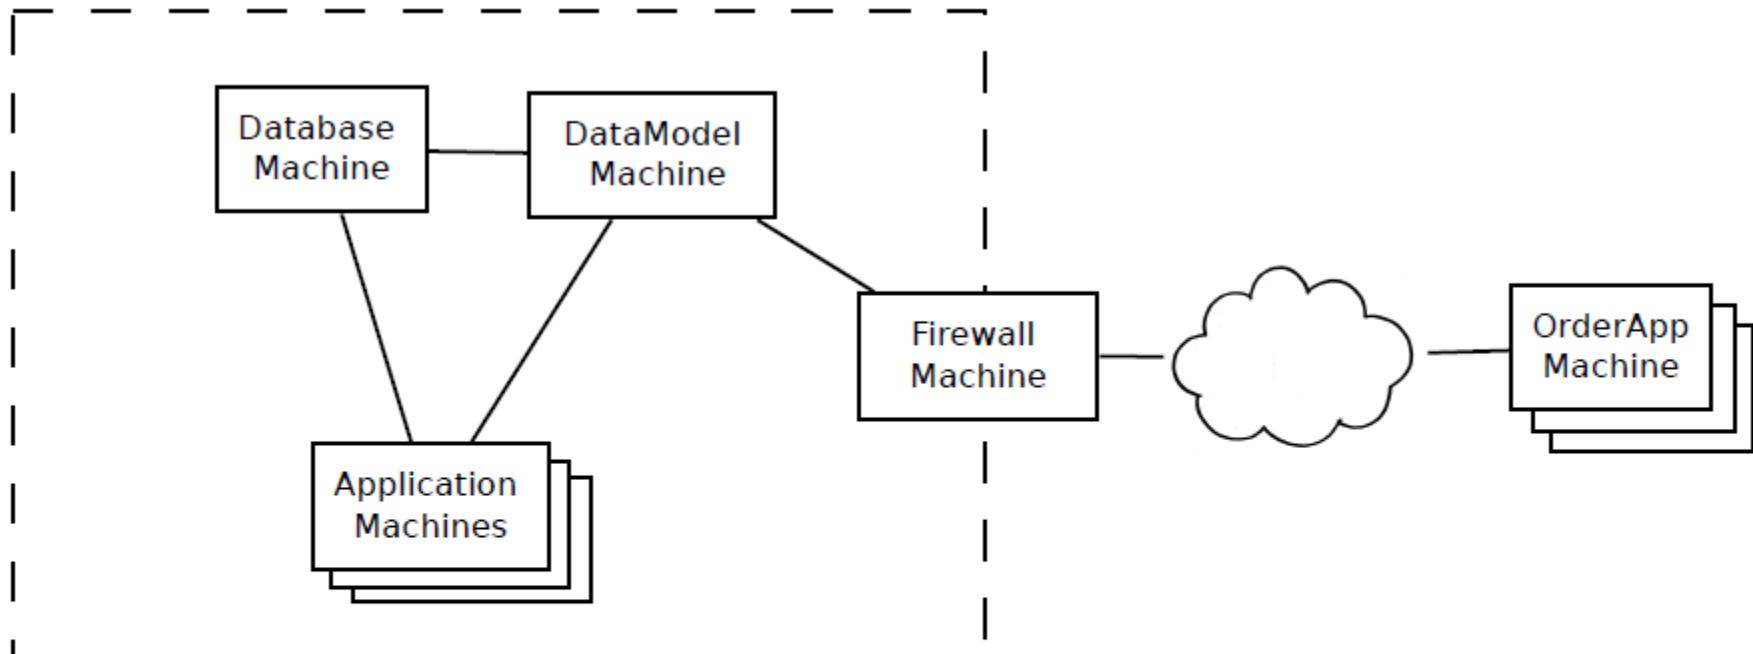
Common Views in Documenting Software Architecture

- Static View
 - Modules (subsystems, structures) and their relations (dependencies, ...)
- Dynamic View
 - Components (processes, runnable entities) and connectors (messages, data flow, ...)
- Physical View (Deployment)
 - Hardware structures and their connections

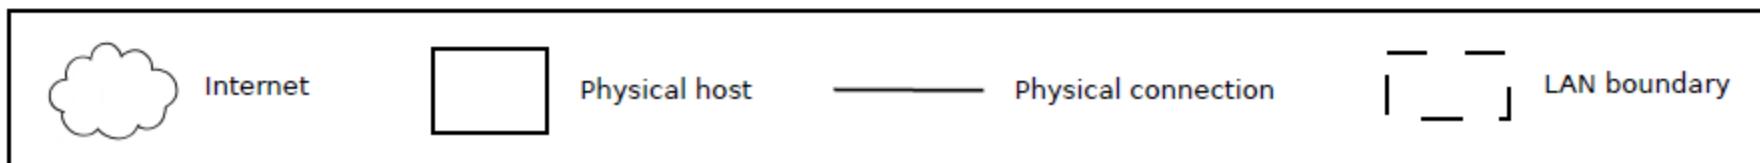








Legend



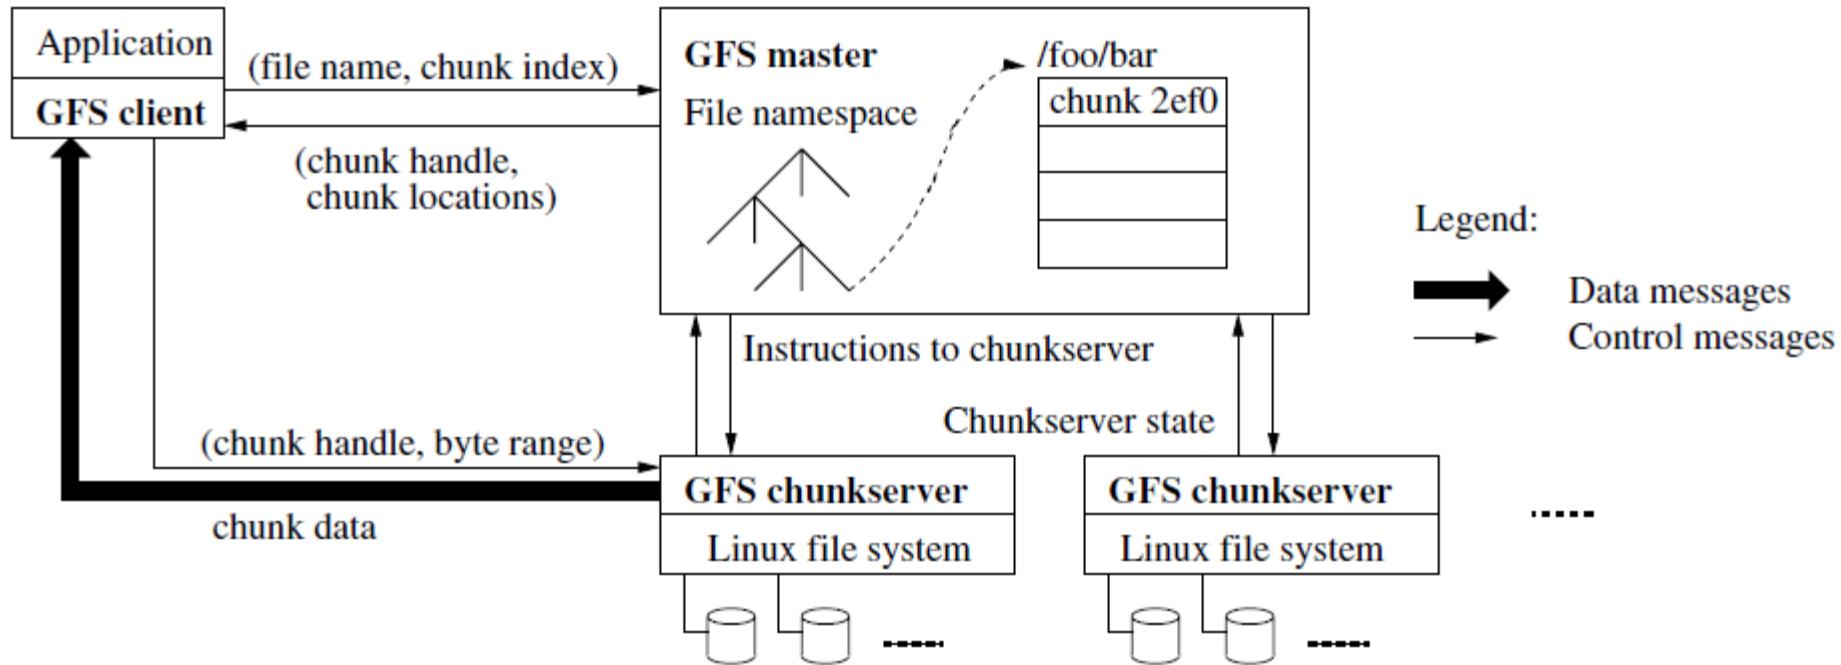


Figure 1: GFS Architecture

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." *ACM SIGOPS operating systems review*. Vol. 37. No. 5. ACM, 2003.

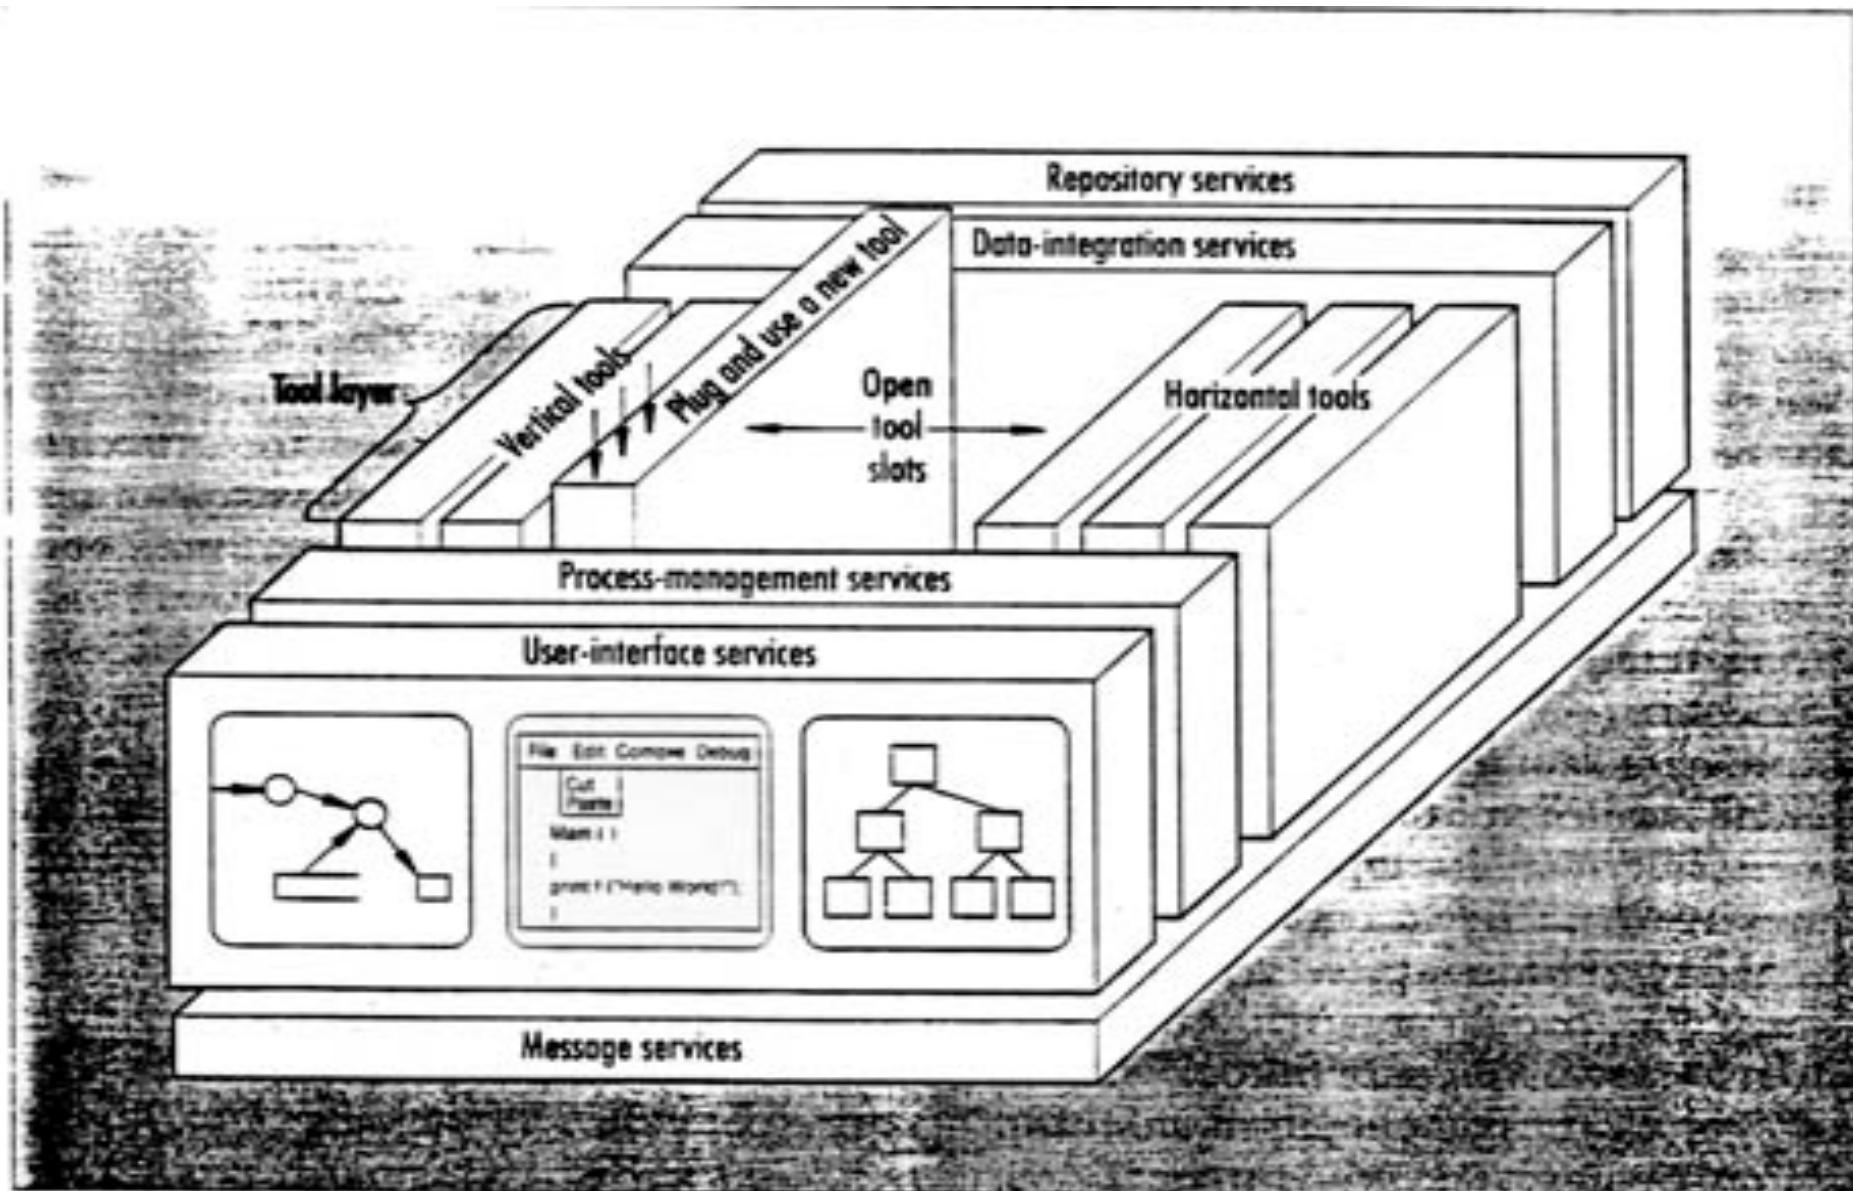


Figure 1. The NIST/ECMA reference model.

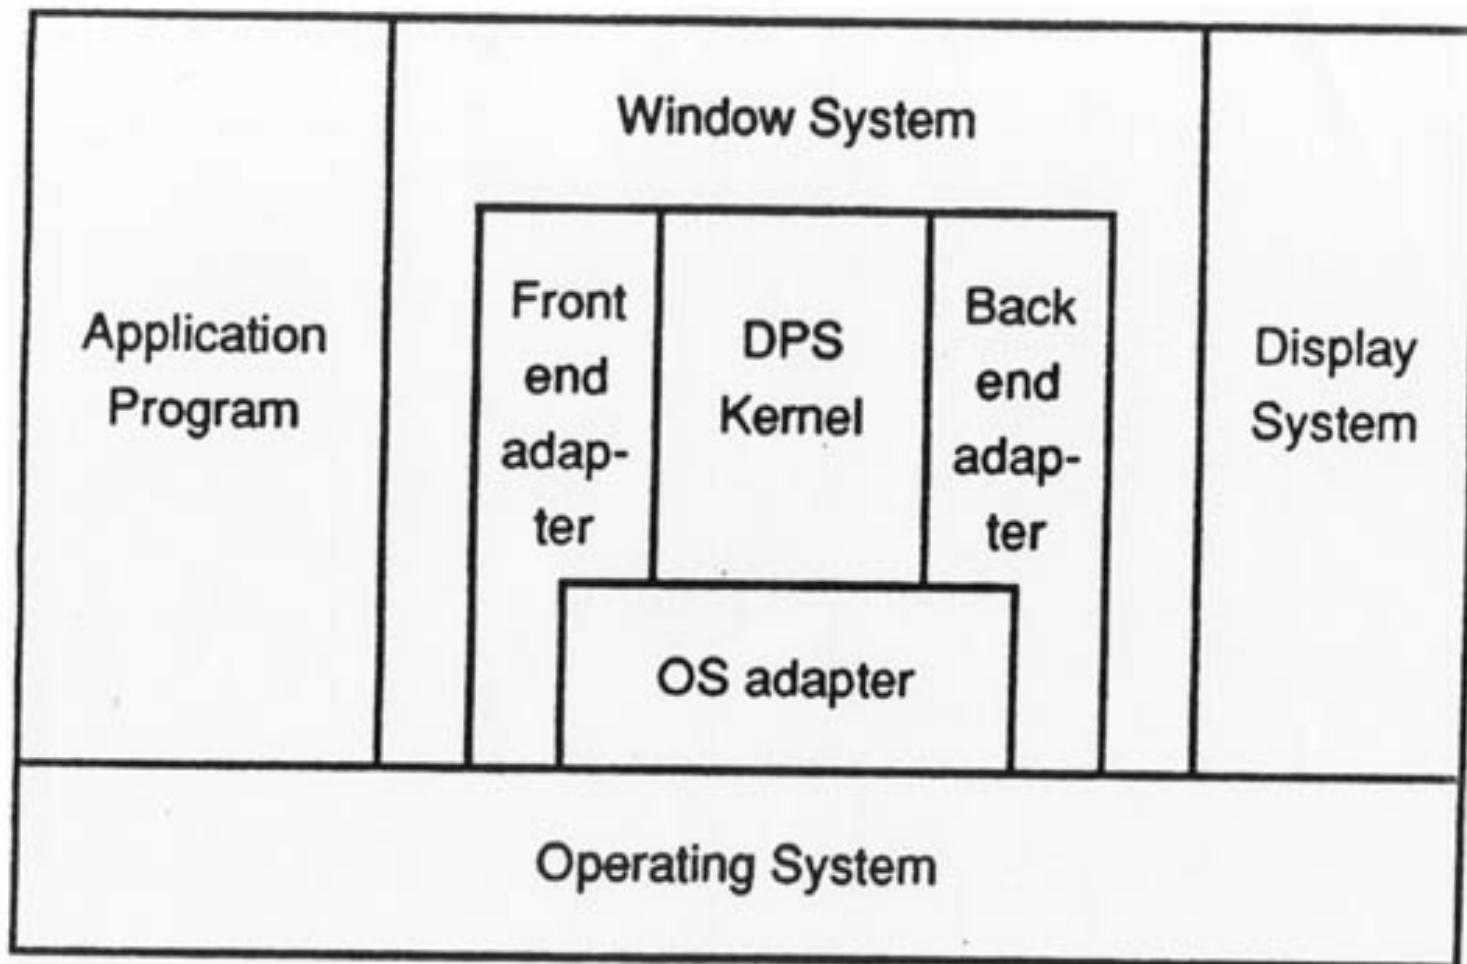


Figure 2. Display PostScript interpreter components.

An Overview of the DISPLAY POSTSCRIPT™ System, Adobe Systems Incorporated, March 16, 1988, P. 10

Client Layer*

Access domain management
Buffering and record-level I/O
Transaction coordination

Agent Layer

Implementation of standard server interface
Logger, agent, and instance tasks

Helix Directories

Path name to FID mapping
Single-file (database) update by one task
Procedural interface for queries

Object (FID directory)

Identification and capability access (via FIDs)
FID to tree-root mapping; table of (FID,root,ref_count)
Existence and deletion (reference counts)
Concurrency control (file interlocking)

Secure Tree

Basic crash-resistant file structure
Conditional commit
Provision of secure array of blocks

System

Commit and restart authority
Disk space allocation
Commit domains

Cache

Caching and performance optimization
Commit support (flush)
Frame allocation (to domains)
Optional disk shadowing

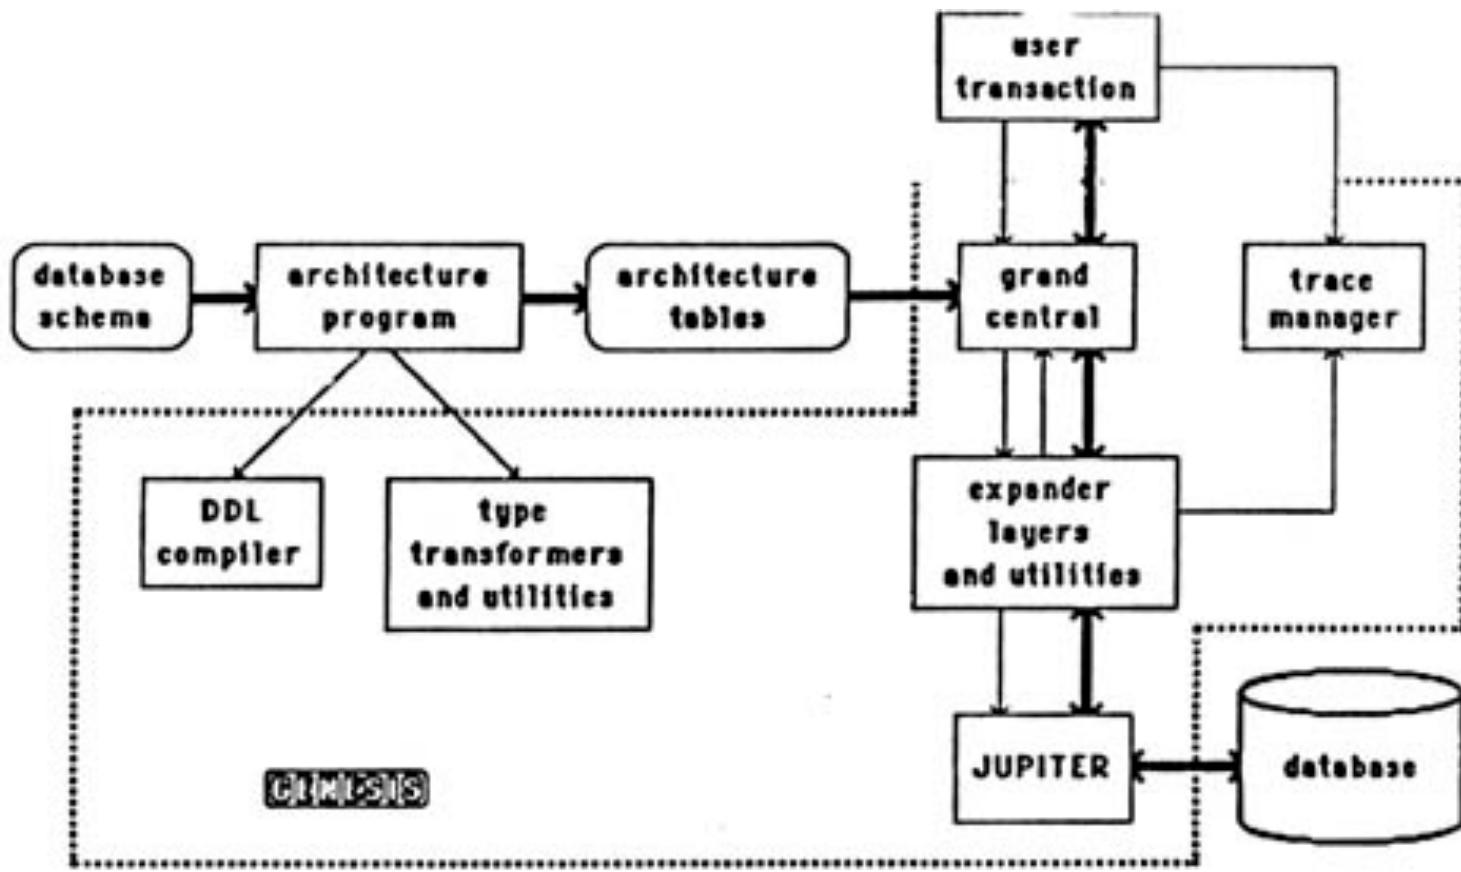
Canonical Disk

Physical disk access

*Also called client Helix.

Figure 2. Abstraction layering.

IEEE Software, "Helix: The architecture of the XMS Distributed File System,
Marek Fridrich and William Olden, May 1985, Vol. 2, No. 3, P. 23



Legend

module or program

$A \rightarrow B$ A calls B

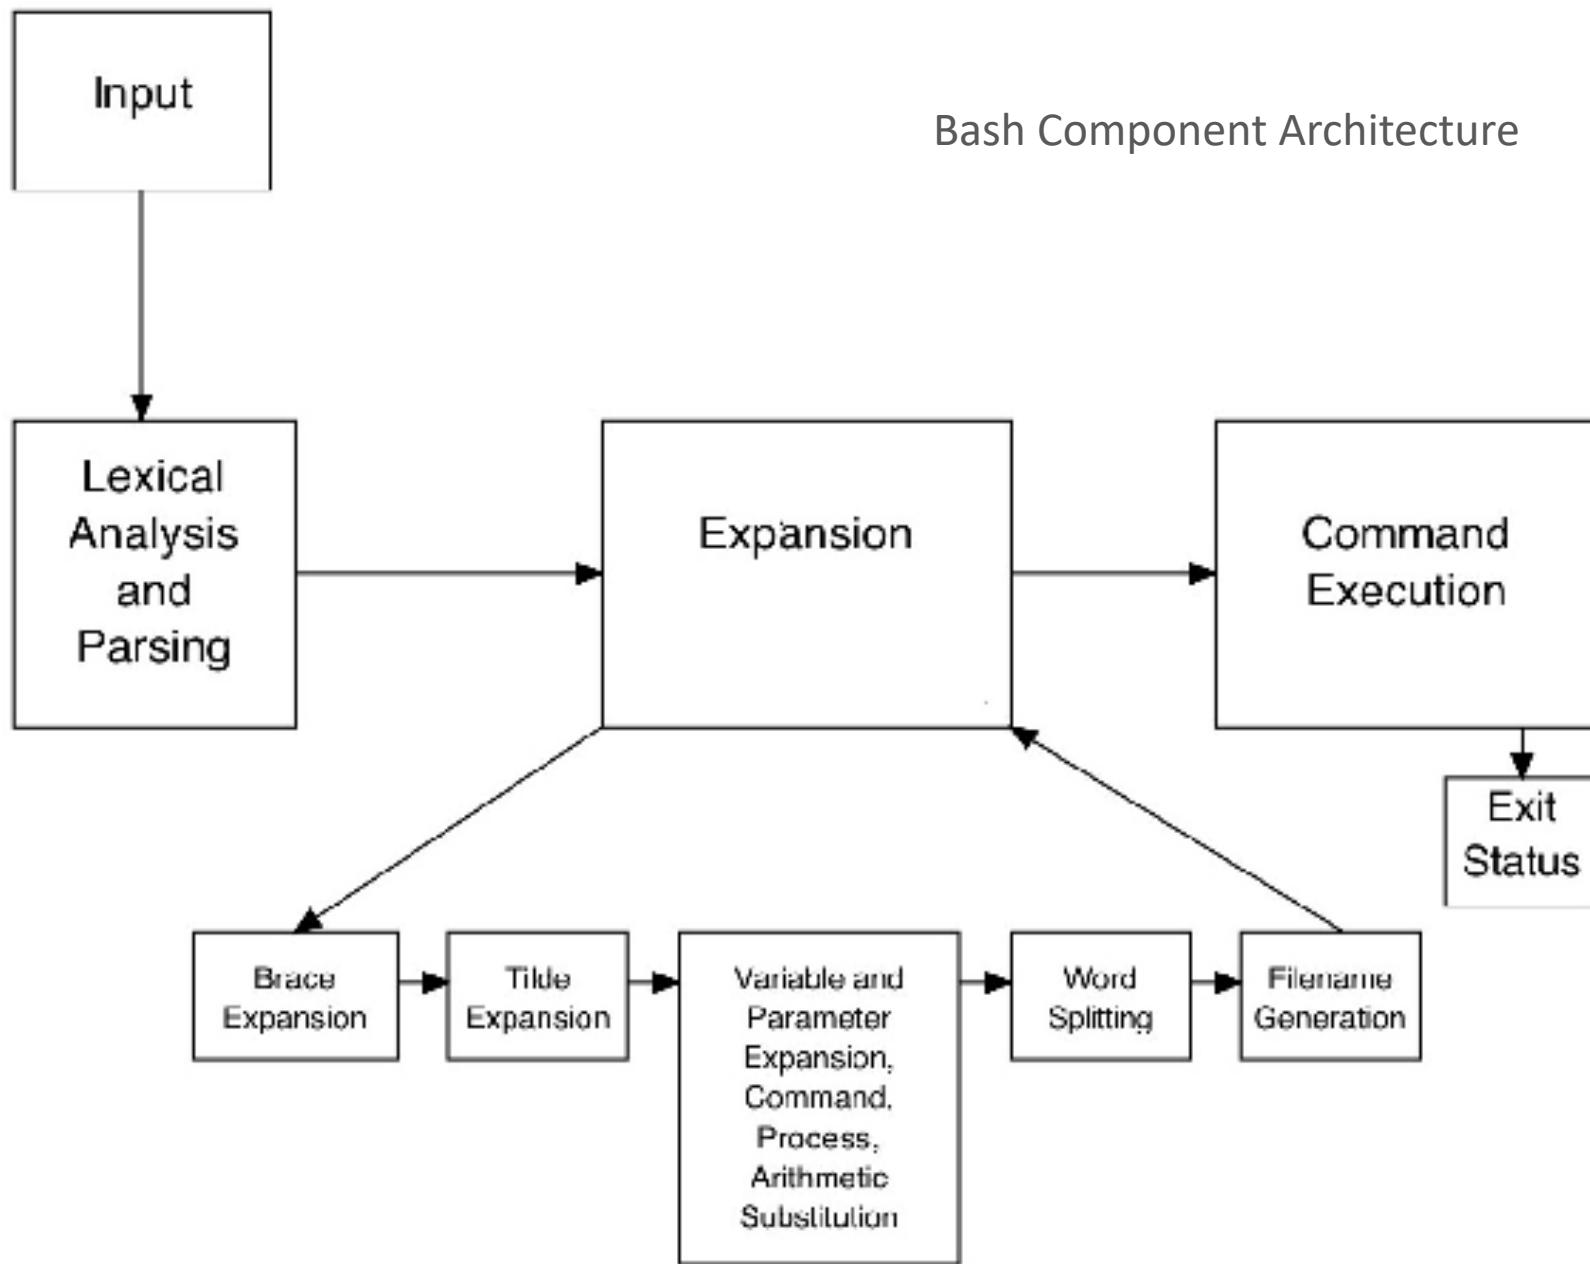
schema or tables

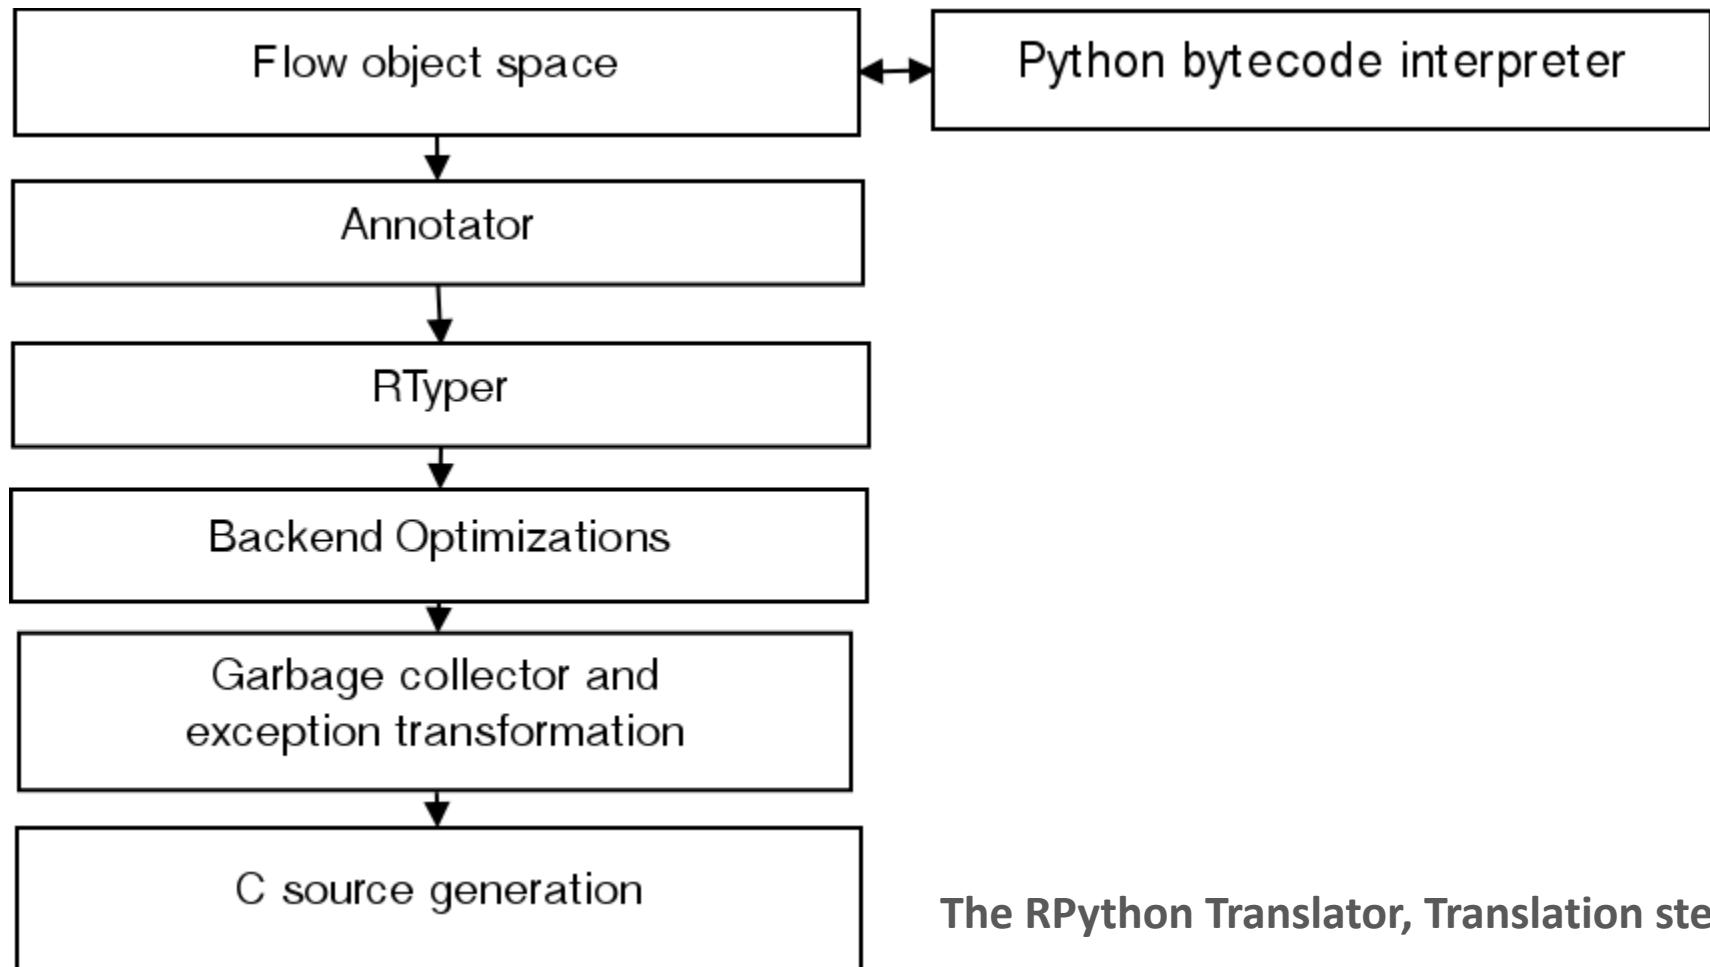
$A \rightarrow B$ data path

Figure 3.1 The Configuration of the GENESIS Prototype

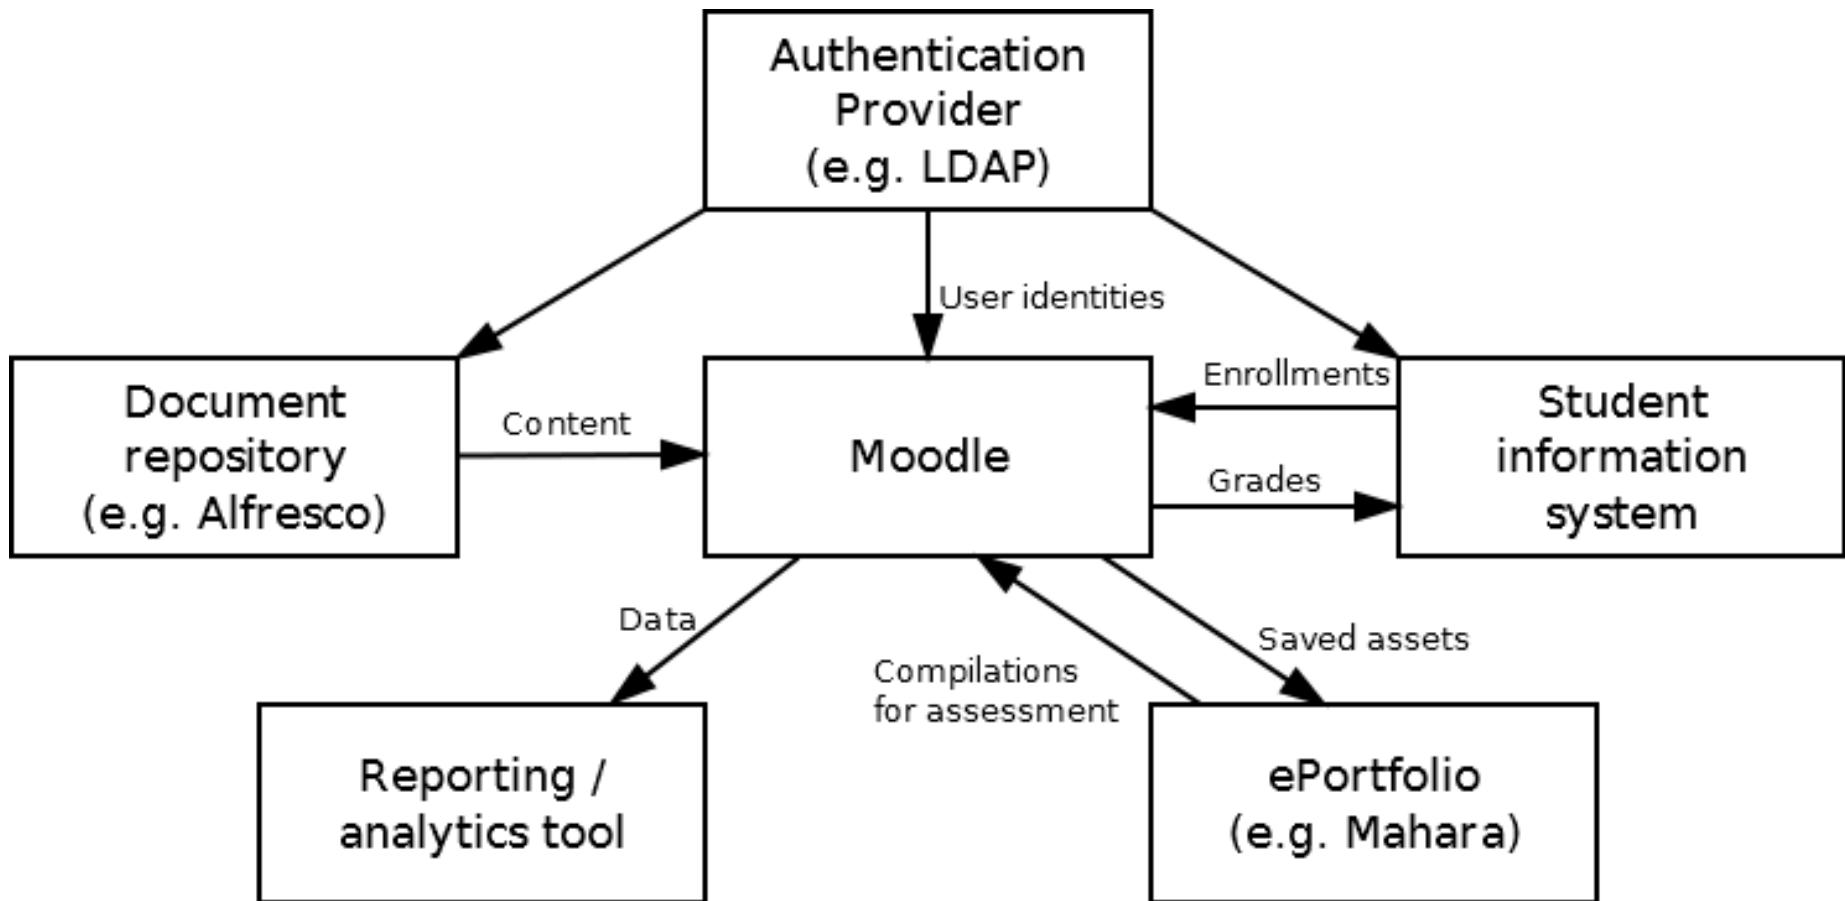
Genesis: A Reconfiguration Database Management System, D. S. Batory, J.R. Barnett, J.F. Garza, K.P. Smith, K. Tsukuda, B.C. Twichell, T.E. Wise, Department of Computer Sciences, University of Texas at Austin,

Bash Component Architecture





The RPython Translator, Translation steps



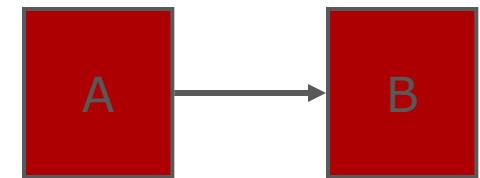
Moodle: Typical university systems architecture – Key subsystems

Bad architectural documentation.

- In practice today's documentation consists of
 - Ambiguous box-and-line diagrams
 - Inconsistent use of notations
 - Confusing combinations of viewtypes
- Many things are left unspecified:
 - What kind of elements?
 - What kind of relations?
 - What do the boxes and arrows mean?
 - What is the significance of the layout?

What could the arrow mean?

- Many possibilities
 - A passes control to B
 - A passes data to B
 - A gets a value from B
 - A streams data to B
 - A sends a message to B
 - A creates B
 - A occurs before B
 - B gets its electricity from A
 - ...



Guidelines: Avoiding Ambiguity

- Always include a legend
- Define precisely what the boxes mean
- Define precisely what the lines mean
- Don't mix viewtypes unintentionally
 - Recall: Module (classes), C&C (components)
- Supplement graphics with explanation
 - Very important: rationale (architectural intent)
- Do not try to do too much in one diagram
 - Each view of architecture should fit on a page
 - Use hierarchy

CASE STUDY: THE GOOGLE FILE SYSTEM

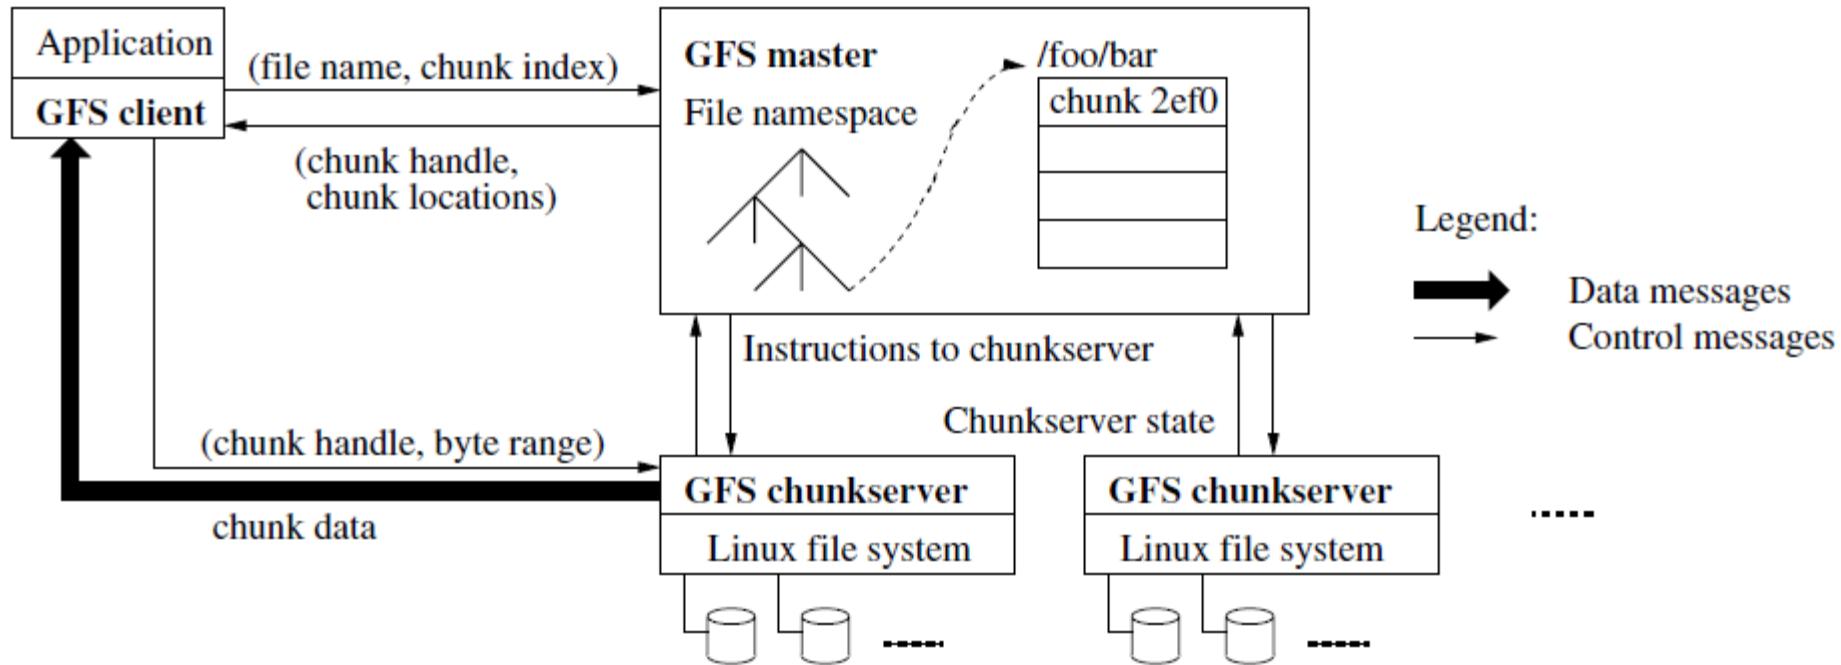


Figure 1: GFS Architecture

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." *ACM SIGOPS operating systems review*. Vol. 37. No. 5. ACM, 2003.

Assumptions

- The system is built from many inexpensive commodity components that often fail.
- The system stores a modest number of large files.
- The workloads primarily consist of two kinds of reads: large streaming reads and small random reads.
- The workloads also have many large, sequential writes that append data to files.
- The system must efficiently implement well-defined semantics for multiple clients that concurrently append to the same file.
- High sustained bandwidth is more important than low latency.

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." *ACM SIGOPS operating systems review*. Vol. 37, No. 5. ACM, 2003.

Qualities:
Scalability
Reliability
Performance
Cost

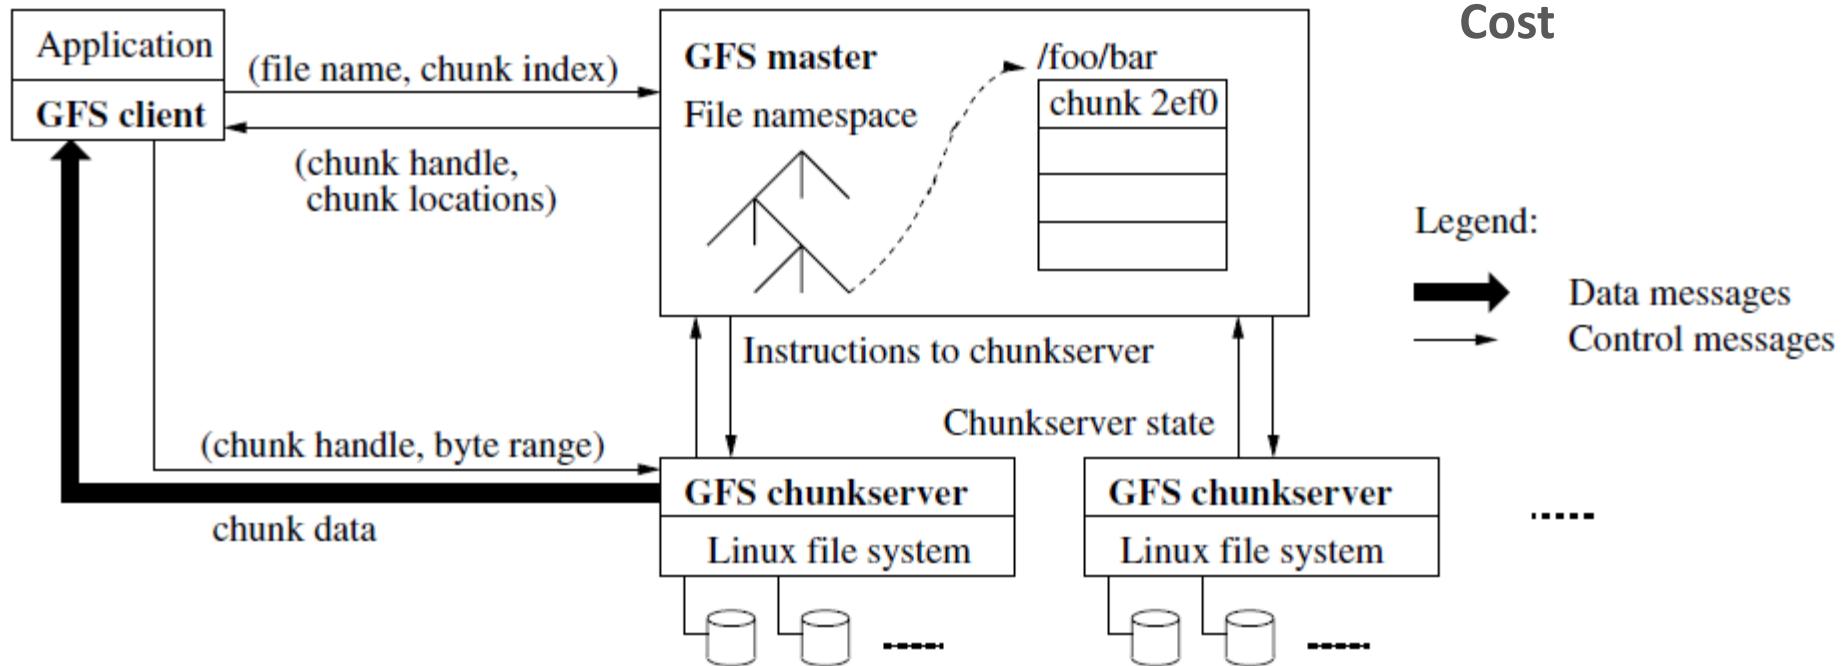


Figure 1: GFS Architecture

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." *ACM SIGOPS operating systems review*. Vol. 37. No. 5. ACM, 2003.

Questions

1. What are the most important quality attributes in the design?
2. How are those quality attributes realized in the design?

Qualities:
Scalability
Reliability
Performance
Cost

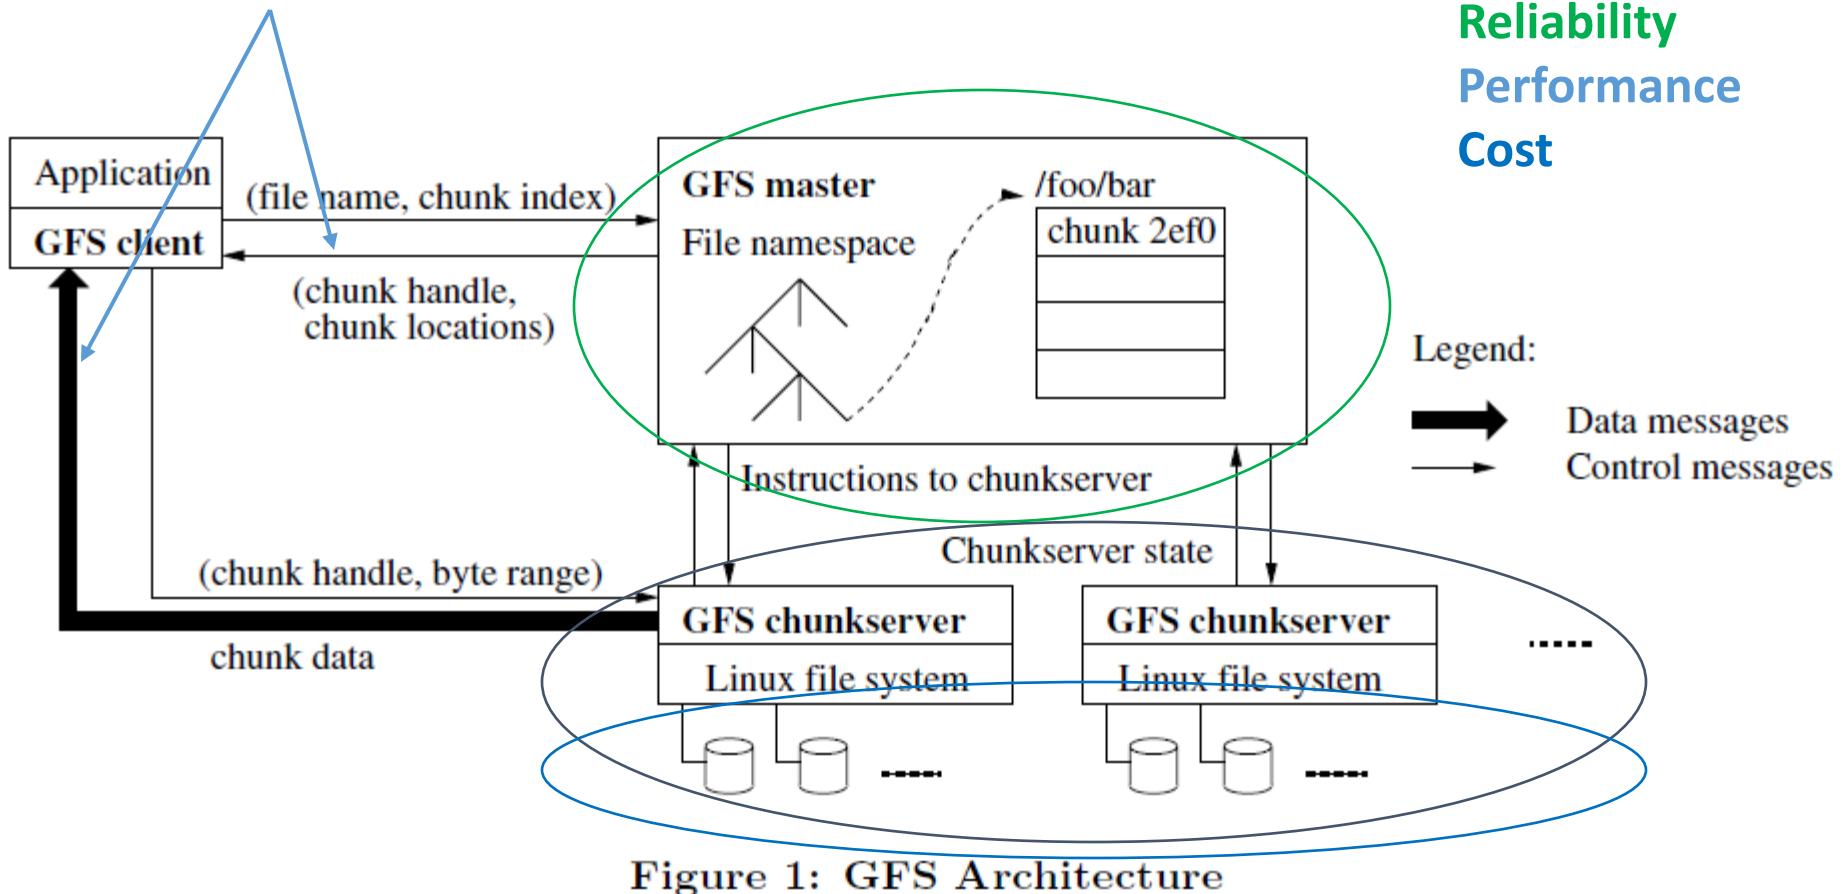


Figure 1: GFS Architecture

Exercise

For the Google File System, create a physical architecture view that addresses a relevant quality attribute

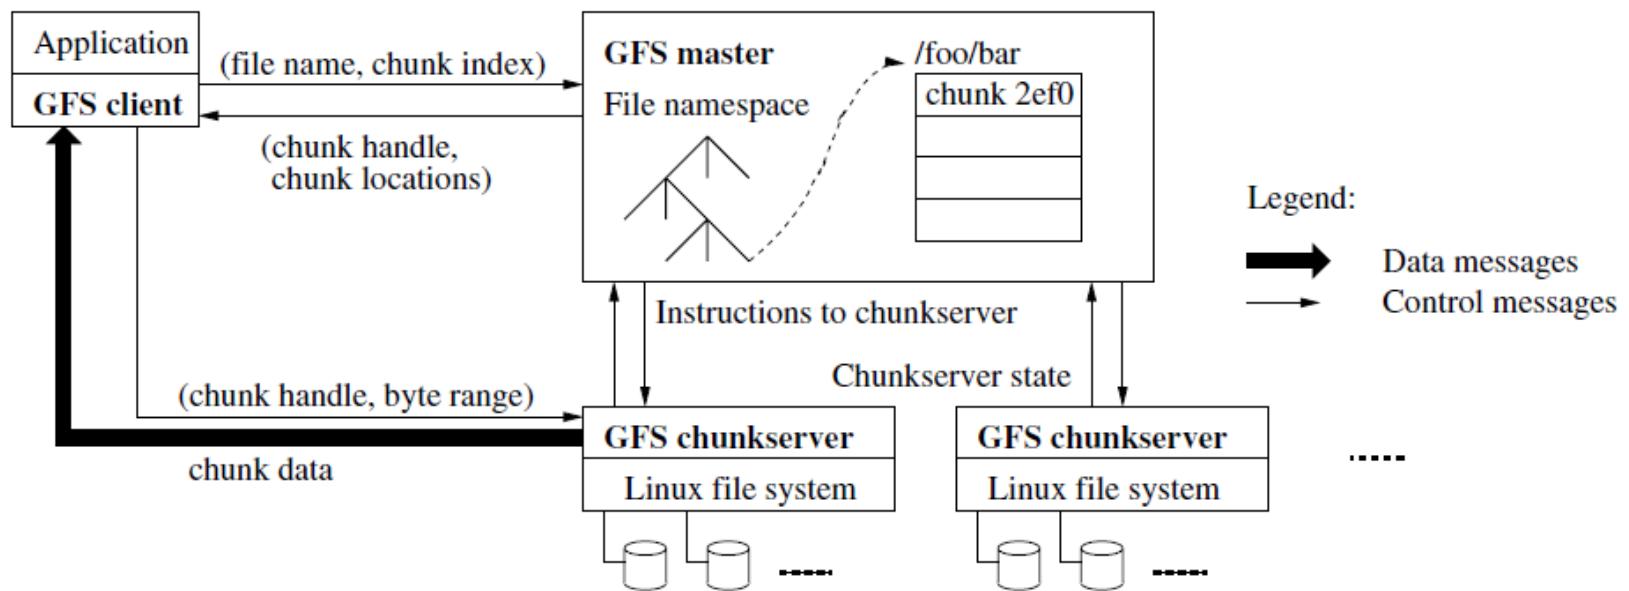
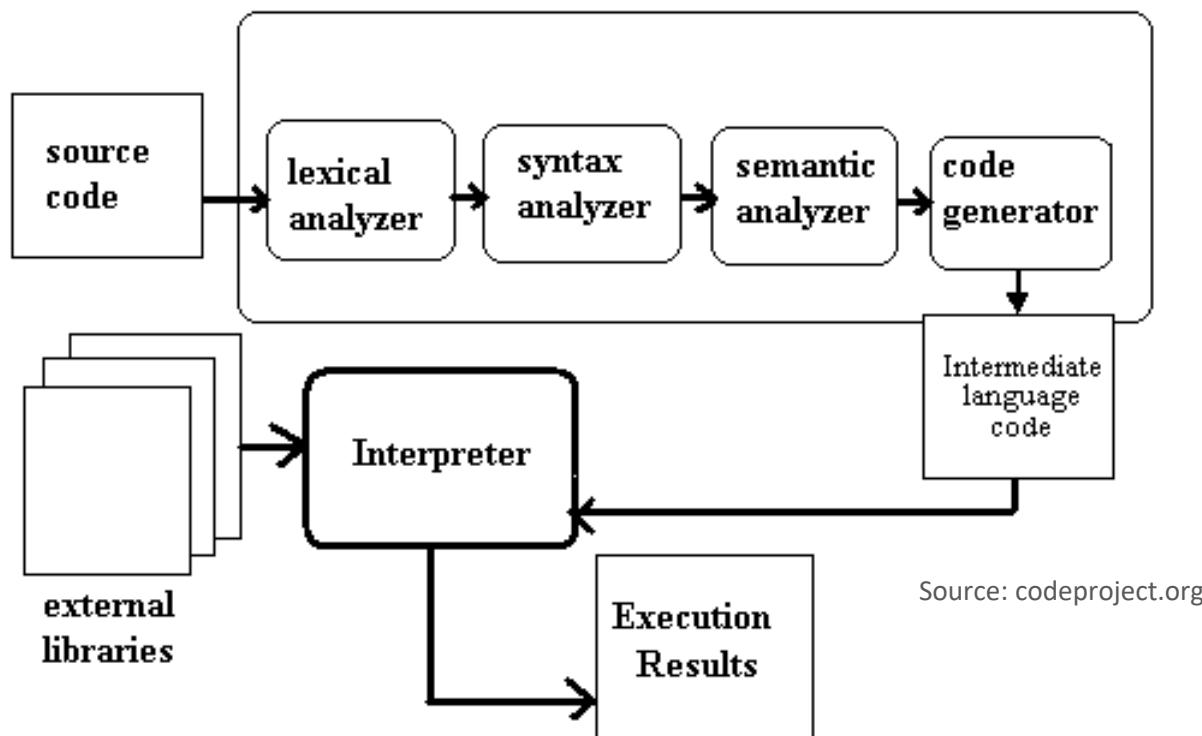


Figure 1: GFS Architecture

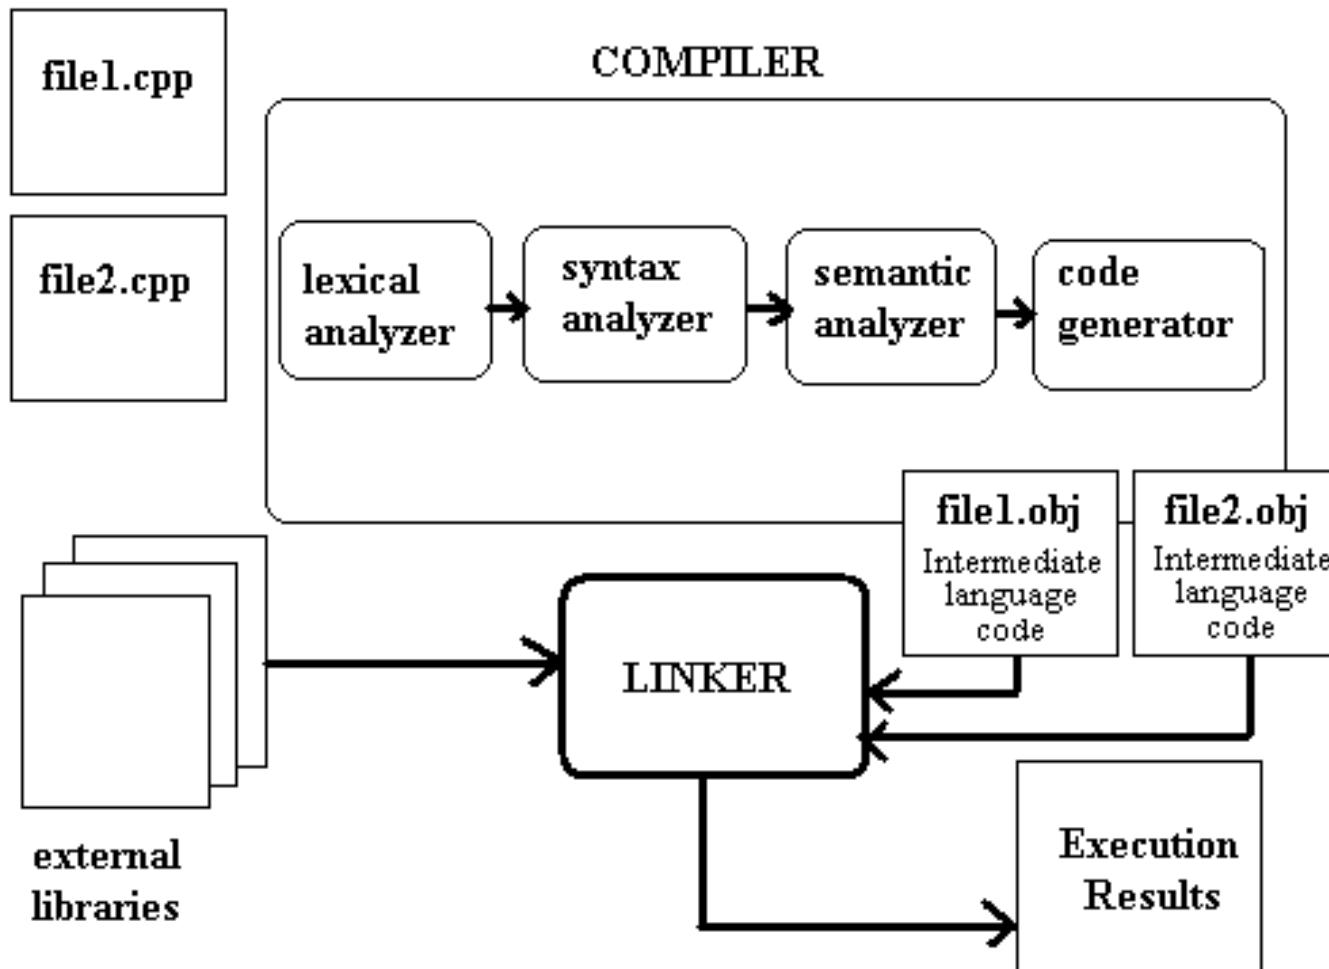
ARCHITECTURAL STYLES AND TACTICS

Architectural style/pattern: broad principle of system organization.

- Describes computational model
 - E.g., pipe and filter, call-return, publish-subscribe, layered, services
- Related to one of common view types
 - Static, dynamic, physical



Architectural style (pattern)



Source: codeproject.org

Example Architectural Patterns

- System organization
 - Repository model
 - Client-server model
 - Layered model
- Modular decomposition
 - Object oriented
 - Function-oriented pipelining
- Control styles
 - Centralized control
 - Event-driven systems

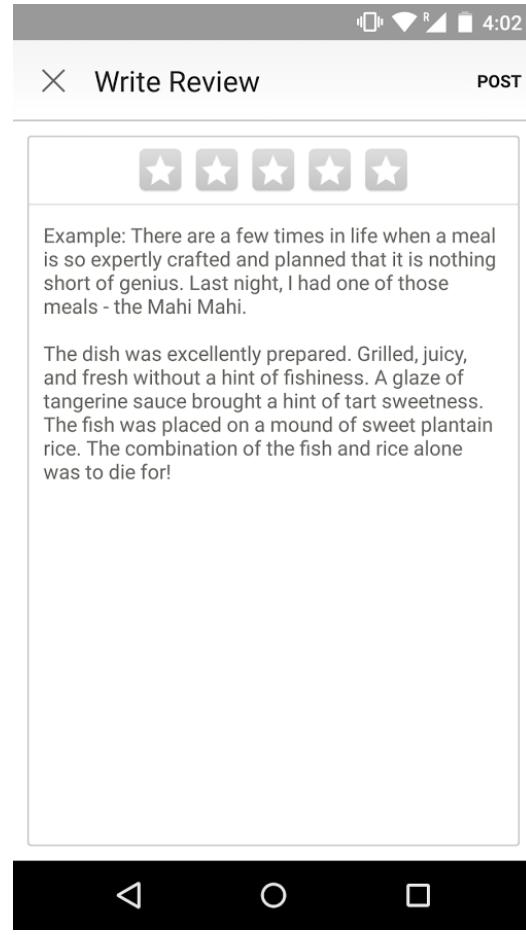
Client

Server

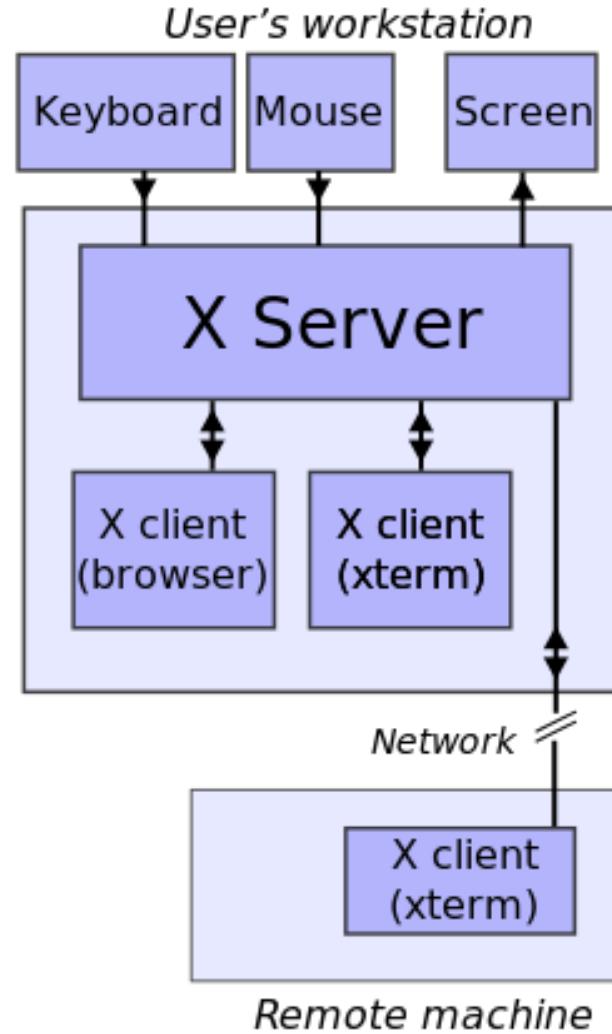
Database

Where to
validate user
input?

Example: Yelp App



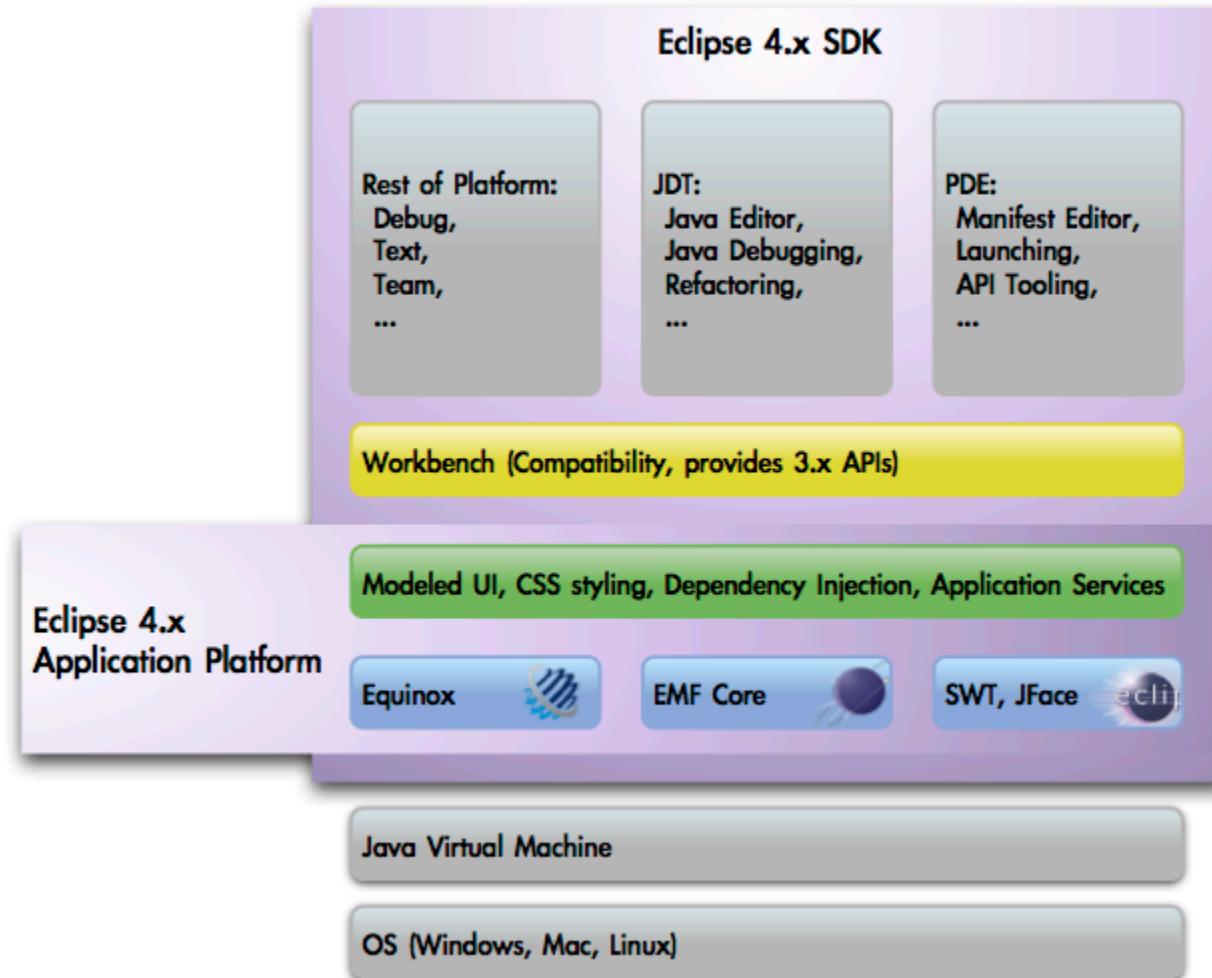
Client-server style



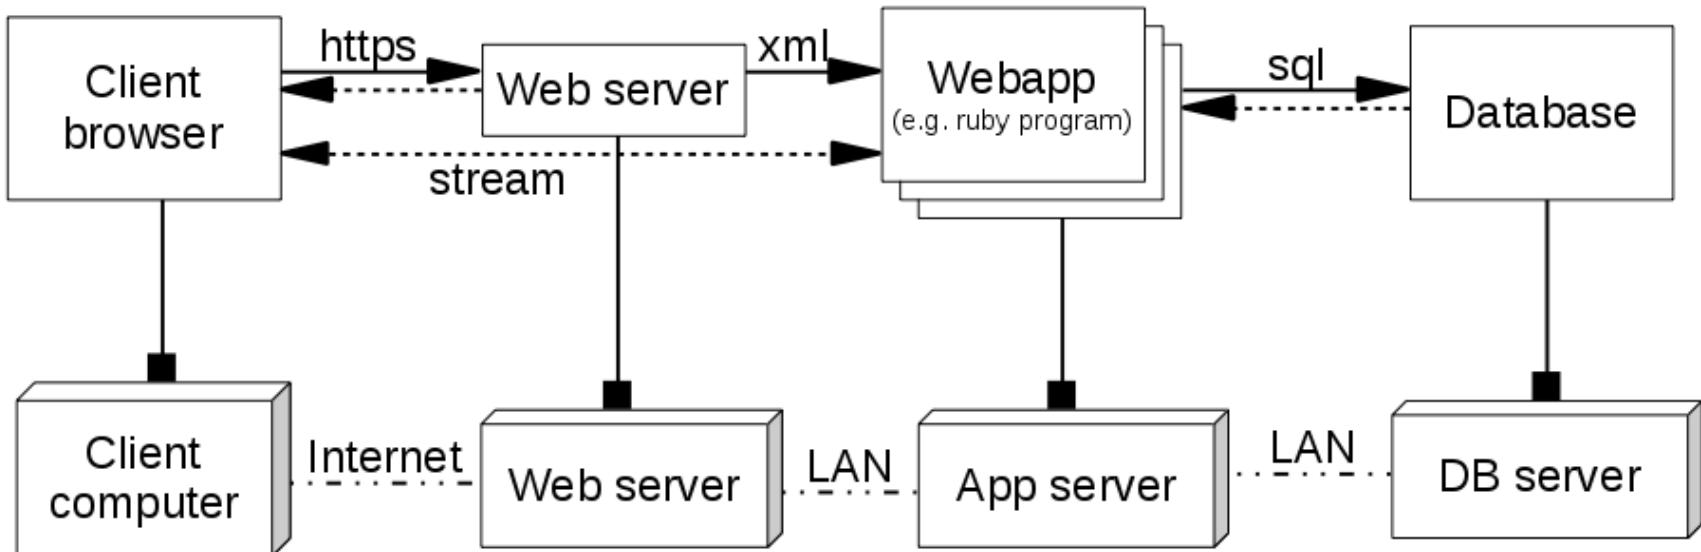
Source: wikimedia commons

Remote machine

Layered system



Tiered architecture



Architectural Style?

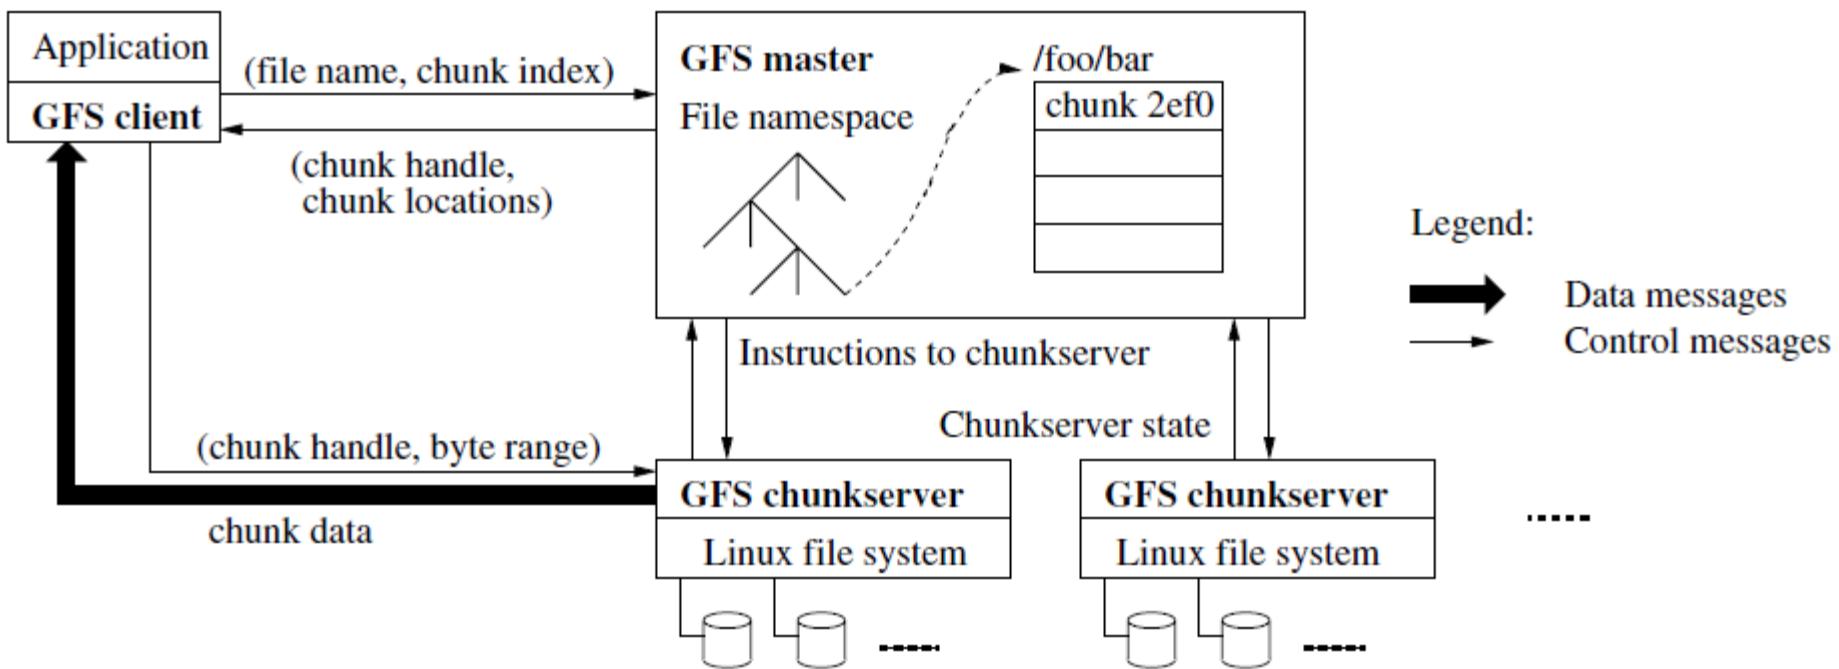


Figure 1: GFS Architecture

Ghemawat, Sanjay, Howard Gobioff, and Shun-Tak Leung. "The Google file system." *ACM SIGOPS operating systems review*. Vol. 37. No. 5. ACM, 2003.

Tactics

- Architectural techniques to achieve qualities
 - More tied to specific context and quality
- Smaller scope than architectural patterns
 - Problem solved by patterns: “How do I structure my (sub)system?”
 - Problem solved by tactics: “How do I get better at quality X?”
- Collection of common strategies and known solutions
 - Resemble OO design patterns

Prioritization
Concurrency
Nondeterminism
Encryption
Separation
Undo
Authentication
Timestamps
Sanitation

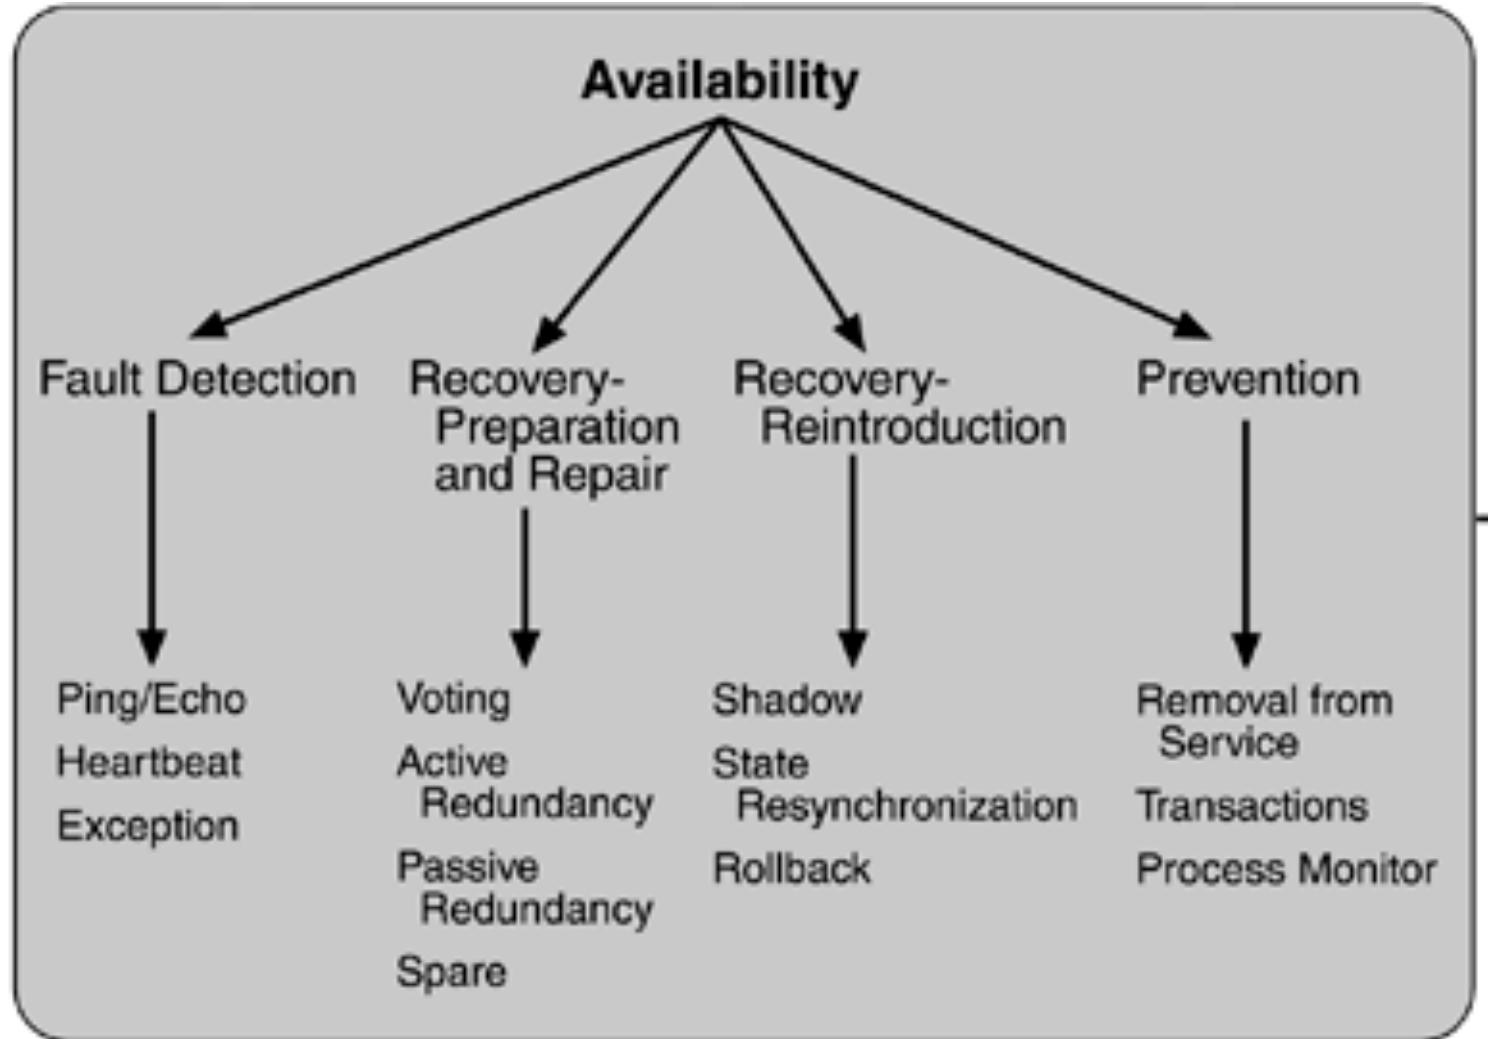
Redundancy
Authorization
Synchronization
Transactions
Coupling
Auditing
Sandboxing
Voting
Cohesion

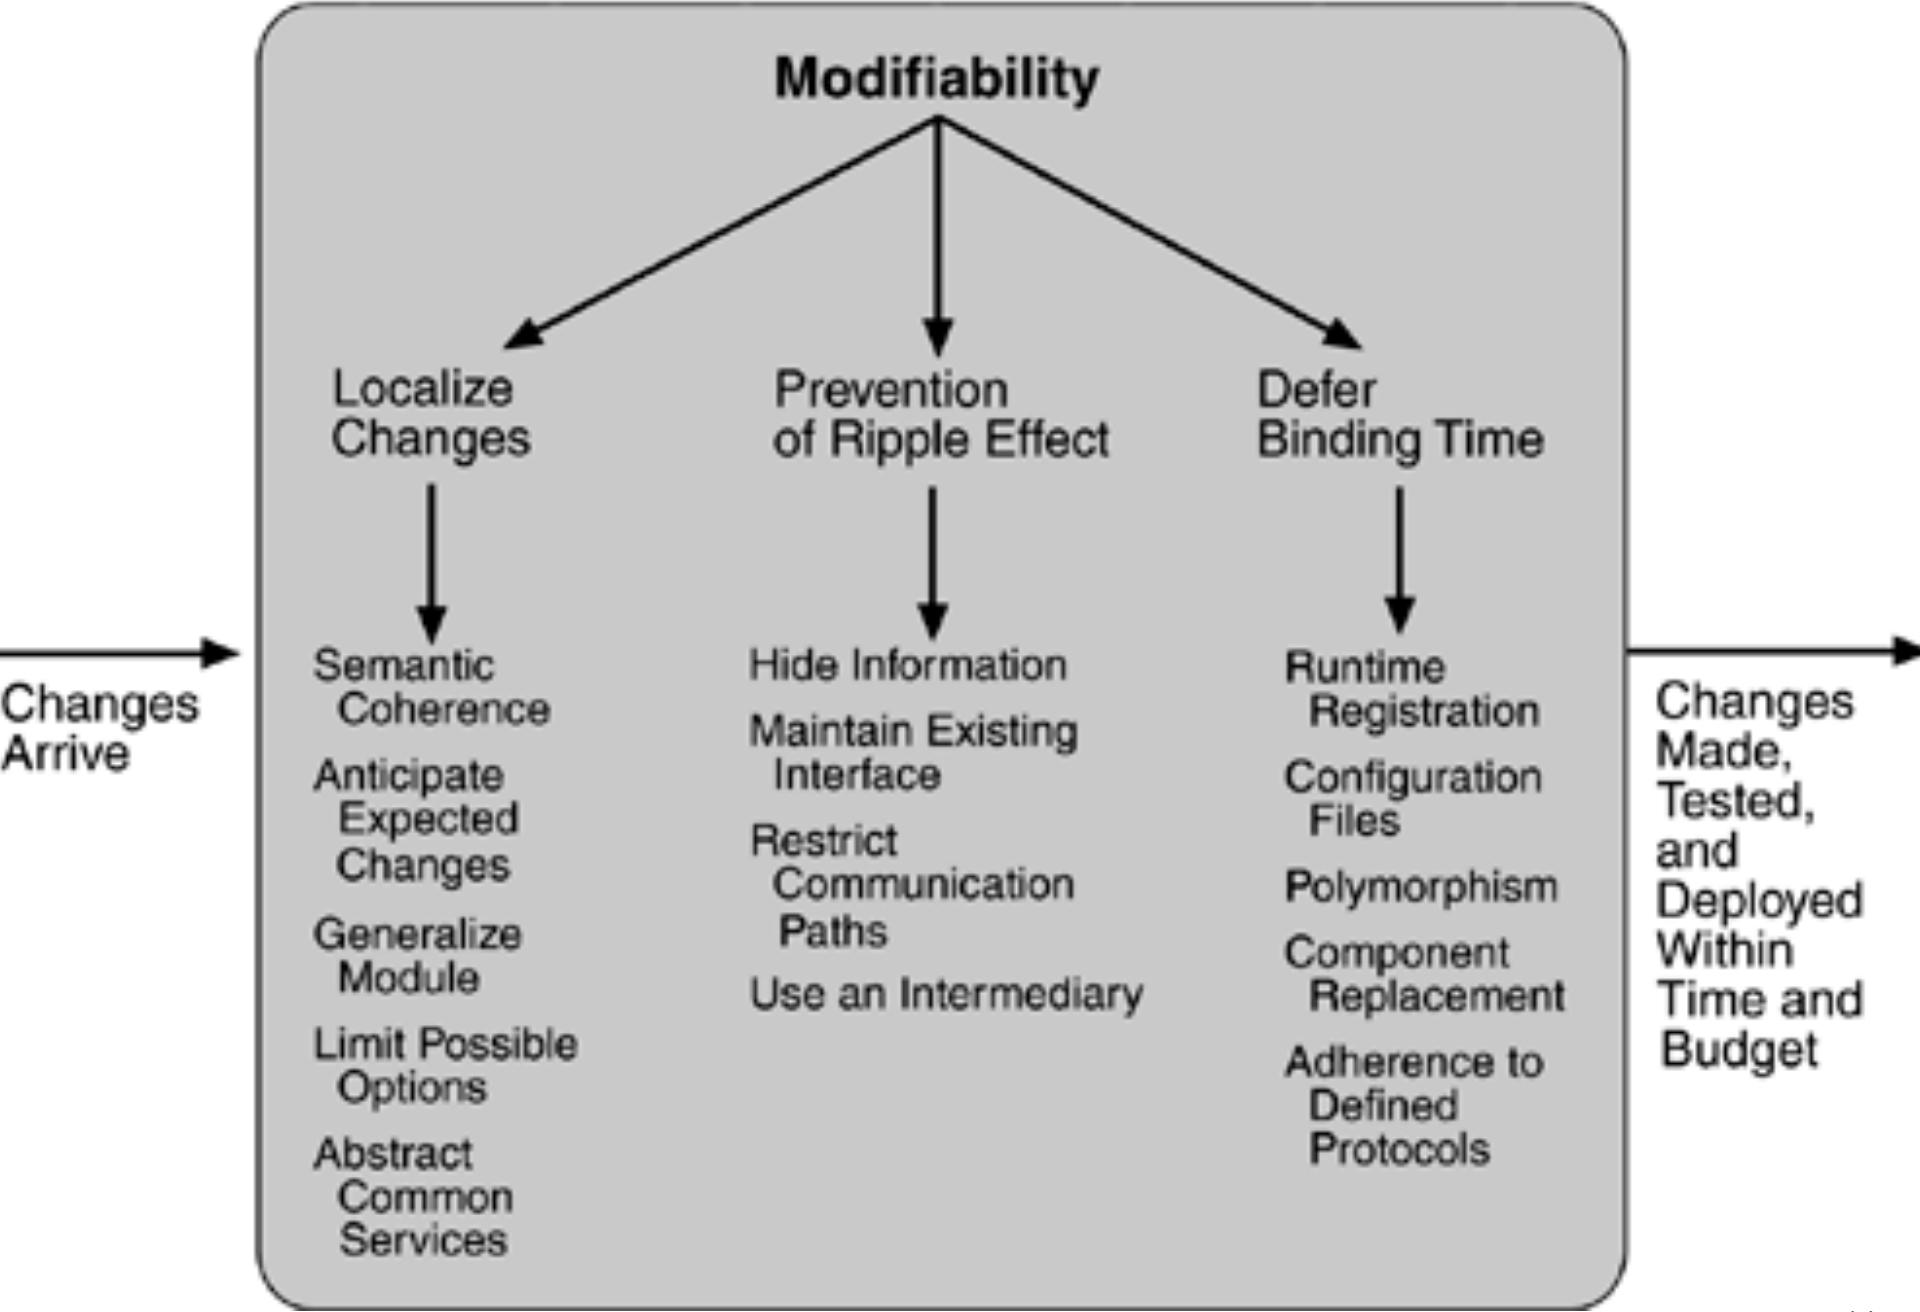
Example Tactic Description: Record/playback

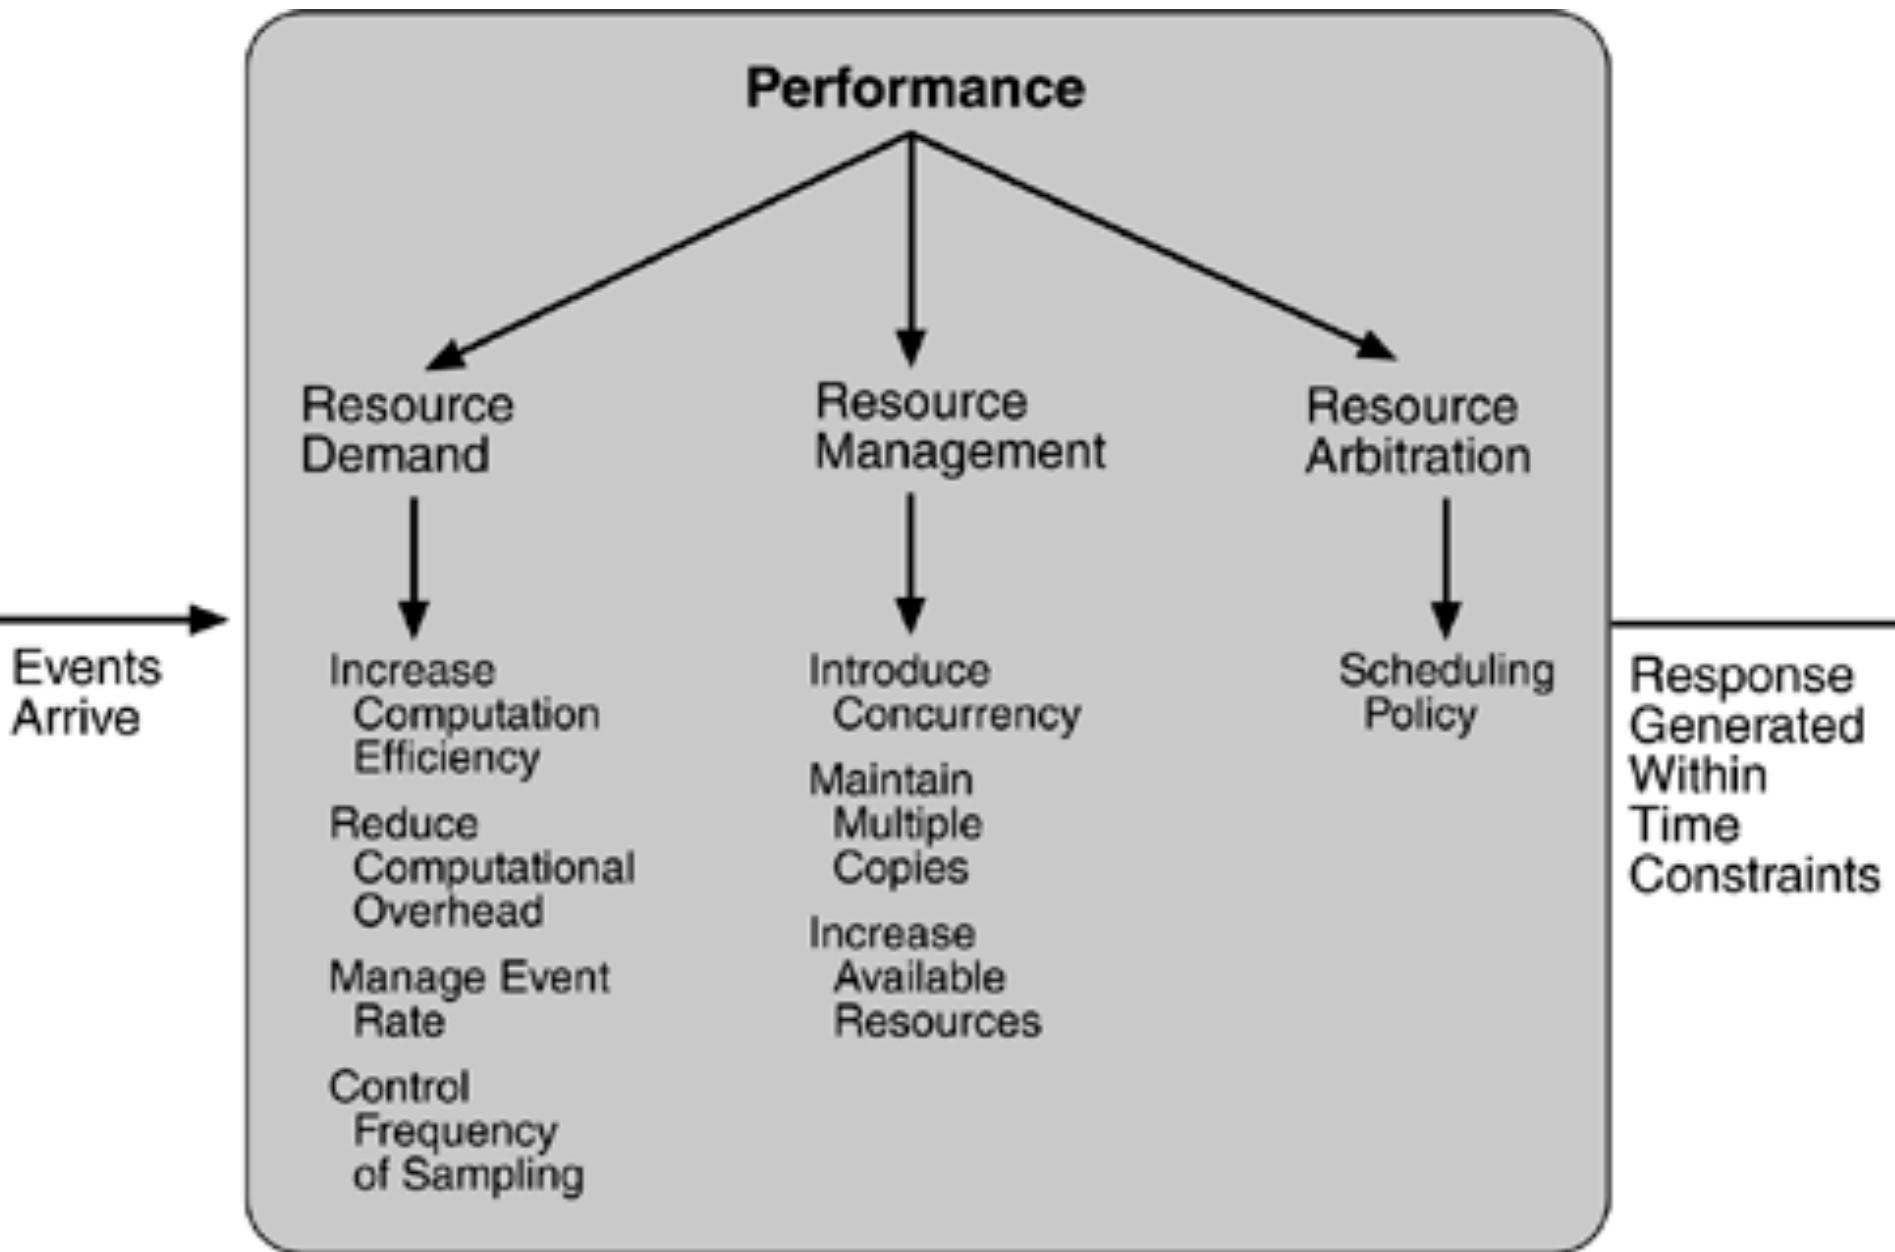
- Record/playback refers to both capturing information crossing an interface and using it as input into the test harness. The information crossing an interface during normal operation is saved in some repository and represents output from one component and input to another. Recording this information allows test input for one of the components to be generated and test output for later comparison to be saved.

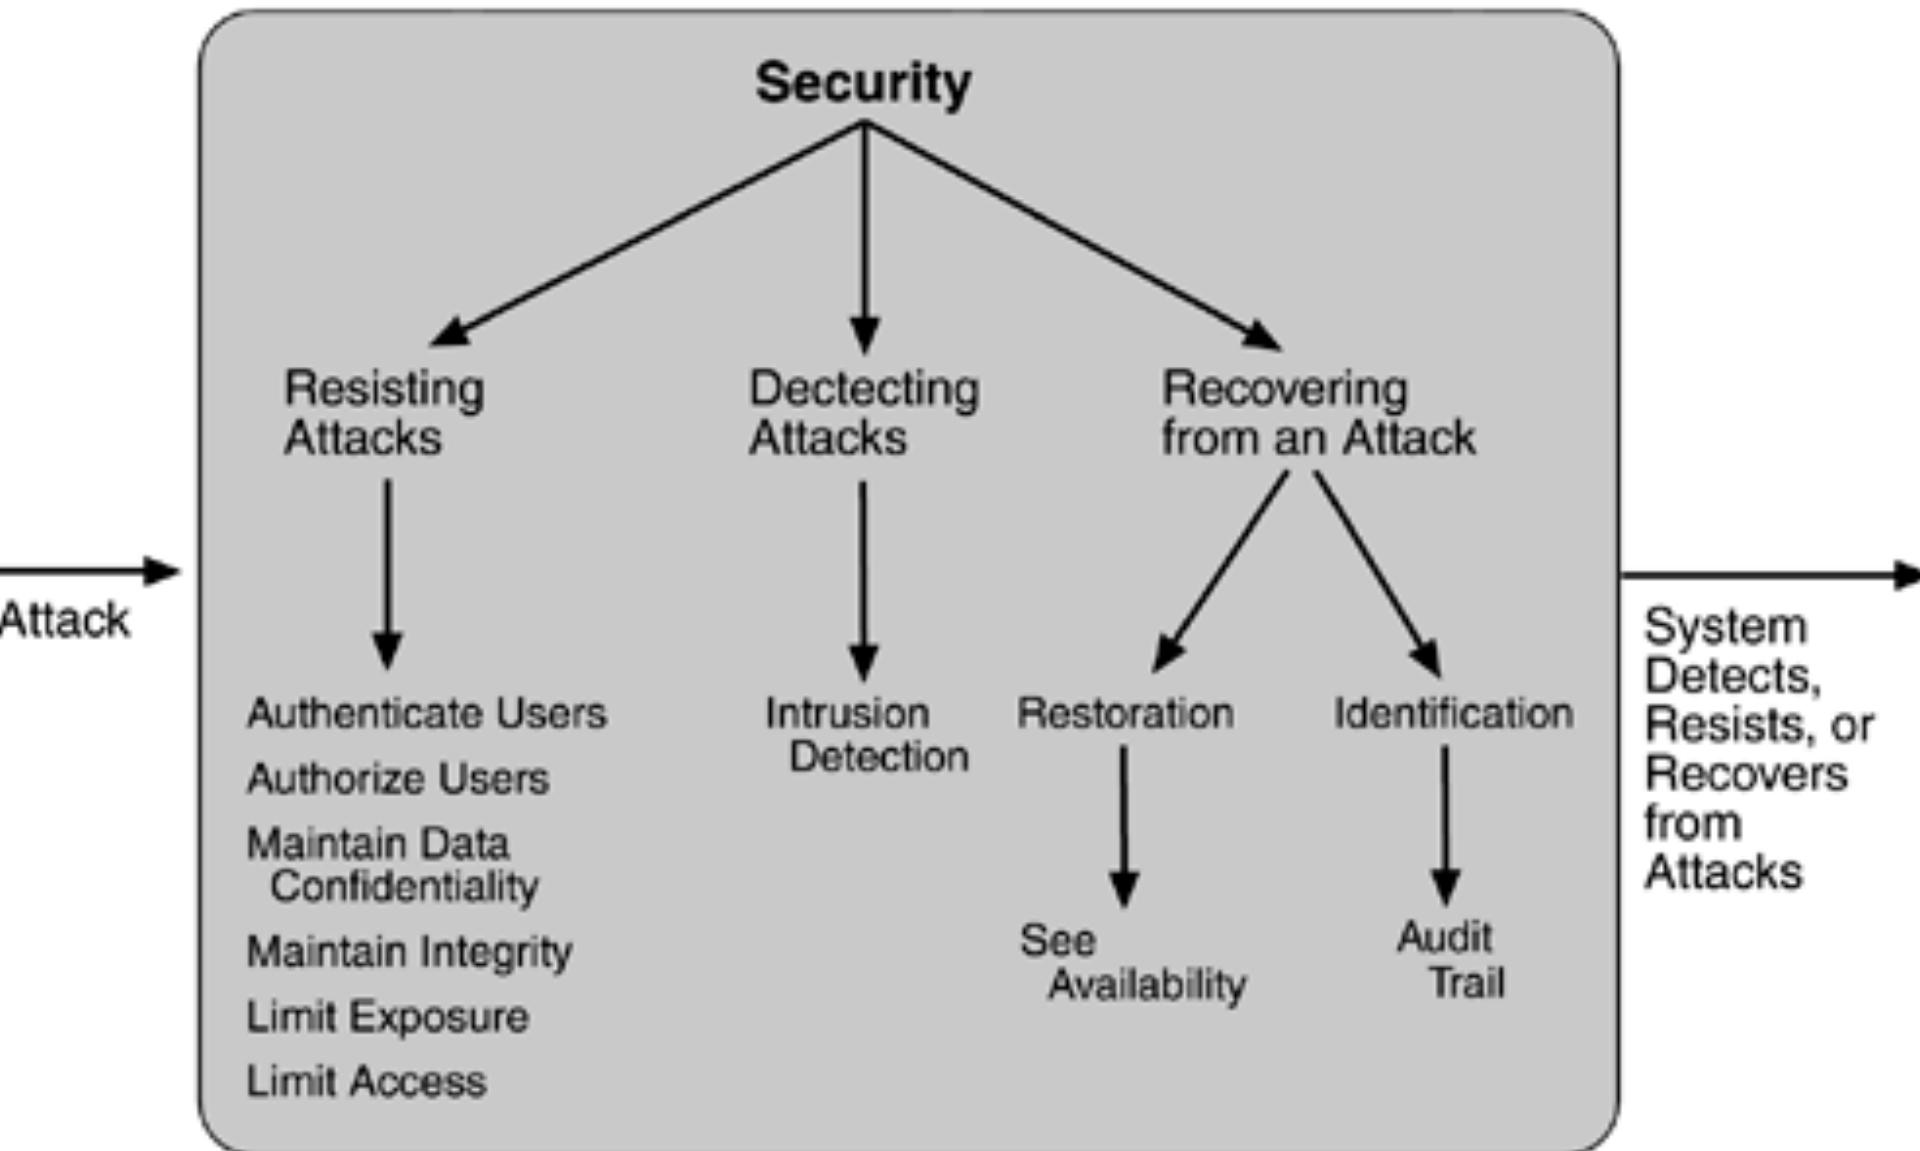
Example Tactic Description: Built-in monitors

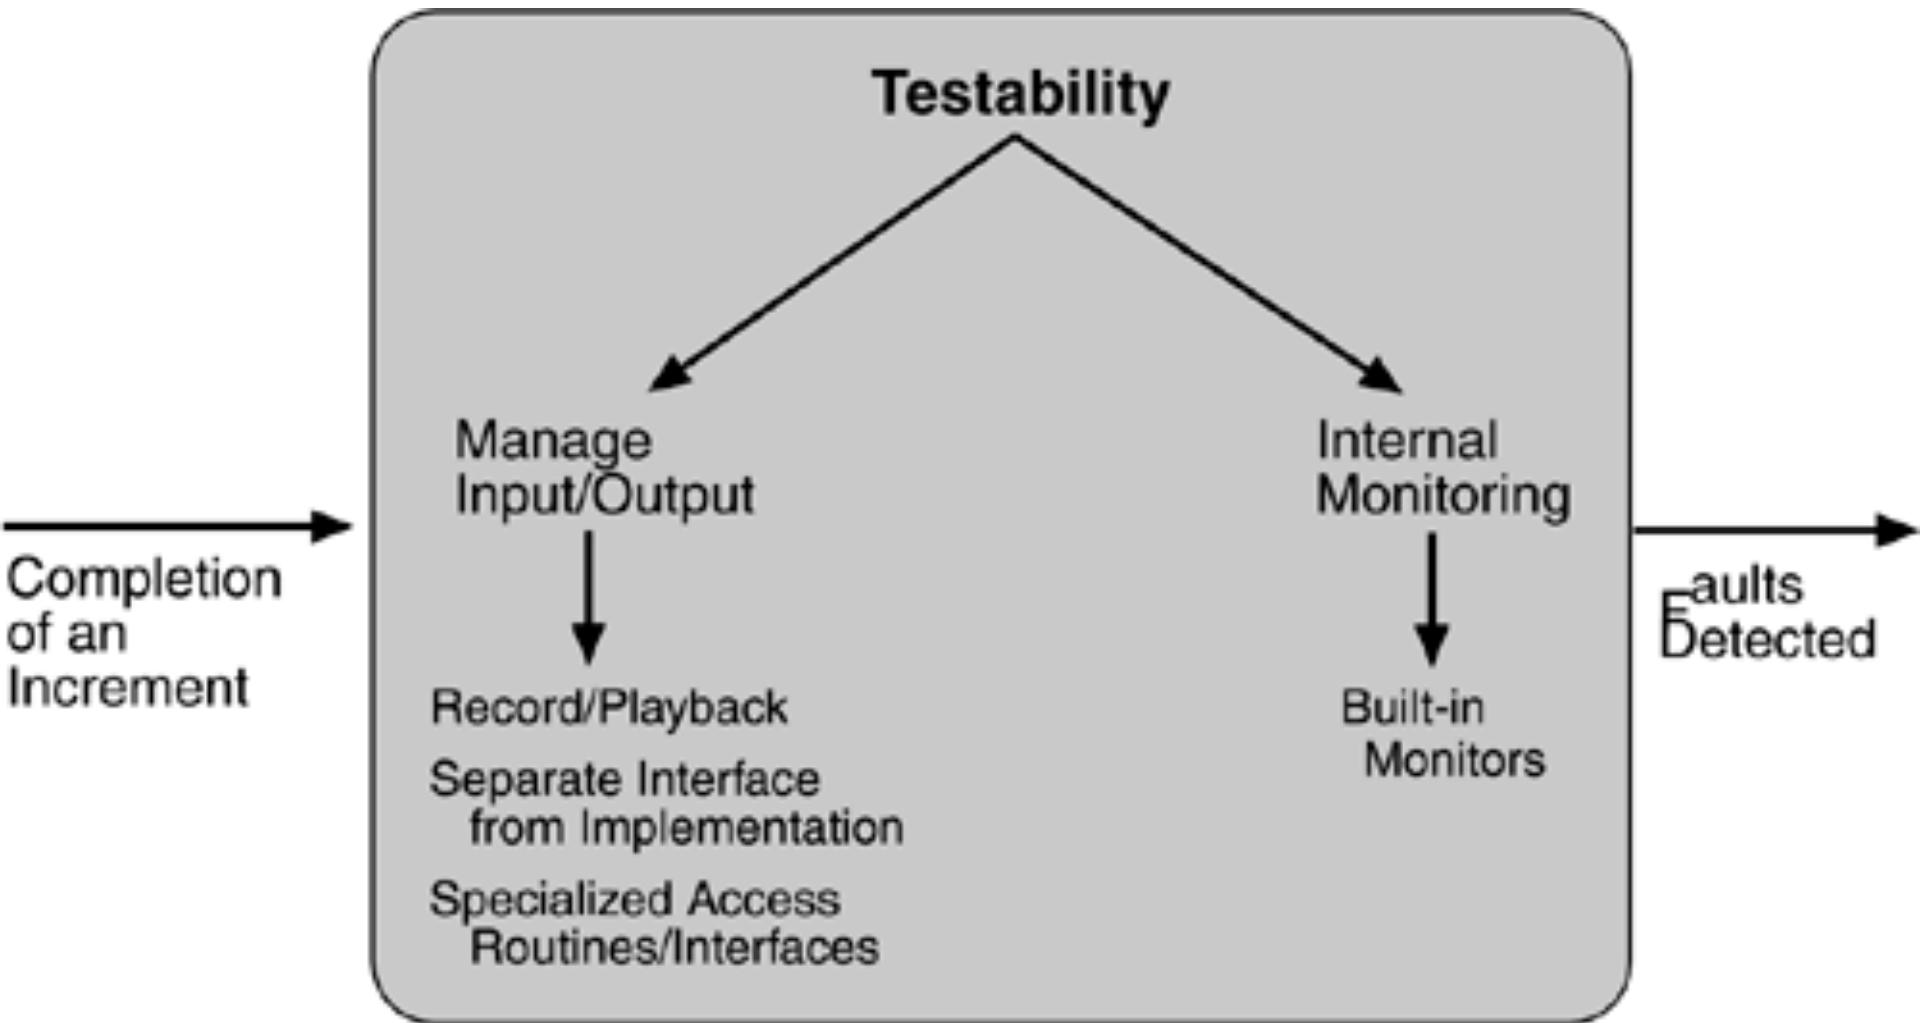
- The component can maintain state, performance load, capacity, security, or other information accessible through an interface. This interface can be a permanent interface of the component or it can be introduced temporarily via an instrumentation technique such as aspect-oriented programming or preprocessor macros. A common technique is to record events when monitoring states have been activated. Monitoring states can actually increase the testing effort since tests may have to be repeated with the monitoring turned off. Increased visibility into the activities of the component usually more than outweigh the cost of the additional testing.

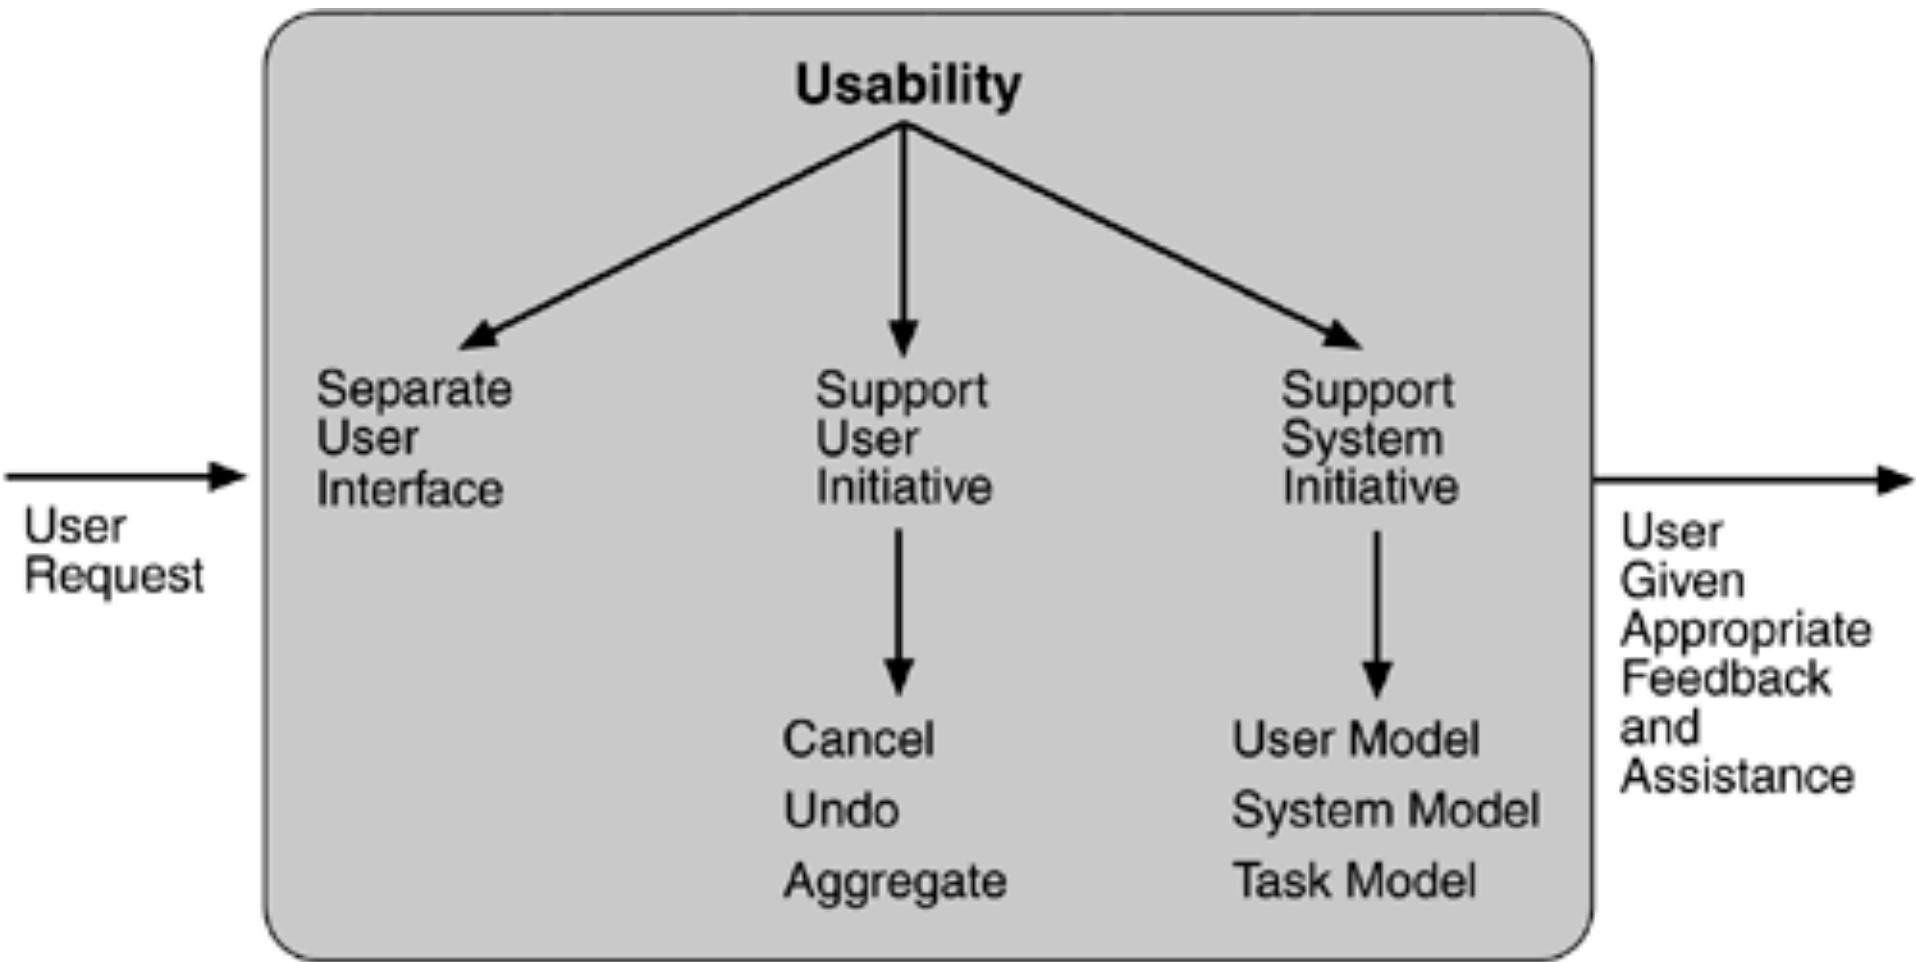


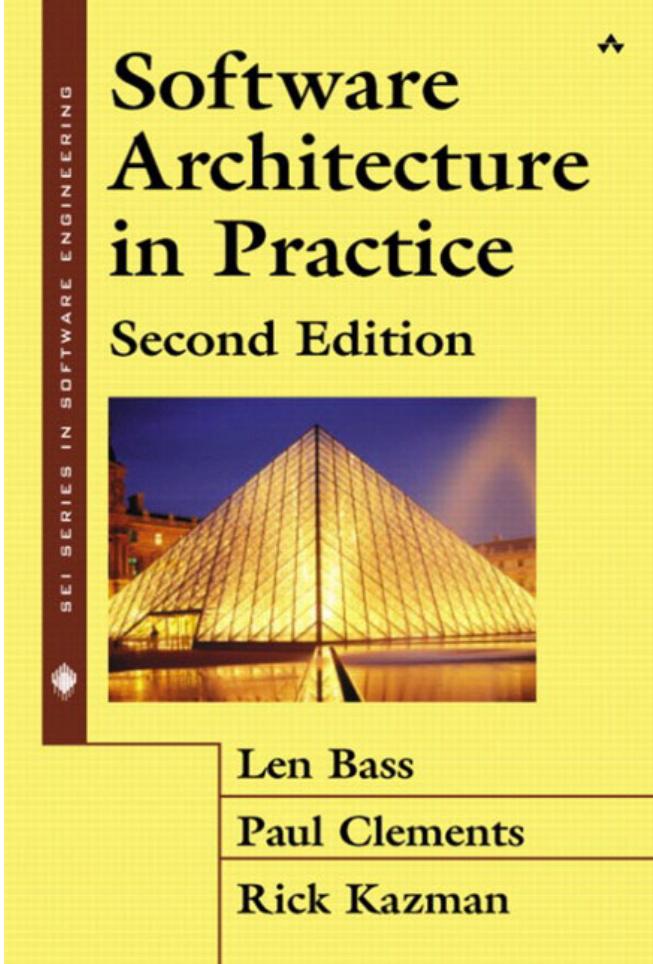












Many tactics
described in Chapter
5

Brief high-level
descriptions (about 1
paragraph per tactic)

Second and more detailed third edition available as ebook

through CMU library.



institute for
SOFTWARE
RESEARCH

Carnegie Mellon University
School of Computer Science

Summary

Architecture as
structures and relations

- Patterns
- Tactics

Architecture as
documentation

- Views
- Rationale

Architecture as process

- Decisions
- Evaluation
- Reconstruction
- Agile

Further Readings

- Bass, Clements, and Kazman. Software Architecture in Practice. Addison-Wesley, 2003.
- Boehm and Turner. Balancing Agility and Discipline: A Guide for the Perplexed, 2003.
- Clements, Bachmann, Bass, Garlan, Ivers, Little, Merson, Nord, Stafford. Documenting Software Architectures: Views and Beyond, 2010.
- Fairbanks. Just Enough Software Architecture. Marshall & Brainerd, 2010.
- Jansen and Bosch. Software Architecture as a Set of Architectural Design Decisions, WICSA 2005.
- Lattanze. Architecting Software Intensive Systems: a Practitioner's Guide, 2009.
- Sommerville. Software Engineering. Edition 7/8, Chapters 11-13
- Taylor, Medvidovic, and Dashofy. Software Architecture: Foundations, Theory, and Practice. Wiley, 2009.